



NYU | TANDON

# Information/Data Visualization

2D/3D Transformation and  
Scientific Visualization

Qi Sun

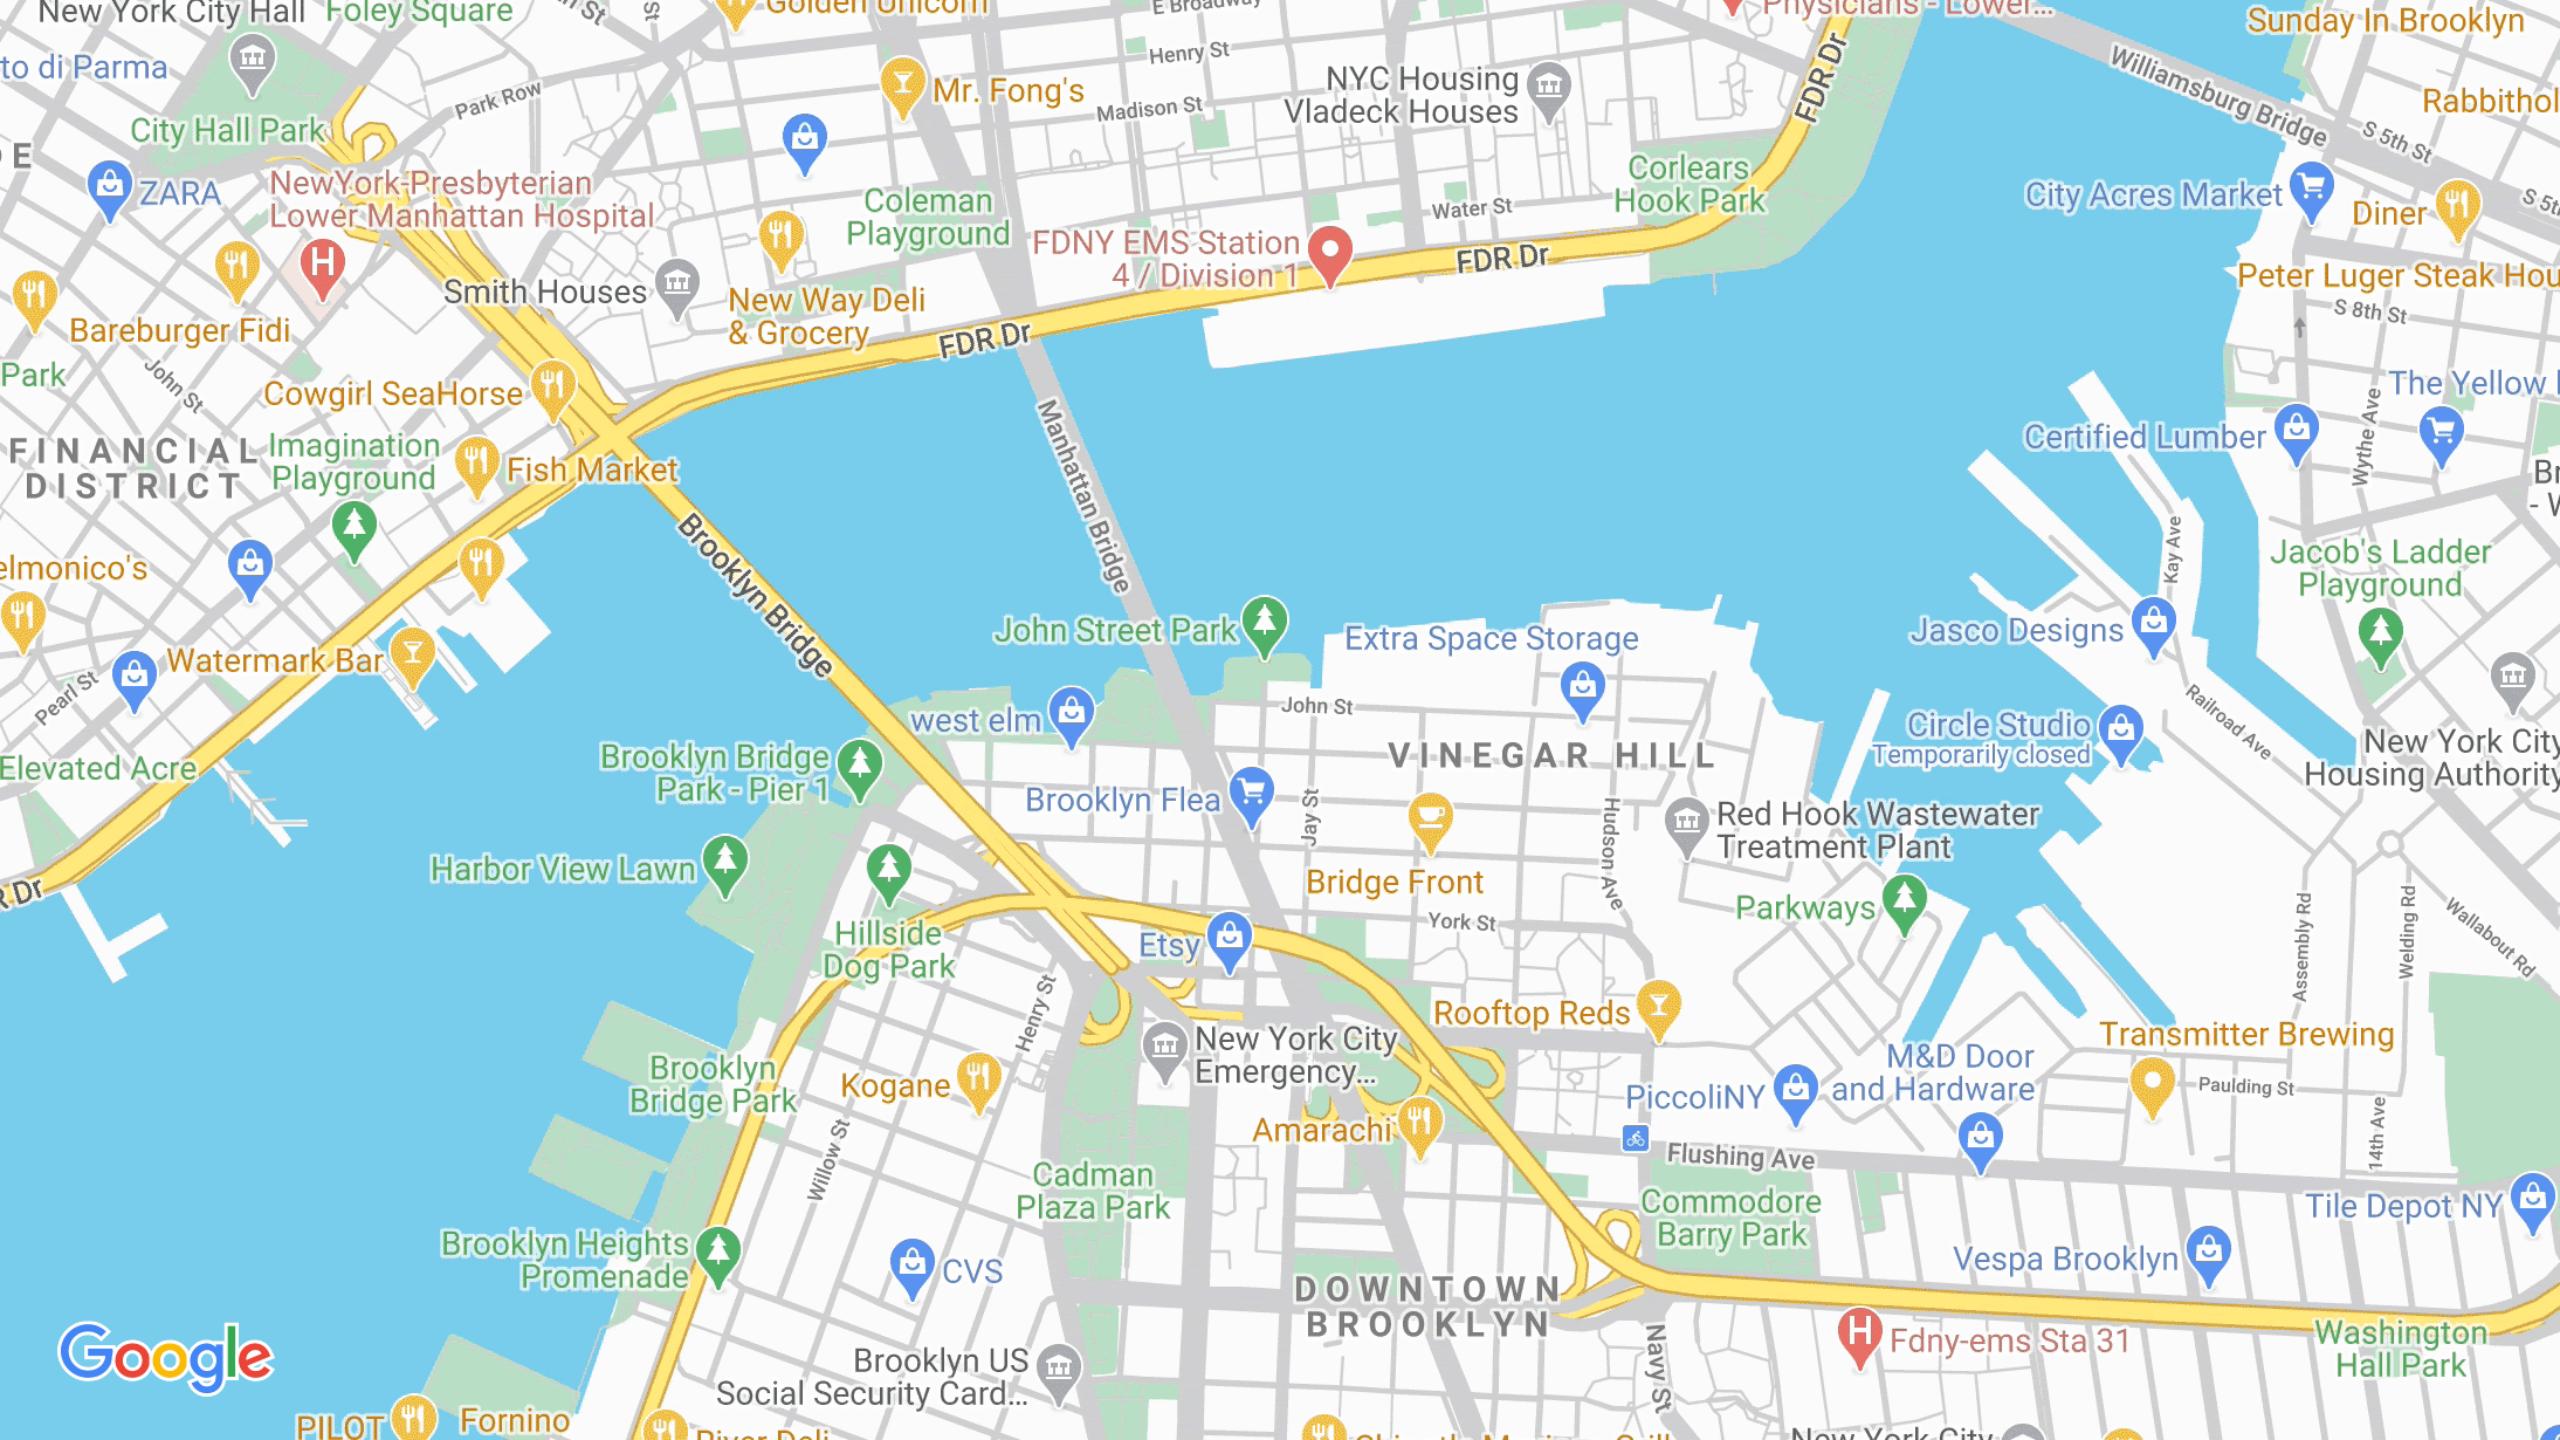
# Important Logistics

- Final Project Presentation
  - 3min each: on-site presentation **vs** pre-recorded video + Q&A
  - engineering volume? complete/end-to-end system >> complexity
  - come to the office hours if you need help!
  - due date: April 30th (video) for May 2nd presentation
- The last guest lecture next week
  - Wenqi (Wendy) Xian (<https://www.cs.cornell.edu/~wenqixian/>) from Cornell Tech
  - AI-based neural video synthesis

P A R T   0 1

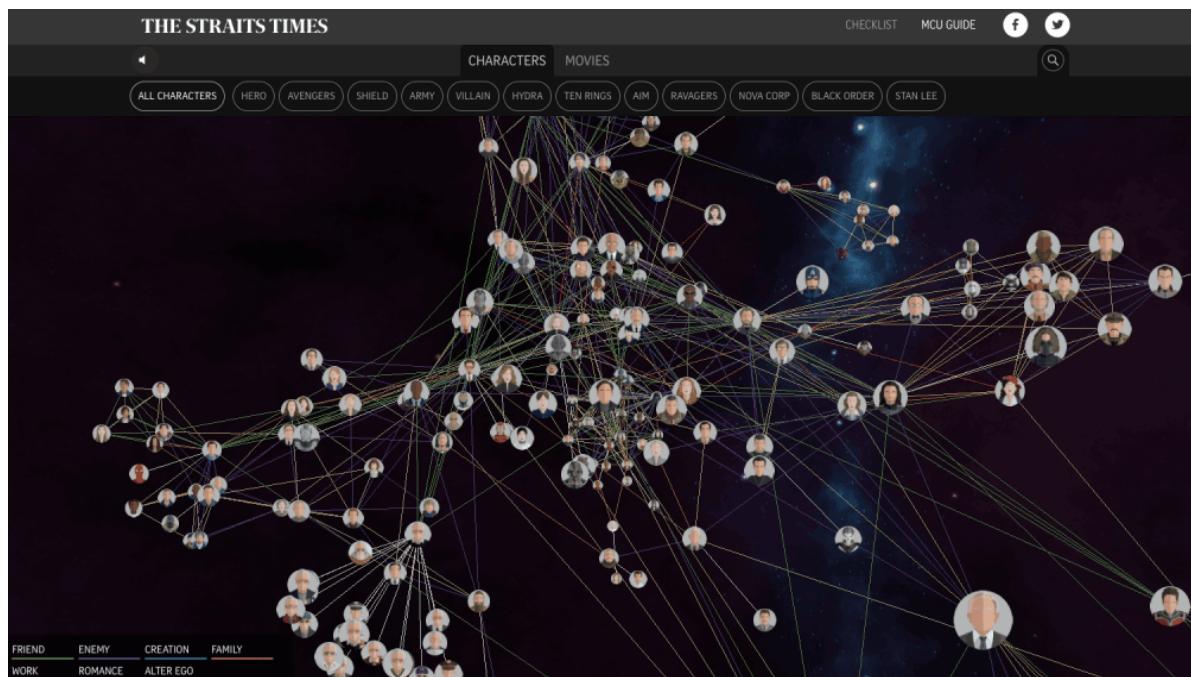
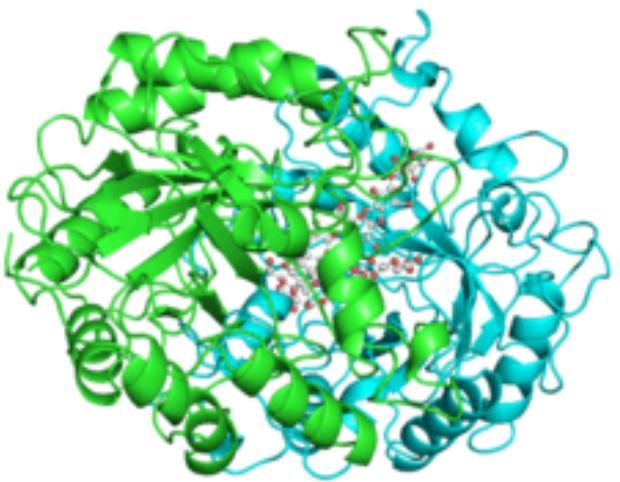
---

# Interaction: **(2D)** Transformation

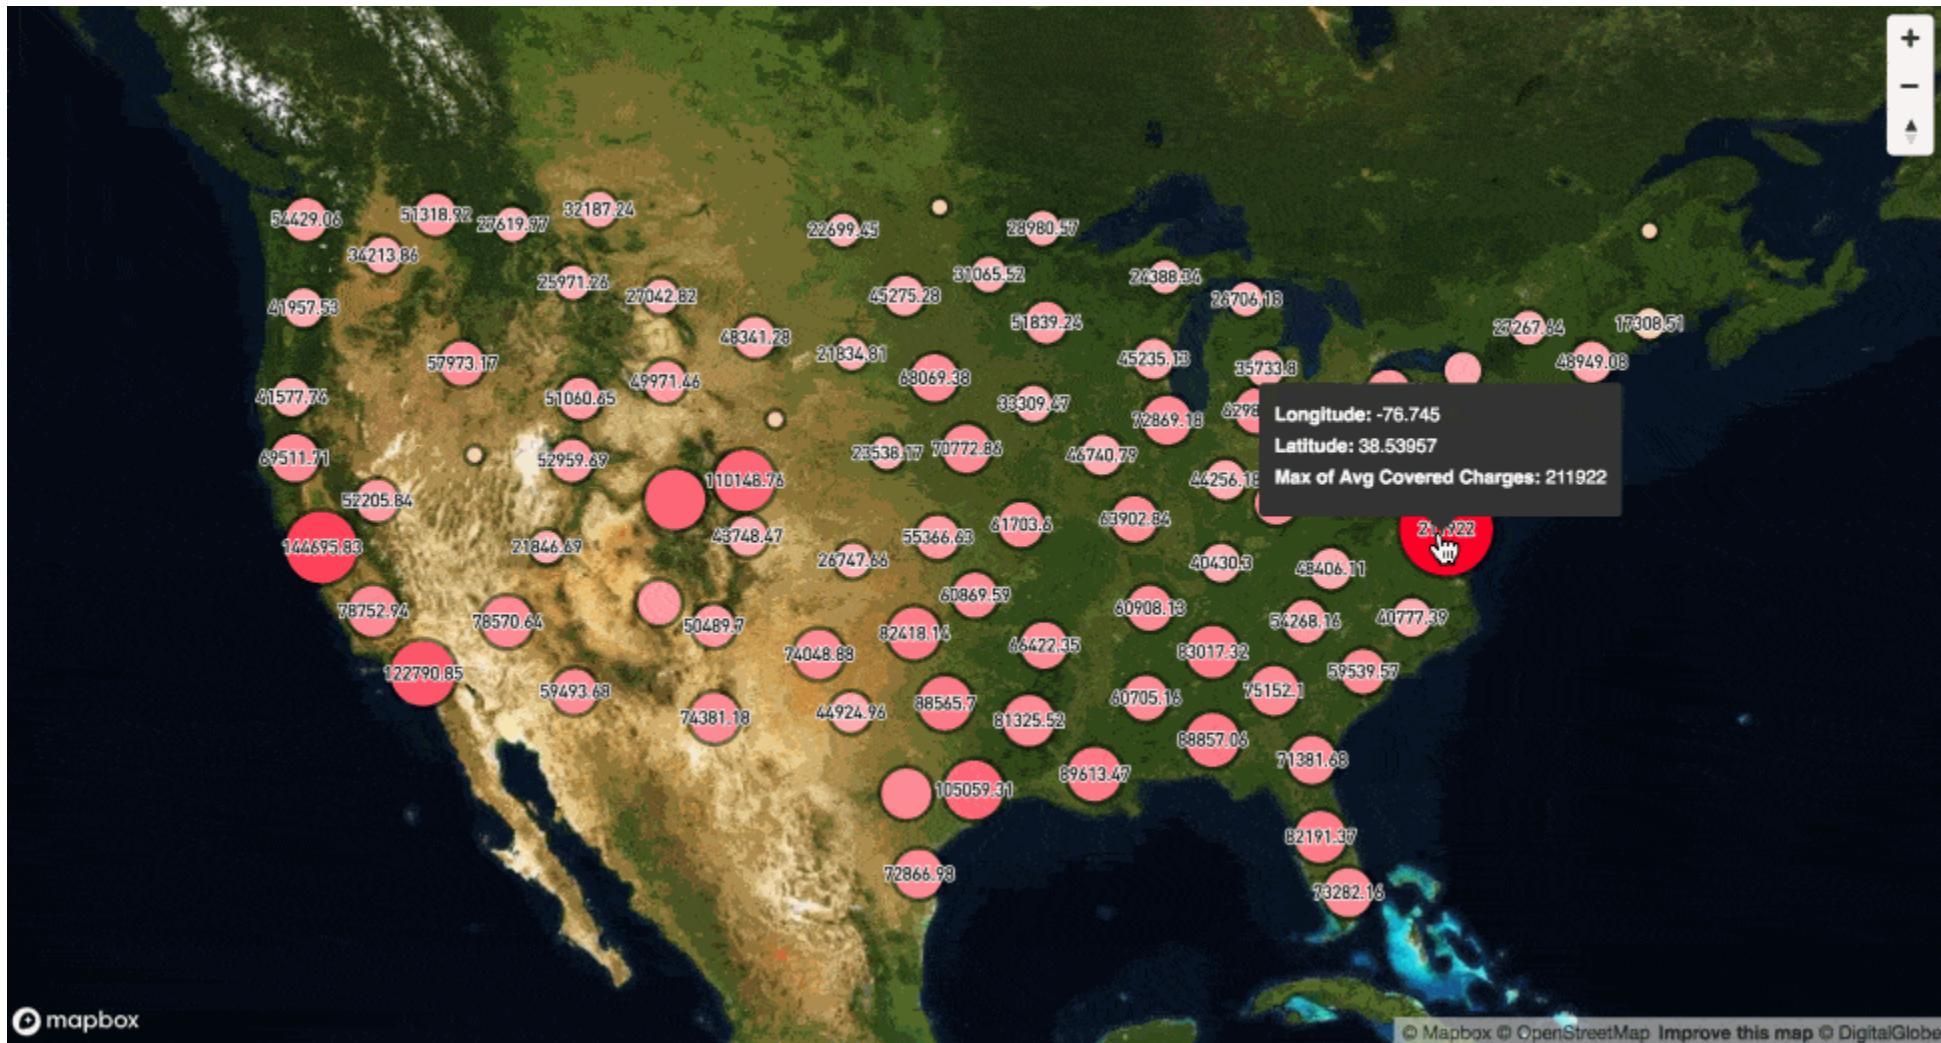


Google

# Transformation - Rotation



# Transformation - Scaling

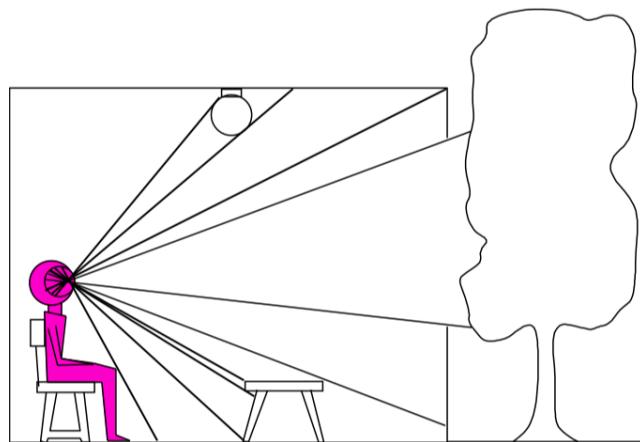


# Transformation – Non-Rigid



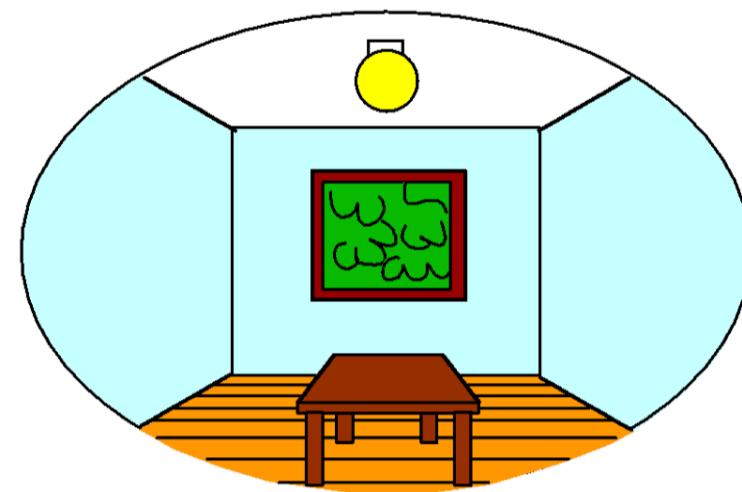
# Transformation

*3D world*



Point of observation

*2D image*

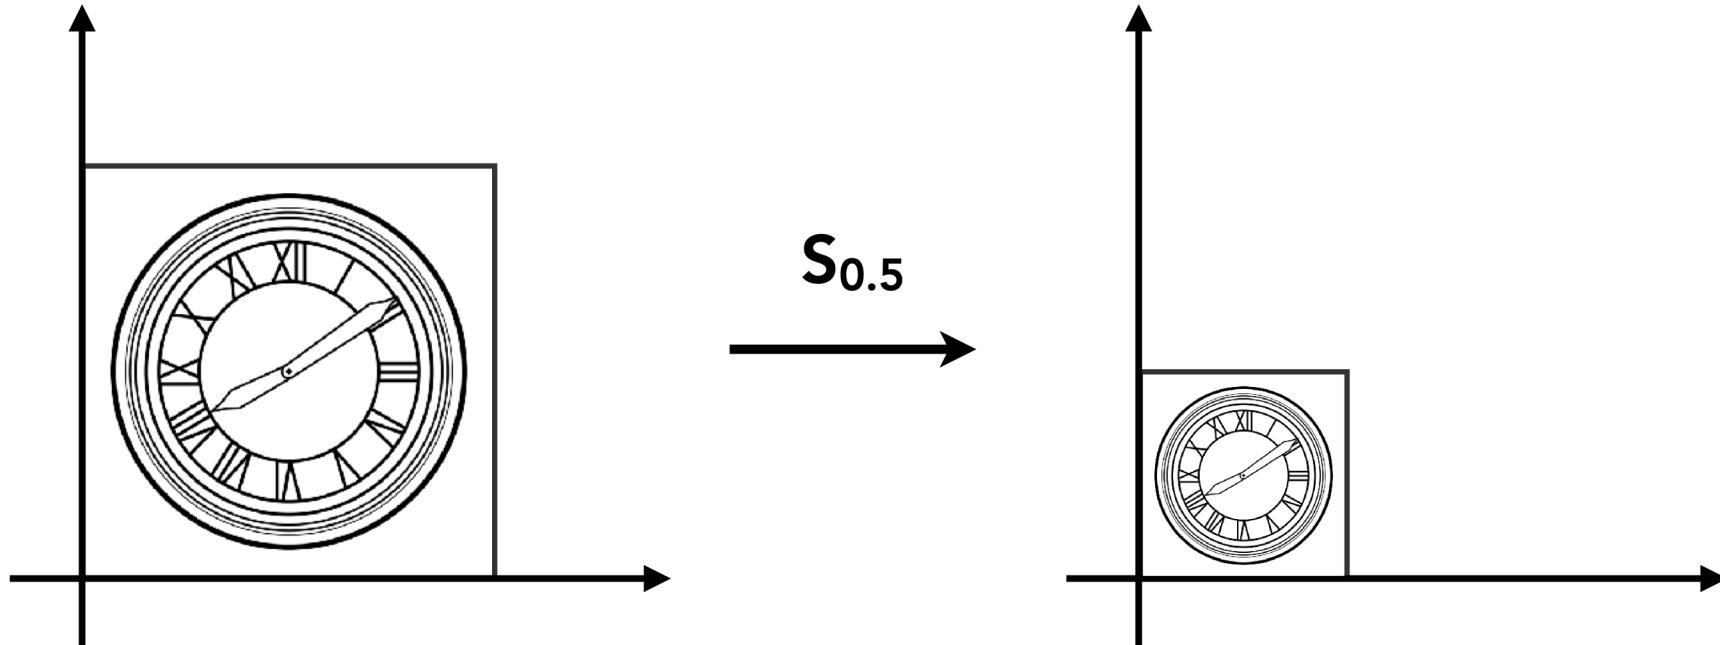


Figures © Stephen E. Palmer, 2002

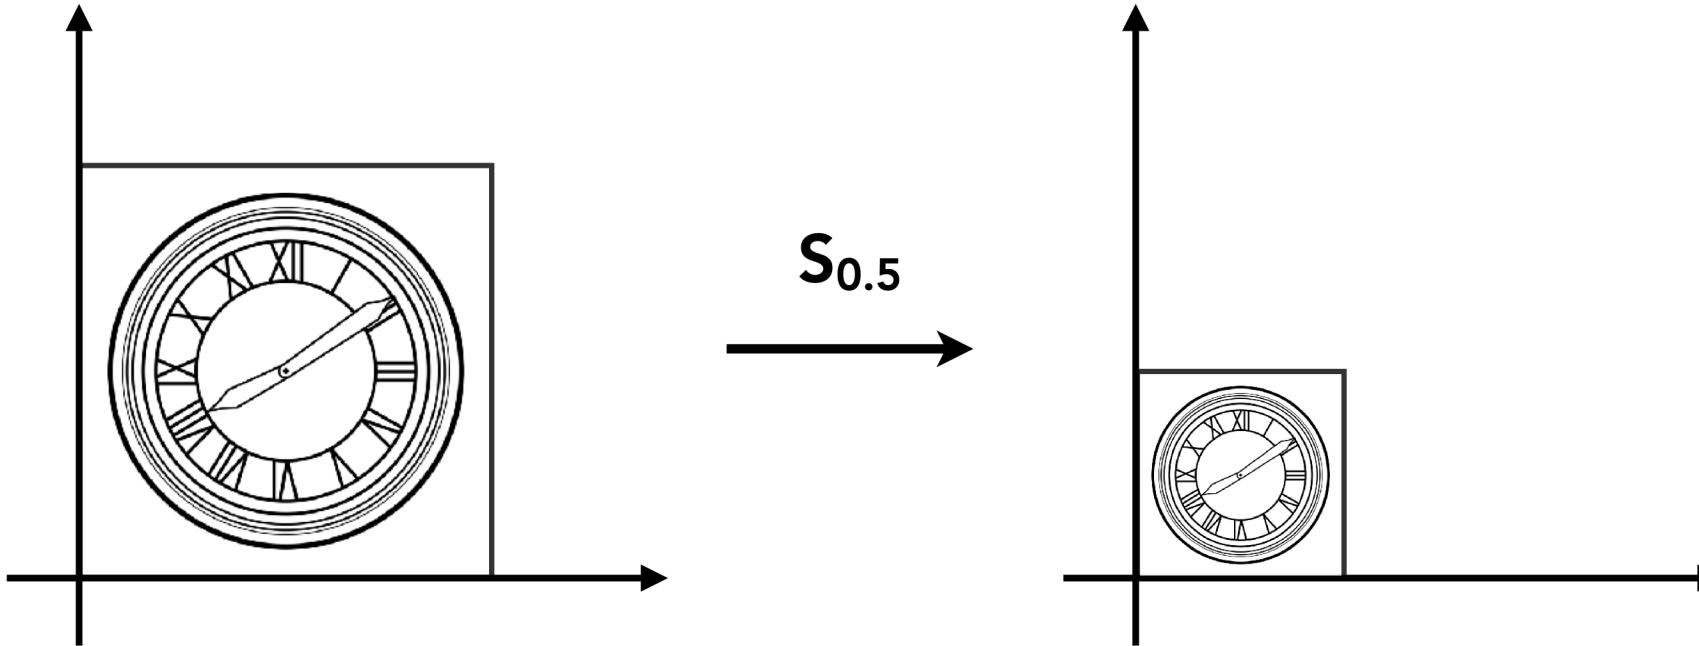
Viewing: (3D to 2D) projection

- 2D transformations
  - Representing transformations using matrices
  - Rotation, scale, shear
- Homogeneous coordinates

# Scale



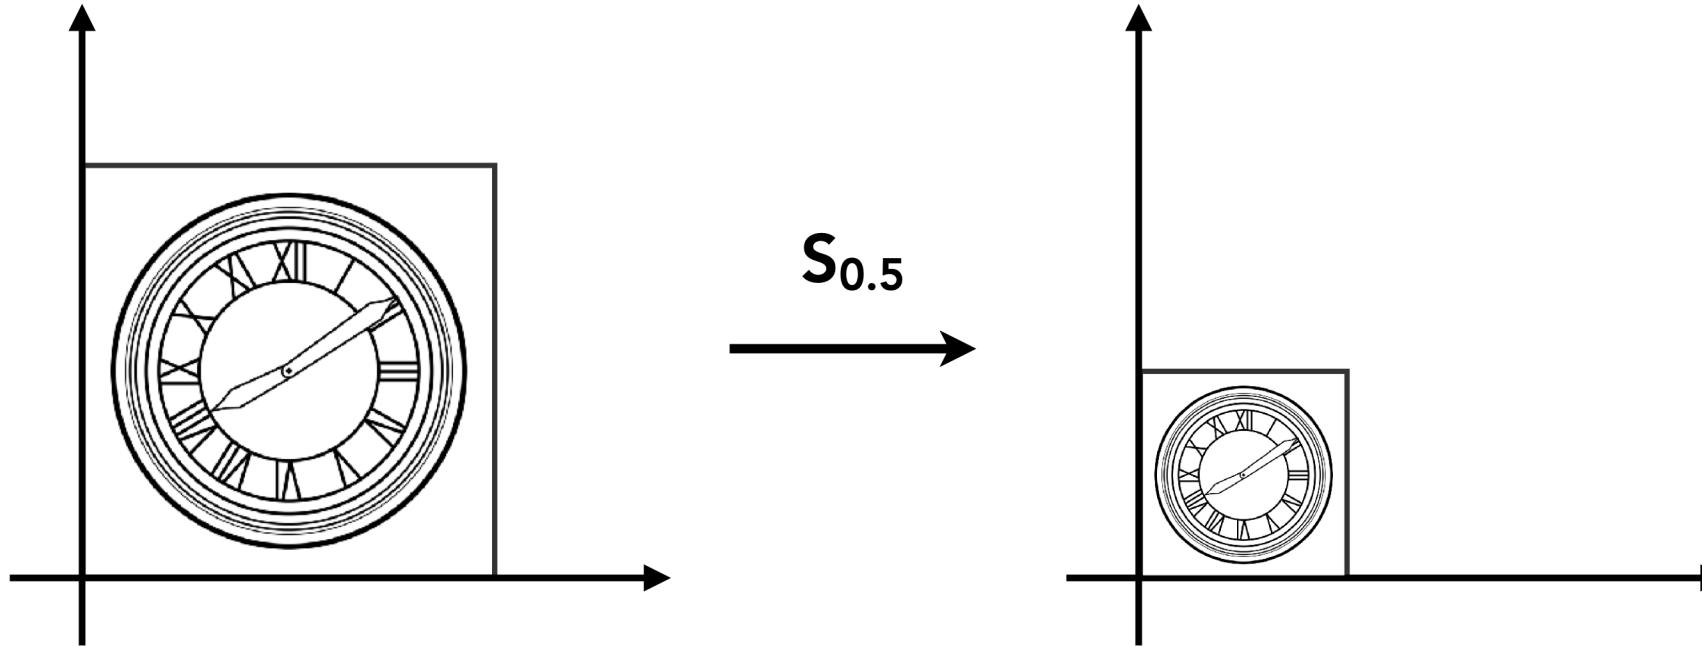
# Scale Transform



$$x' = sx$$

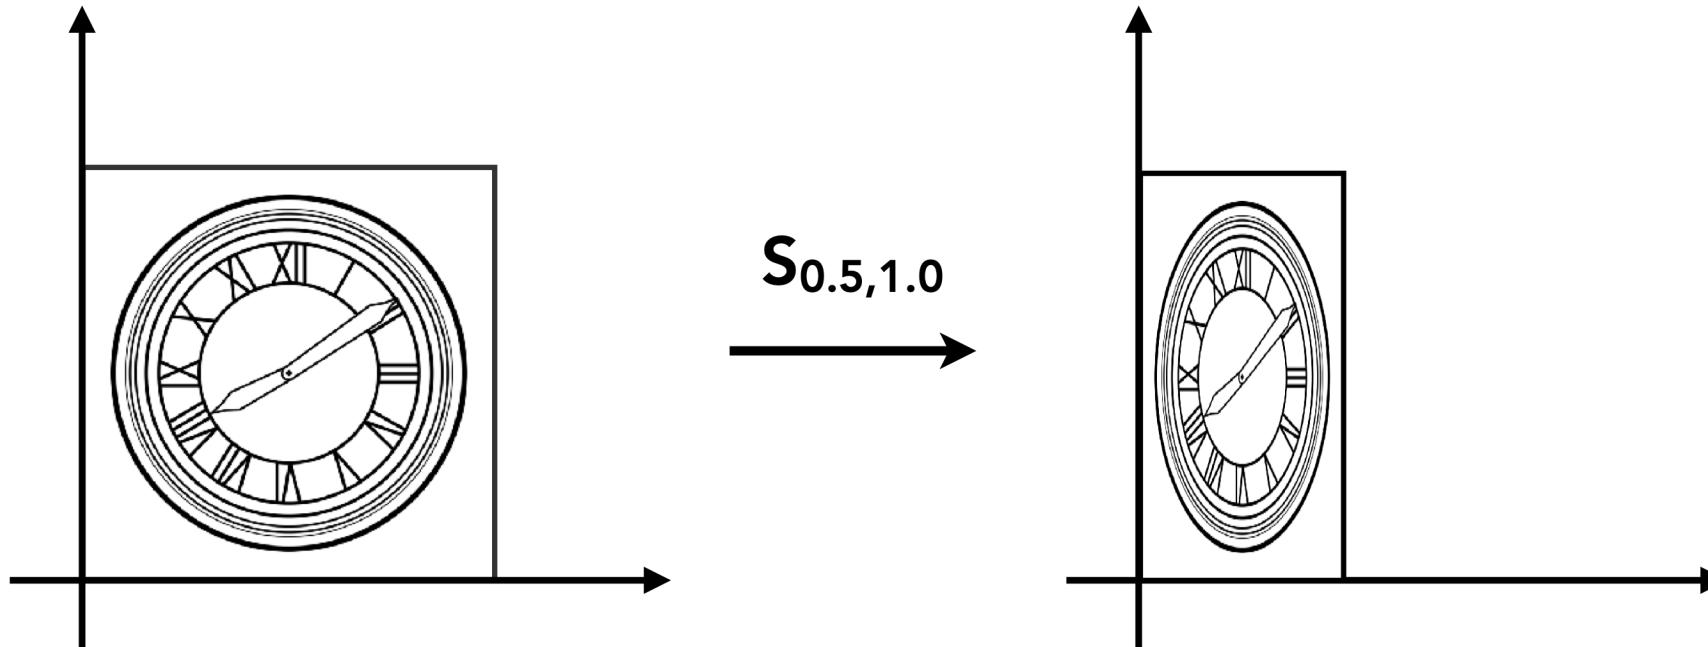
$$y' = sy$$

# Scale Matrix



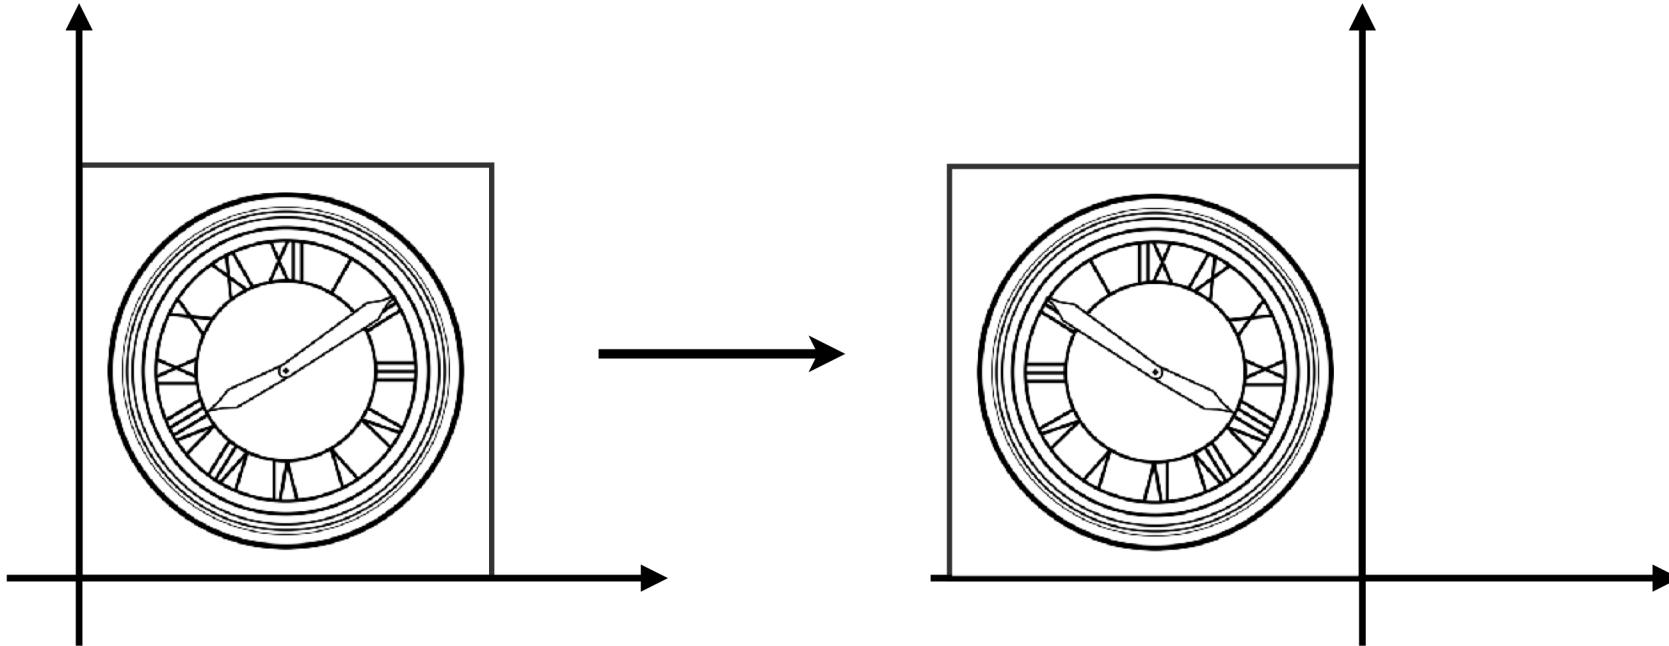
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s & 0 \\ 0 & s \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

# Scale (Non-Uniform)



$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

# Reflection Matrix



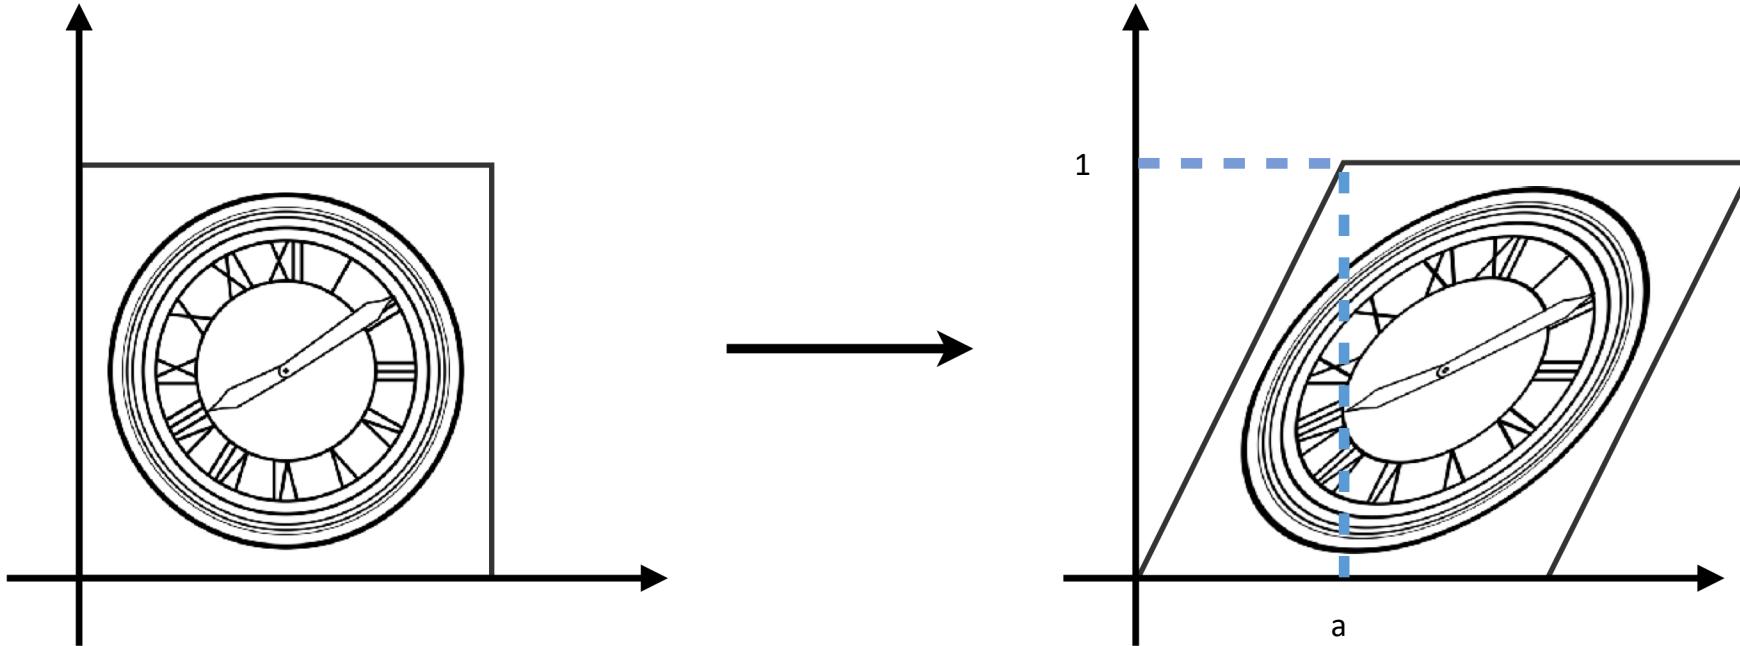
Horizontal reflection:

$$x' = -x$$

$$y' = y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

# Shear Matrix



Hints:

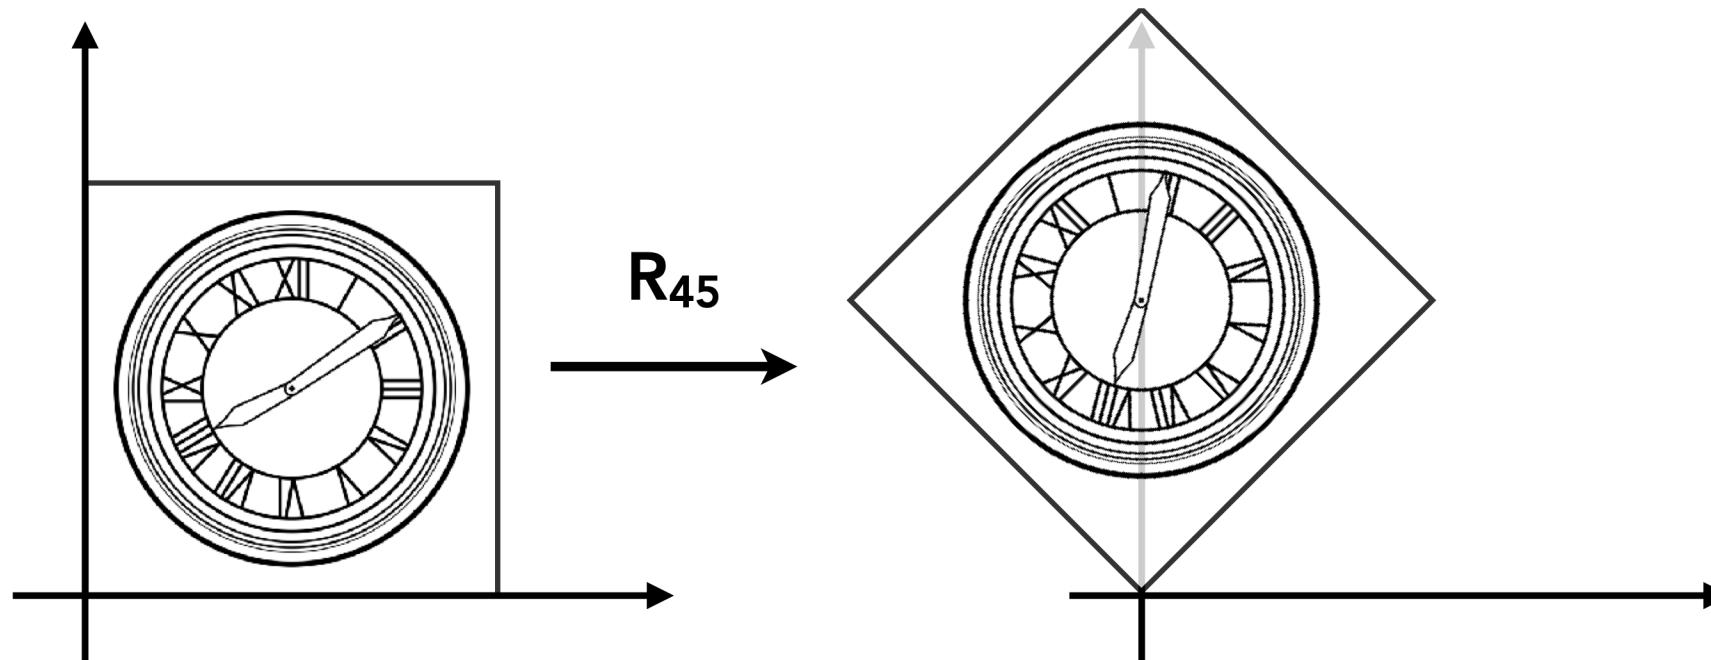
Horizontal shift is 0 at  $y=0$

Horizontal shift is  $a$  at  $y=1$

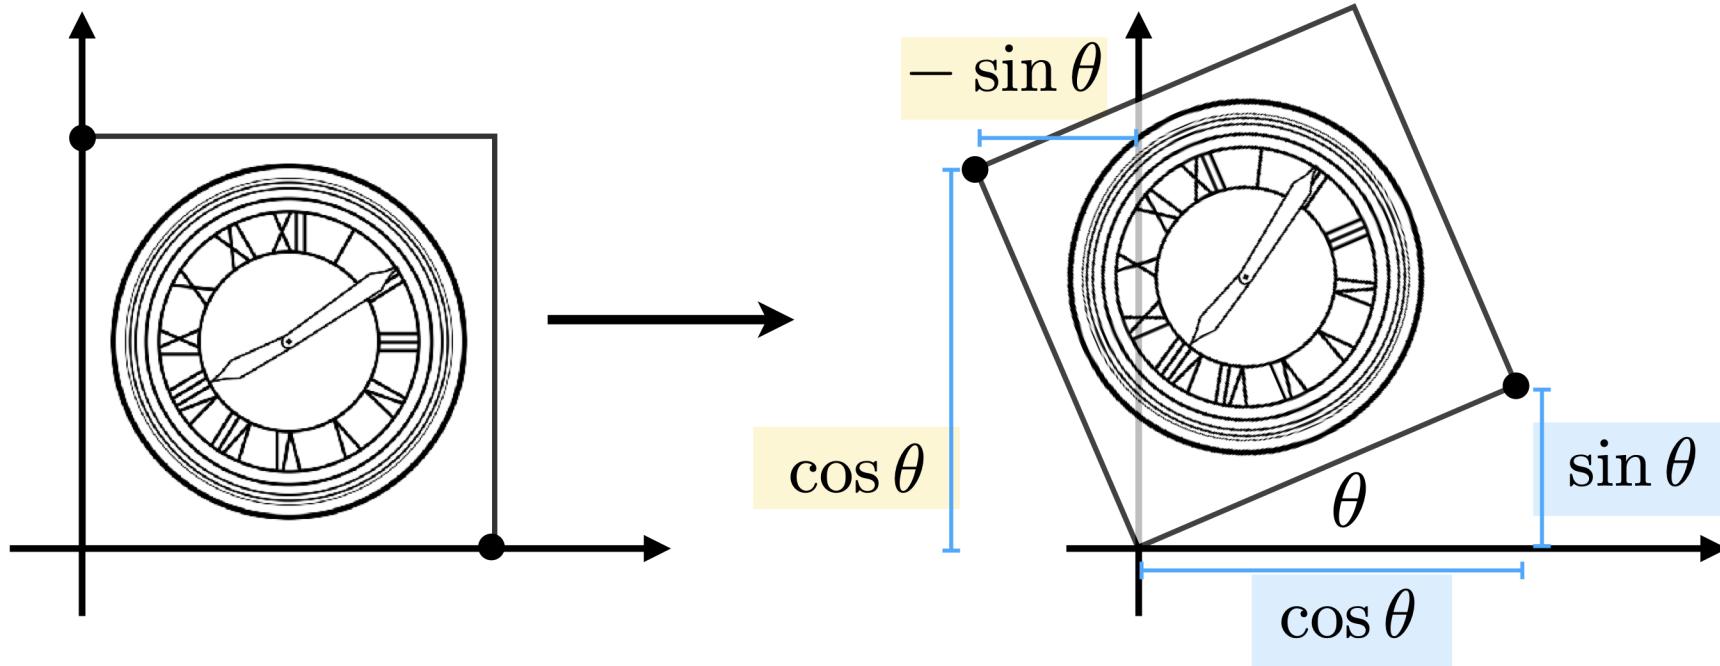
Vertical shift is always 0

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & a \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

# Rotate



# Rotation Matrix



$$\mathbf{R}_\theta = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

# Linear Transforms = Matrices

(of the same dimension)

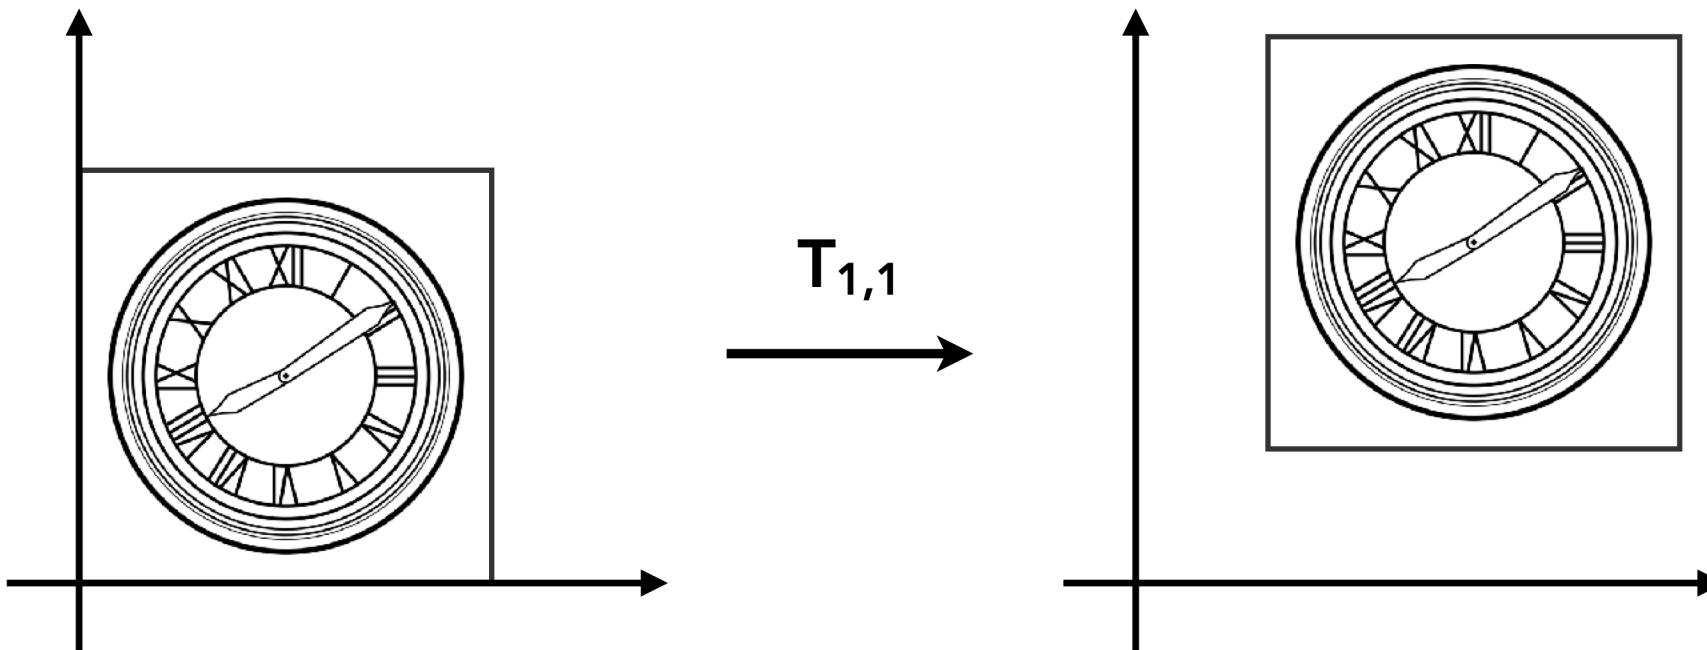
$$x' = a x + b y$$

$$y' = c x + d y$$

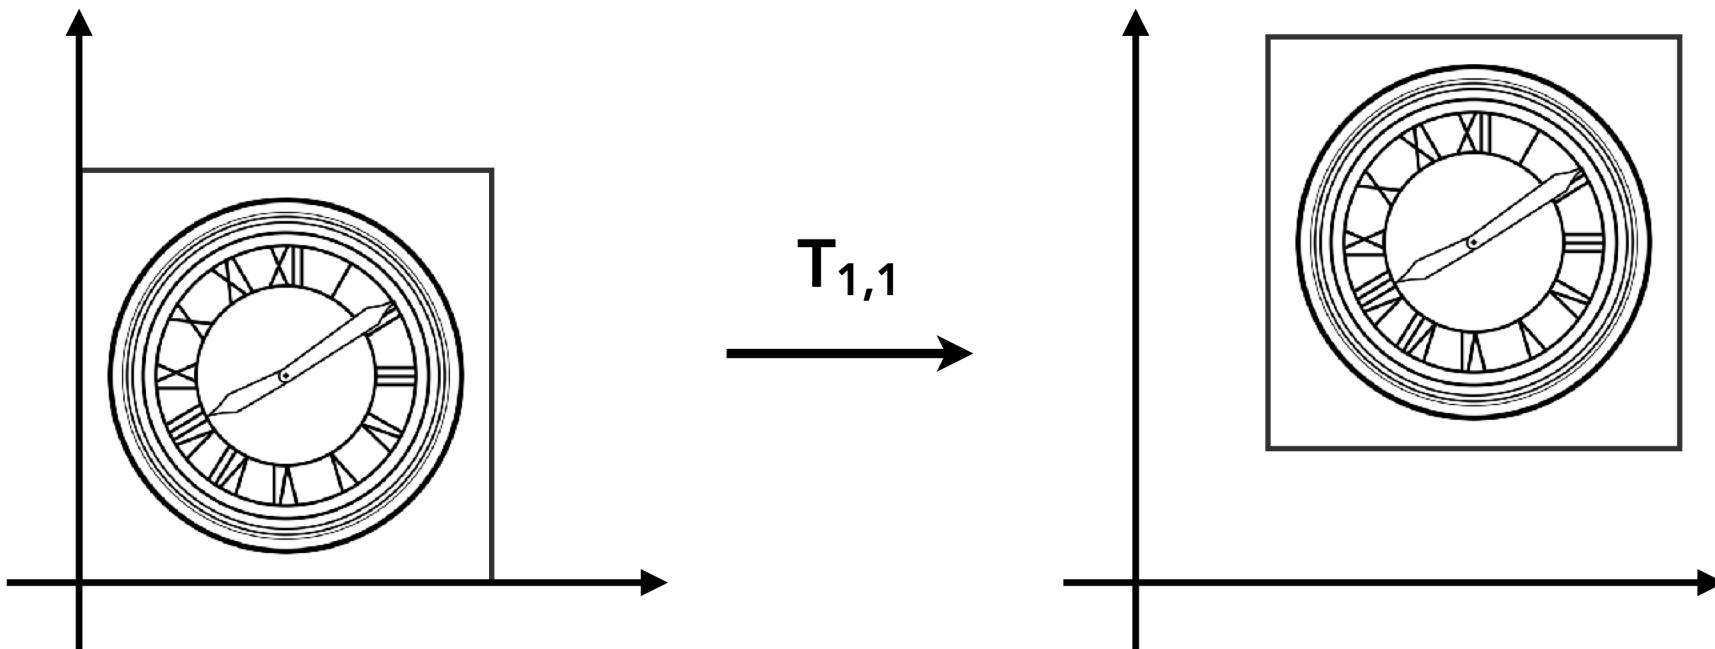
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

$$\mathbf{x}' = \mathbf{M} \mathbf{x}$$

# Translate



# Translation??



$$x' = x + t_x$$

$$y' = y + t_y$$

# Why Homogeneous Coordinates

- Translation cannot be represented in matrix form

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

(So, translation is NOT linear transform)

- But we don't want translation to be a special case
- Is there a unified way to represent all transformations?  
(and what's the cost?)

# Solution: Homogenous Coordinates

Add a third coordinate (**w-coordinate**)

- 2D point =  $(x, y, 1)^T$
- 2D vector =  $(x, y, 0)^T$

Matrix representation of translations

$$\begin{pmatrix} x' \\ y' \\ w' \end{pmatrix} = \begin{pmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = \begin{pmatrix} x + t_x \\ y + t_y \\ 1 \end{pmatrix}$$

# Affine Transformations

**Affine map = linear map + translation**

$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \cdot \begin{pmatrix} x \\ y \end{pmatrix} + \begin{pmatrix} t_x \\ t_y \end{pmatrix}$$

**Using homogenous coordinates:**

$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} a & b & t_x \\ c & d & t_y \\ 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

# 2D Transformations

## Scale

$$\mathbf{S}(s_x, s_y) = \begin{pmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

## Rotation

$$\mathbf{R}(\alpha) = \begin{pmatrix} \cos \alpha & -\sin \alpha & 0 \\ \sin \alpha & \cos \alpha & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

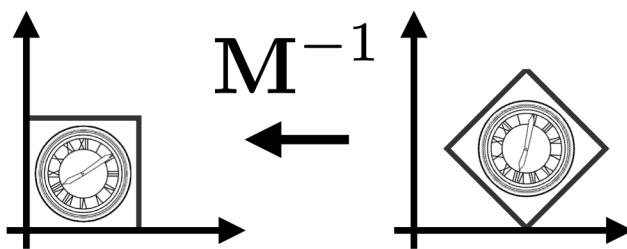
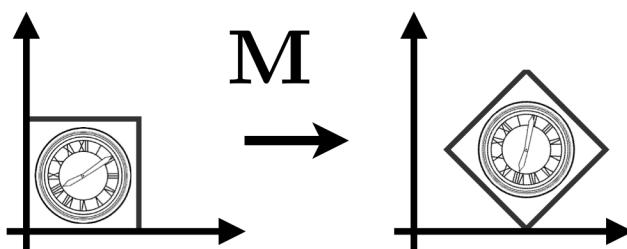
## Translation

$$\mathbf{T}(t_x, t_y) = \begin{pmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{pmatrix}$$

# Inverse Transform

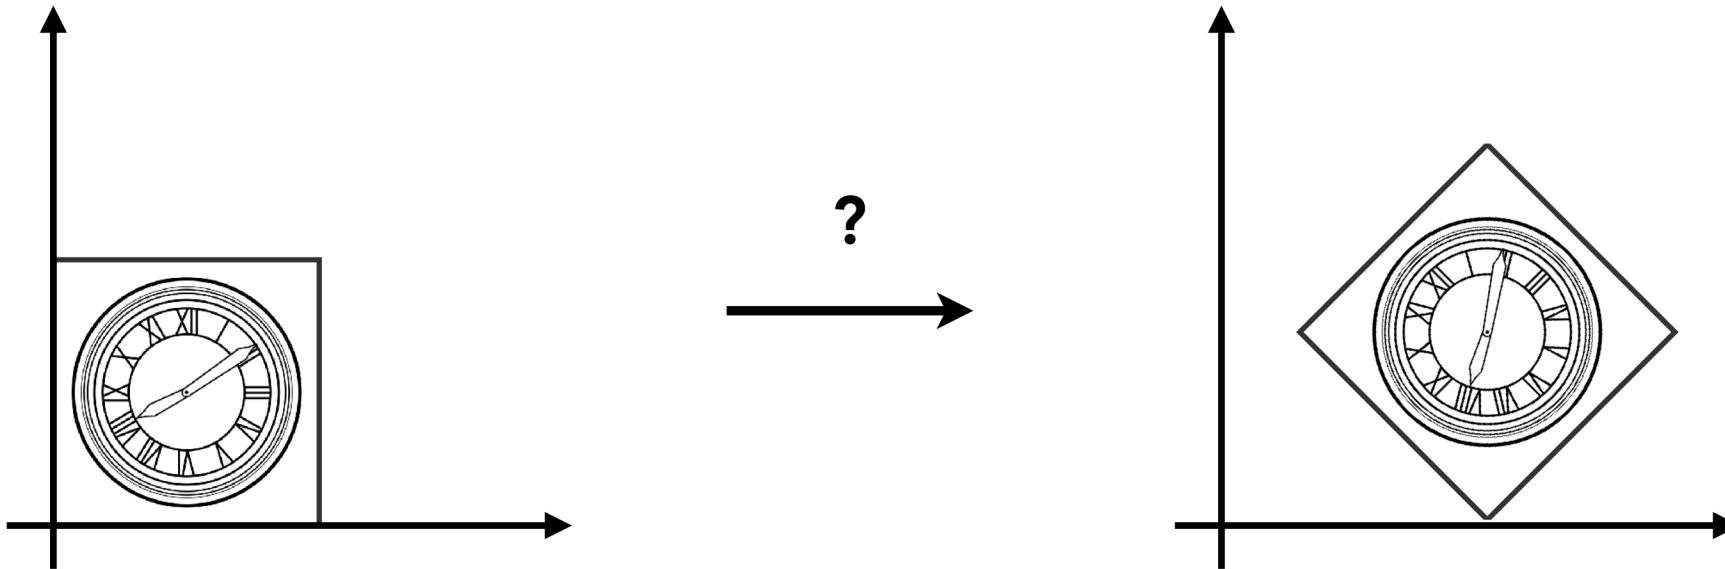
$$M^{-1}$$

$M^{-1}$  is the inverse of transform  $M$  in both a matrix and geometric sense

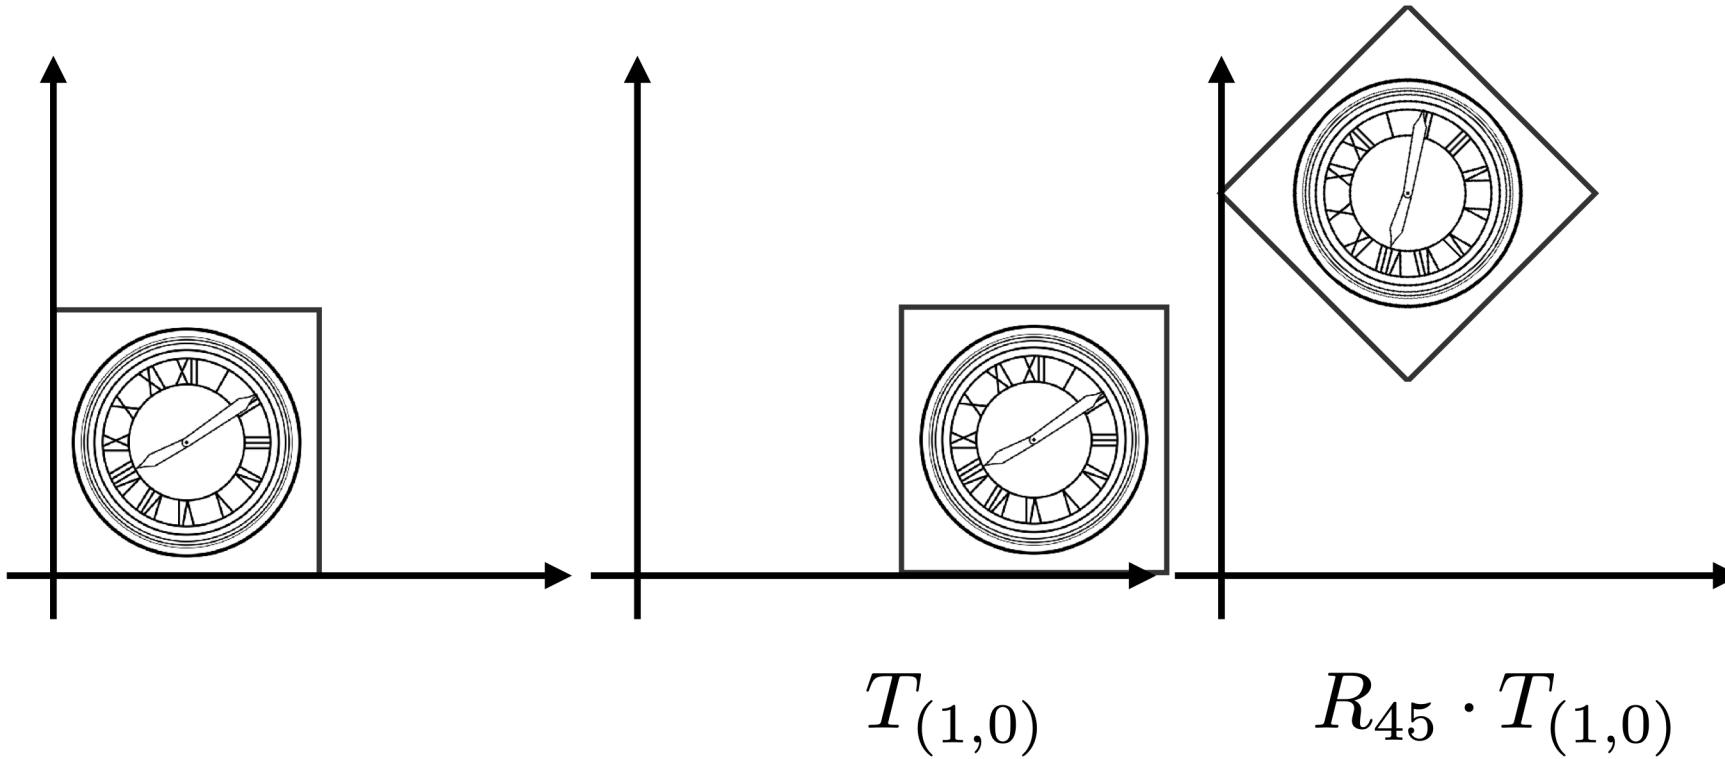


# **Composing Transforms**

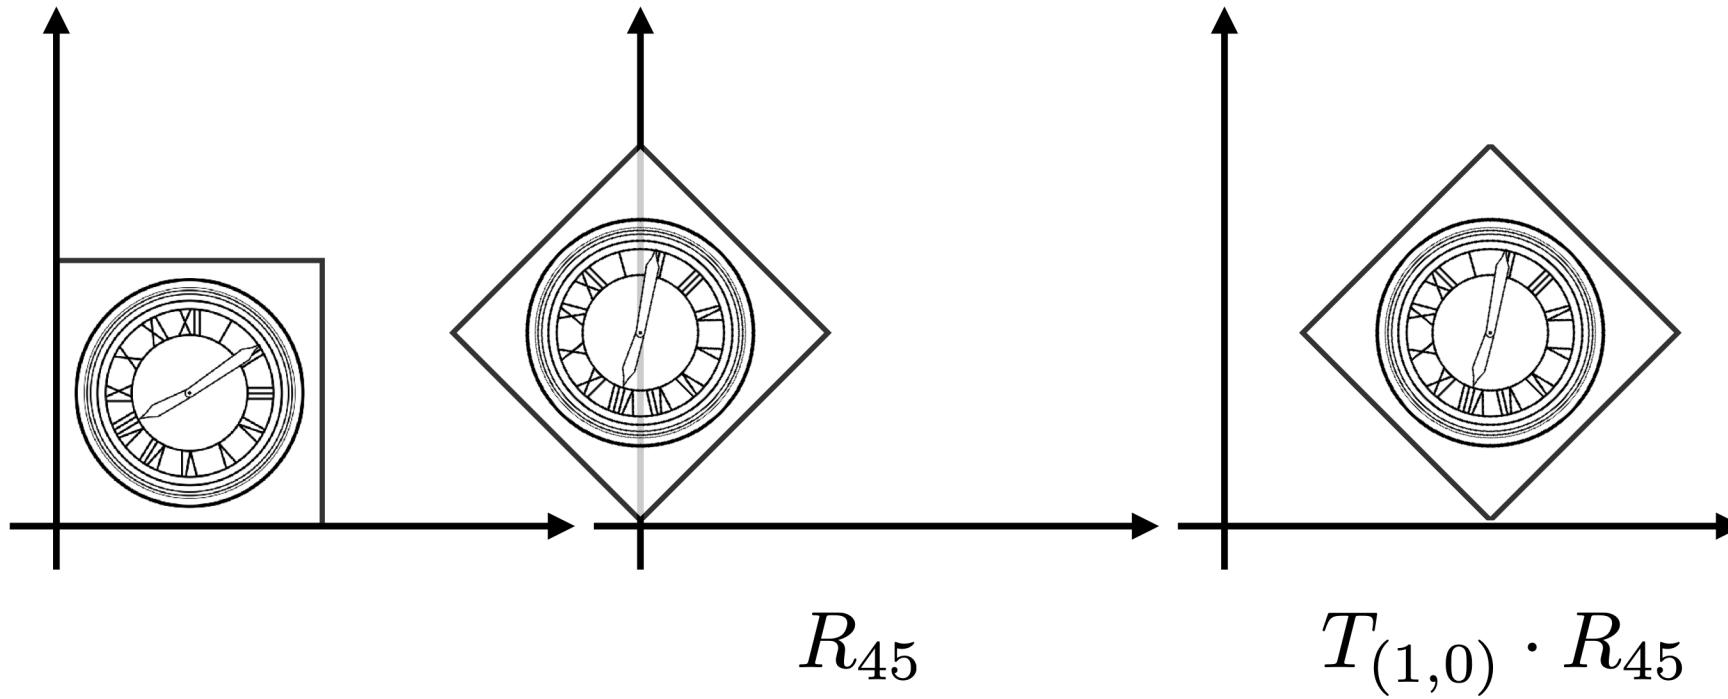
# Composite Transform



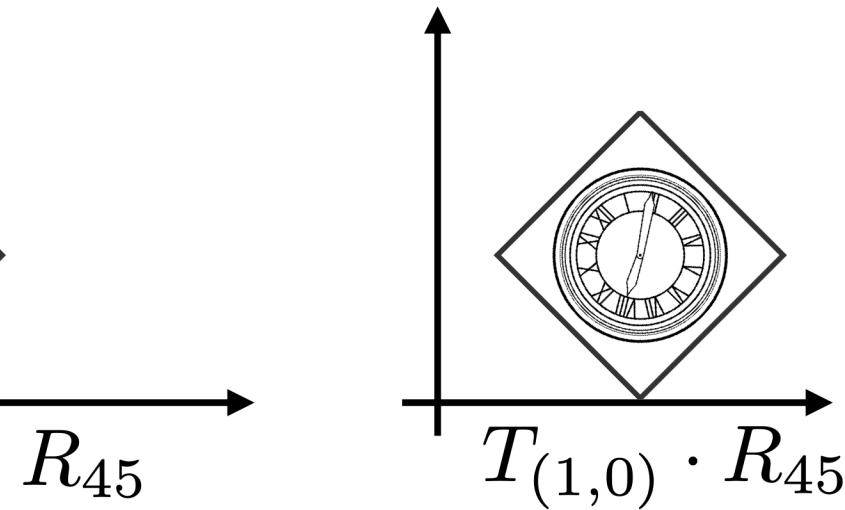
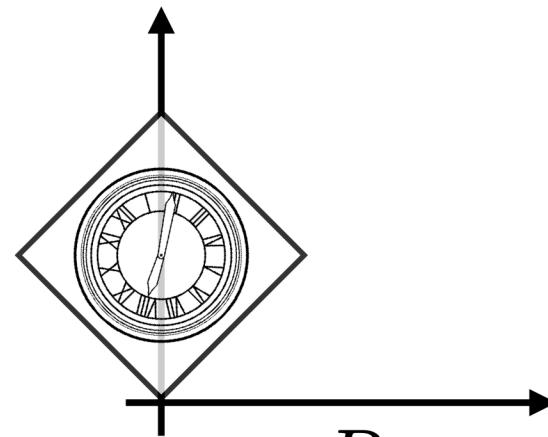
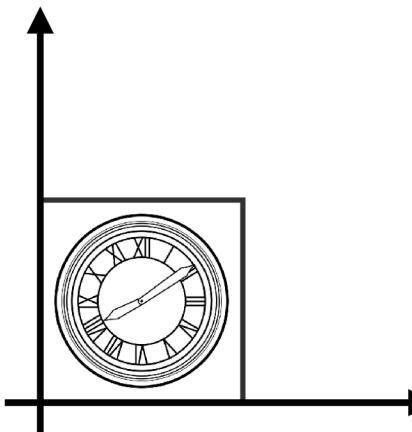
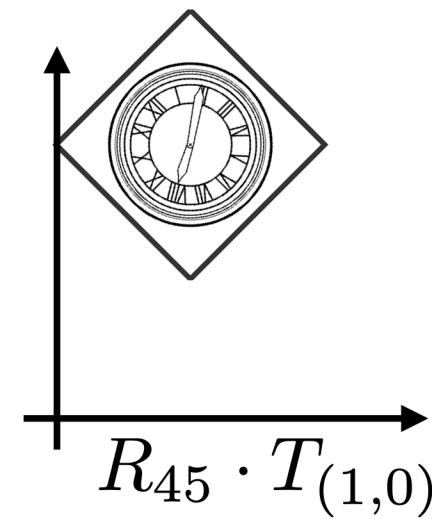
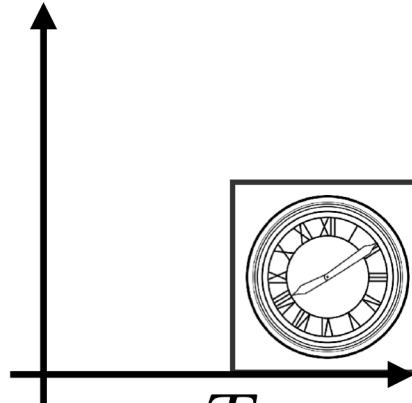
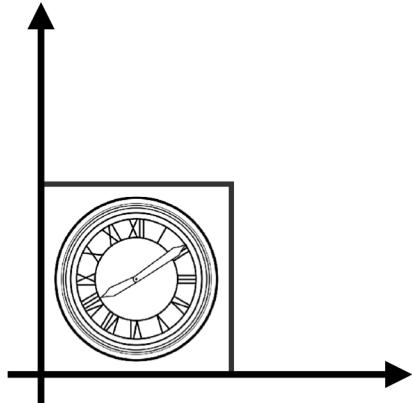
# Translate Then Rotate?



# Rotate Then Translate



# Transform Ordering Matters!



# Transform Ordering Matters!

Matrix multiplication is not commutative

$$R_{45} \cdot T_{(1,0)} \neq T_{(1,0)} \cdot R_{45}$$

Recall the matrix math represented by these symbols:

$$\begin{bmatrix} \cos 45^\circ & -\sin 45^\circ & 0 \\ \sin 45^\circ & \cos 45^\circ & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \neq \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \cos 45^\circ & -\sin 45^\circ & 0 \\ \sin 45^\circ & \cos 45^\circ & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Note that matrices are applied right to left:

$$T_{(1,0)} \cdot R_{45} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \cos 45^\circ & -\sin 45^\circ & 0 \\ \sin 45^\circ & \cos 45^\circ & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Composing Transforms

Sequence of affine transforms  $A_1, A_2, A_3, \dots$

- Compose by matrix multiplication
  - Very important for performance!

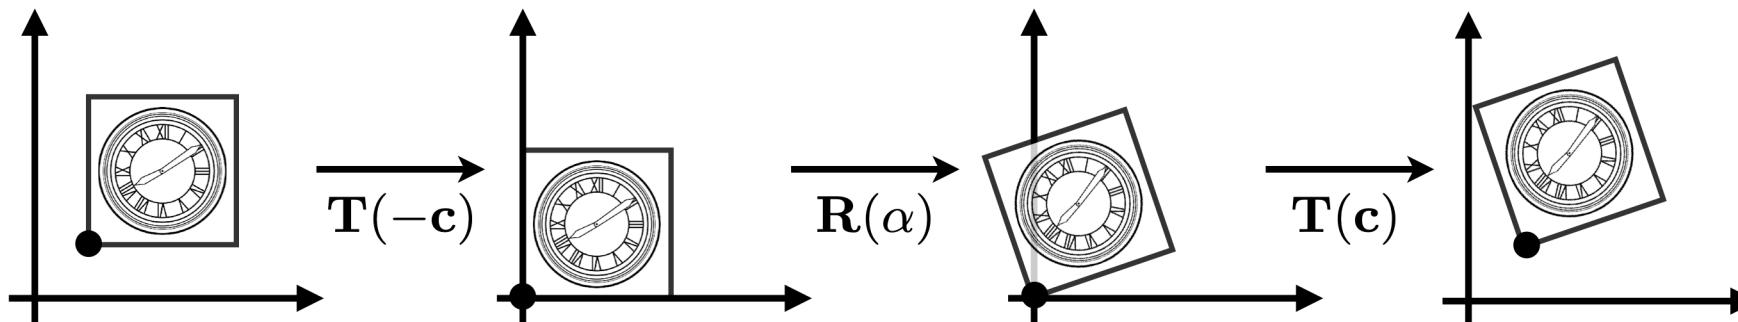
$$A_n(\dots A_2(A_1(\mathbf{x}))) = \mathbf{A}_n \cdots \mathbf{A}_2 \cdot \mathbf{A}_1 \cdot \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$


Pre-multiply  $n$  matrices to obtain a single matrix representing combined transform

# Decomposing Complex Transforms

How to rotate around a given point  $c$ ?

1. Translate center to origin
2. Rotate
3. Translate back



Matrix representation?

$$T(c) \cdot R(\alpha) \cdot T(-c)$$

# Where We Are

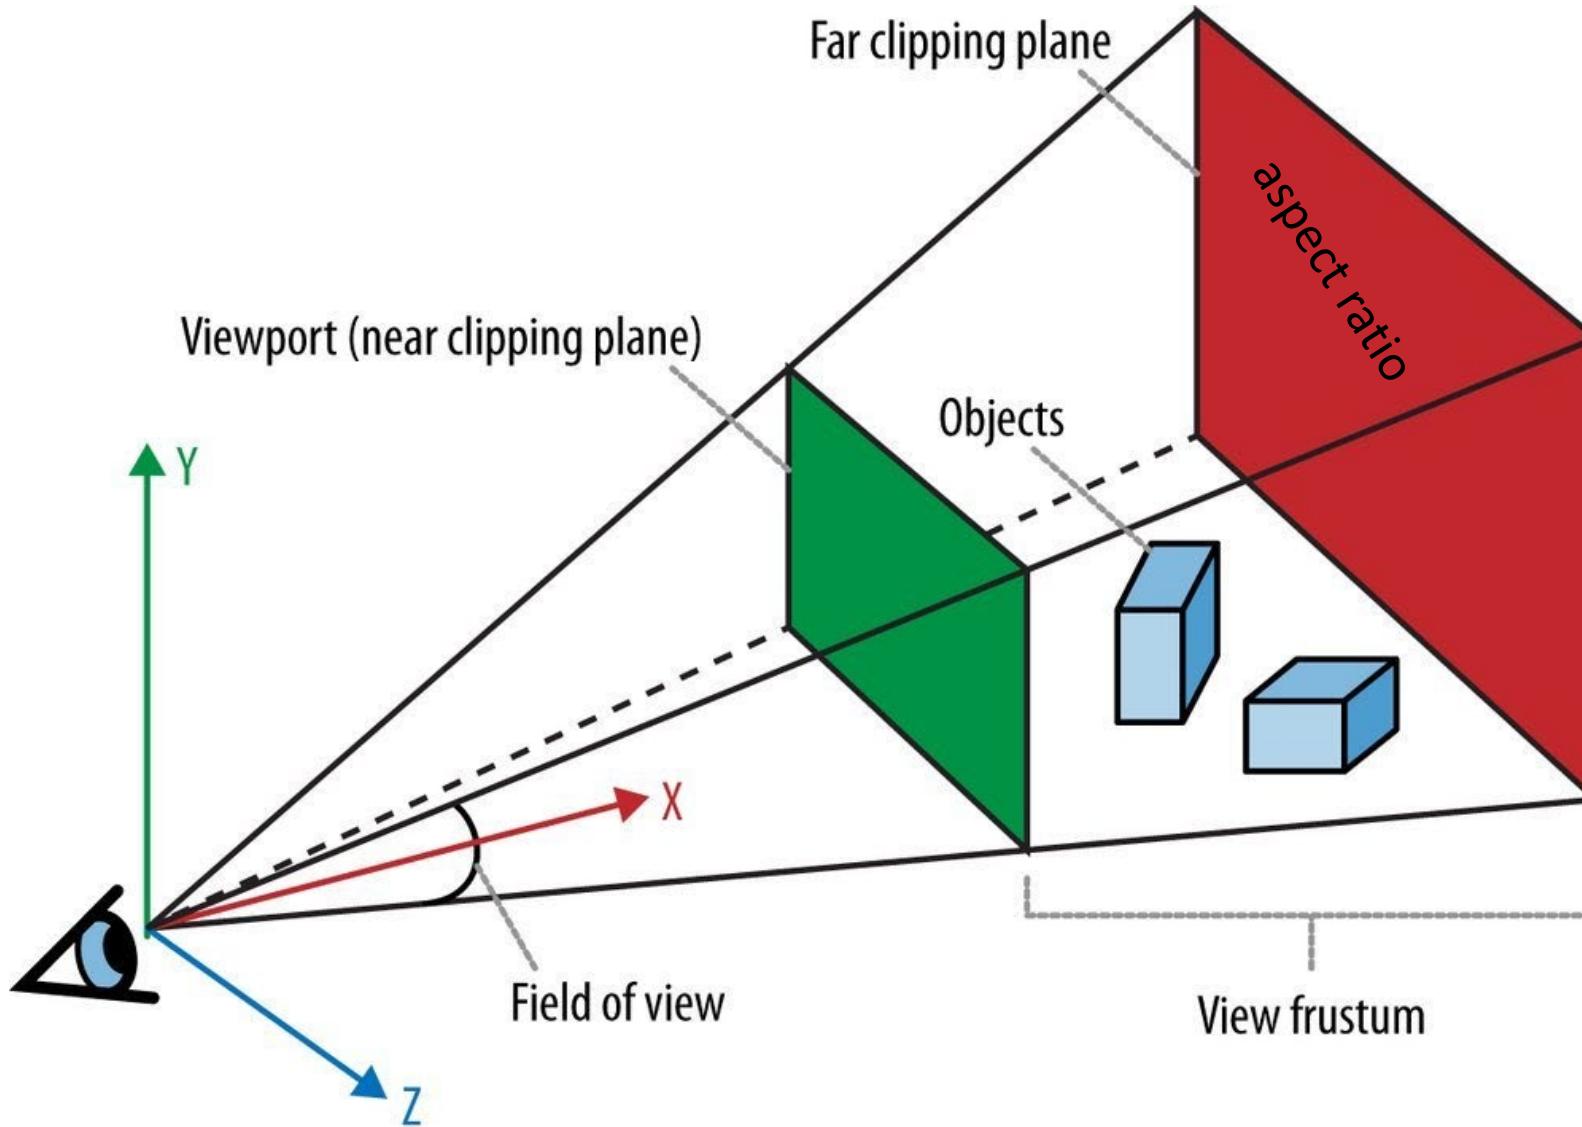
- Basic mathematical tools
- Color: definition, space, map, and design
- 2D visualization: data -> spatial, temporal, network -> interaction in 2D
- Interaction in 3D
- Scientific and 3D visualization
- Application case studies

P A R T   0 2

---

# Interaction: **(3D)** Transformation

# “Camera” in 3D Representation



# Why 3D Transformation?



# 3D Transformations

Use homogeneous coordinates again:

- 3D point =  $(x, y, z, 1)^T$
- 3D vector =  $(x, y, z, 0)^T$

Use  $4 \times 4$  matrices for affine transformations

$$\begin{pmatrix} x' \\ y' \\ z' \\ 1 \end{pmatrix} = \begin{pmatrix} a & b & c & t_x \\ d & e & f & t_y \\ g & h & i & t_z \\ 0 & 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$$

# 3D Transformations

**Scale**

$$\mathbf{S}(s_x, s_y, s_z) = \begin{pmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

**Translation**

$$\mathbf{T}(t_x, t_y, t_z) = \begin{pmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

**Coordinate Change  
(Frame-to-world)**

$$\mathbf{F}(\mathbf{u}, \mathbf{v}, \mathbf{w}, \mathbf{o}) = \begin{bmatrix} \mathbf{u} & \mathbf{v} & \mathbf{w} & \mathbf{o} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

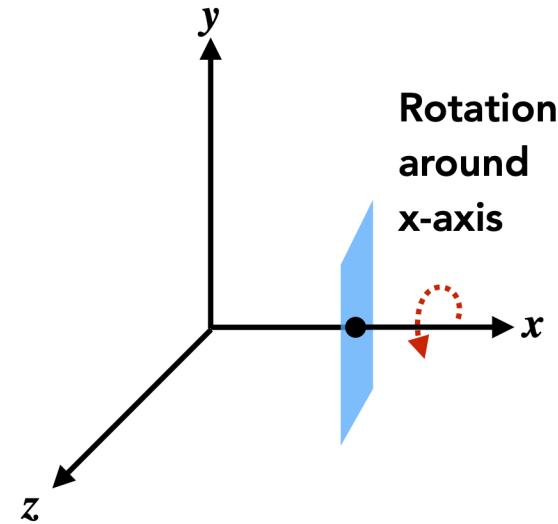
# 3D Transformations

Rotation around x-, y-, or z-axis

$$\mathbf{R}_x(\alpha) = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos \alpha & -\sin \alpha & 0 \\ 0 & \sin \alpha & \cos \alpha & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

$$\mathbf{R}_y(\alpha) = \begin{pmatrix} \cos \alpha & 0 & \sin \alpha & 0 \\ 0 & 1 & 0 & 0 \\ -\sin \alpha & 0 & \cos \alpha & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

$$\mathbf{R}_z(\alpha) = \begin{pmatrix} \cos \alpha & -\sin \alpha & 0 & 0 \\ \sin \alpha & \cos \alpha & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

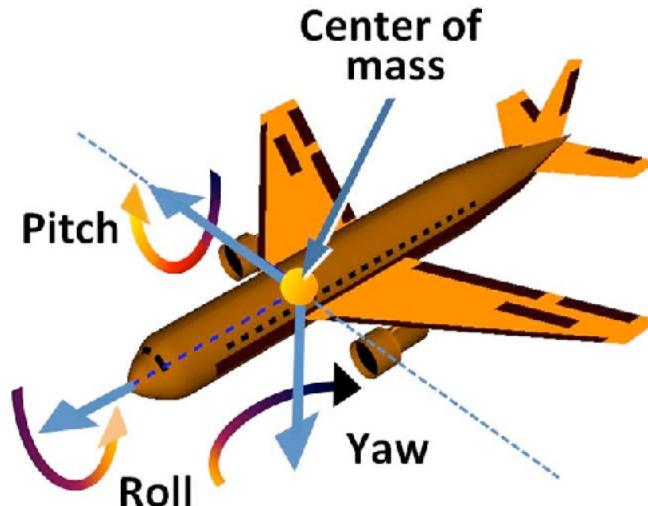


# 3D Rotations

Compose any 3D rotation from  $\mathbf{R}_x, \mathbf{R}_y, \mathbf{R}_z$ ?

$$\mathbf{R}_{xyz}(\alpha, \beta, \gamma) = \mathbf{R}_x(\alpha) \mathbf{R}_y(\beta) \mathbf{R}_z(\gamma)$$

- So-called *Euler angles*
- Often used in flight simulators: roll, pitch, yaw



# 3D Rotations

Compose any 3D rotation from  $\mathbf{R}_x, \mathbf{R}_y, \mathbf{R}_z$ ?

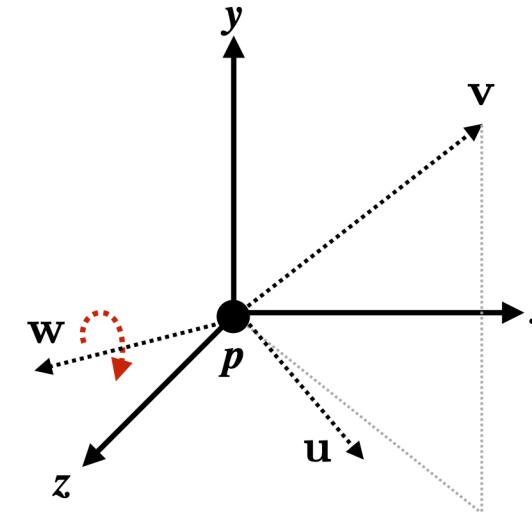
$$\mathbf{R}_{xyz}(\alpha, \beta, \gamma) = \mathbf{R}_x(\alpha) \mathbf{R}_y(\beta) \mathbf{R}_z(\gamma)$$

- So-called *Euler angles*
- Often used in flight simulators: roll, pitch, yaw

# 3D Rotation Around Arbitrary Axis

Construct orthonormal frame transformation  $F$  with  $p, u, v, w$ , where  $p$  and  $w$  match the rotation axis

Apply the transform  $(F R_z(\theta) F^{-1})$



Interpretations (both valid):

- Move to Z axis, rotate, then move back
- Cast w-axis rotation in new coordinate frame

# Rodrigues' Rotation Formula

Rotation by angle  $\alpha$  around axis  $\mathbf{n}$

$$\mathbf{R}(\mathbf{n}, \alpha) = \cos(\alpha) \mathbf{I} + (1 - \cos(\alpha)) \mathbf{n}\mathbf{n}^T + \sin(\alpha) \underbrace{\begin{pmatrix} 0 & -n_z & n_y \\ n_z & 0 & -n_x \\ -n_y & n_x & 0 \end{pmatrix}}_{\mathbf{N}}$$

How to prove this magic formula?

- Matrix  $\mathbf{N}$  computes a cross-product:  $\mathbf{N} \mathbf{x} = \mathbf{n} \times \mathbf{x}$
- Assume orthonormal system  $\mathbf{e}_1, \mathbf{e}_2, \mathbf{n}$

$$\mathbf{R}\mathbf{n} = \mathbf{n}$$

$$\mathbf{R}\mathbf{e}_1 = \cos \alpha \mathbf{e}_1 + \sin \alpha \mathbf{e}_2$$

$$\mathbf{R}\mathbf{e}_2 = -\sin \alpha \mathbf{e}_1 + \cos \alpha \mathbf{e}_2$$

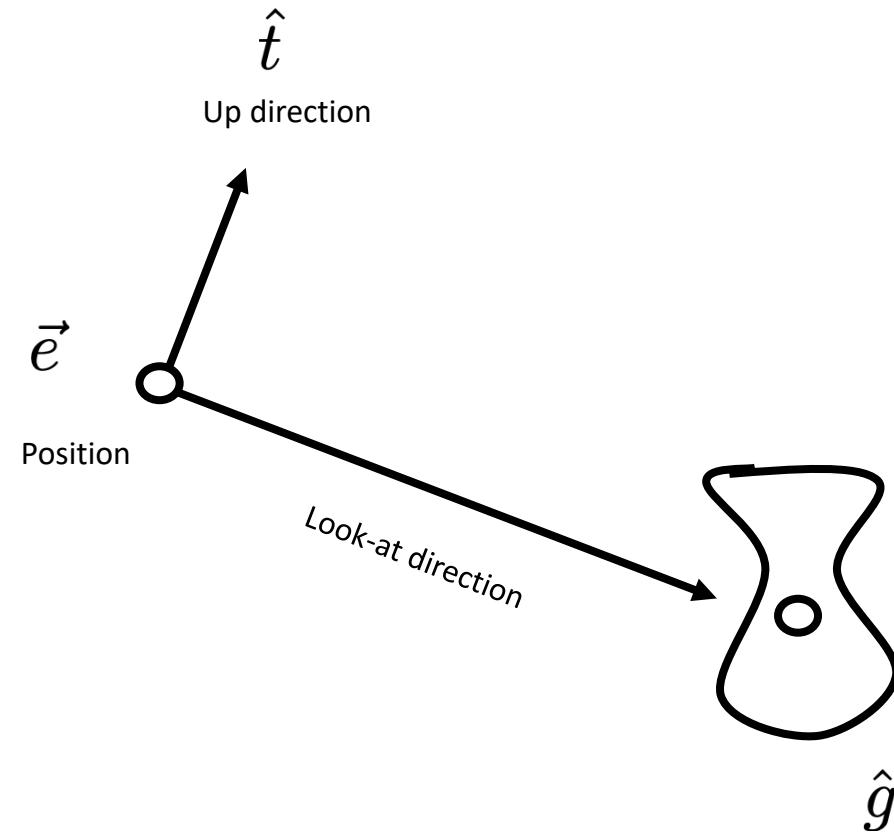
# 3D Transforms

Model-View-Projection

- While taking a selfie
  - Find a good place to stand (**model** transformation)
  - Find a good “angle” to place the camera (**view** transformation)
  - Cheeeese! (**projection** transformation)

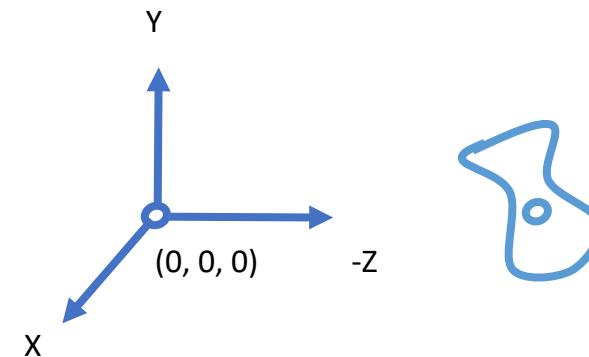
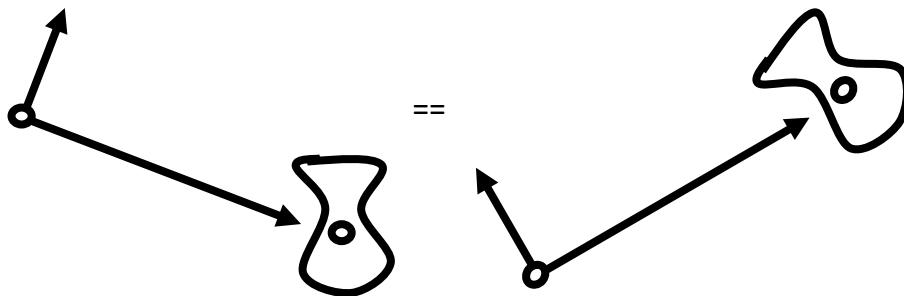
# View / Camera Transformation

- How to perform view transformation?
- Define the camera first
  - Position
  - Look-at / gaze direction
  - Up direction (assuming perp. to look-at)



# View / Camera Transformation

- Key observation
  - If the camera and objects move together, the “photo” will be the same



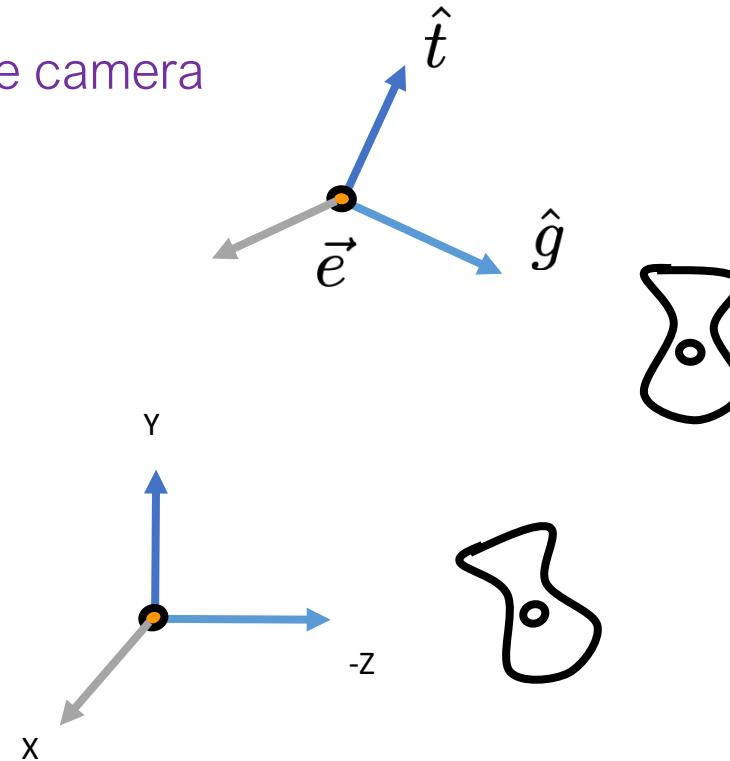
- How about that we move the camera to
  - The origin, up at Y, look at -z
  - Move the objects along with the camera

# View / Camera Transformation

- Transform the camera by  $M_{view}$ 
  - So that it's located at the origin, up at Y, look at -z
  - Transform the objects along with the camera

- $M_{view}$  in math?

- Translates e to origin
- Rotates g to -z
- Rotates t to Y
- Rotates  $(g \times t)$  To X



# View / Camera Transformation

- $M_{view}$  in math?

- Let's write  $M_{view} = R_{view}T_{view}$
- Translate e to origin

$$T_{view} = \begin{bmatrix} 1 & 0 & 0 & -x_e \\ 0 & 1 & 0 & -y_e \\ 0 & 0 & 1 & -z_e \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- Rotate g to -Z, t to Y,  $(g \times t)$  To X
- Consider its **inverse** rotation: X to  $(g \times t)$ , Y to t, Z to -g

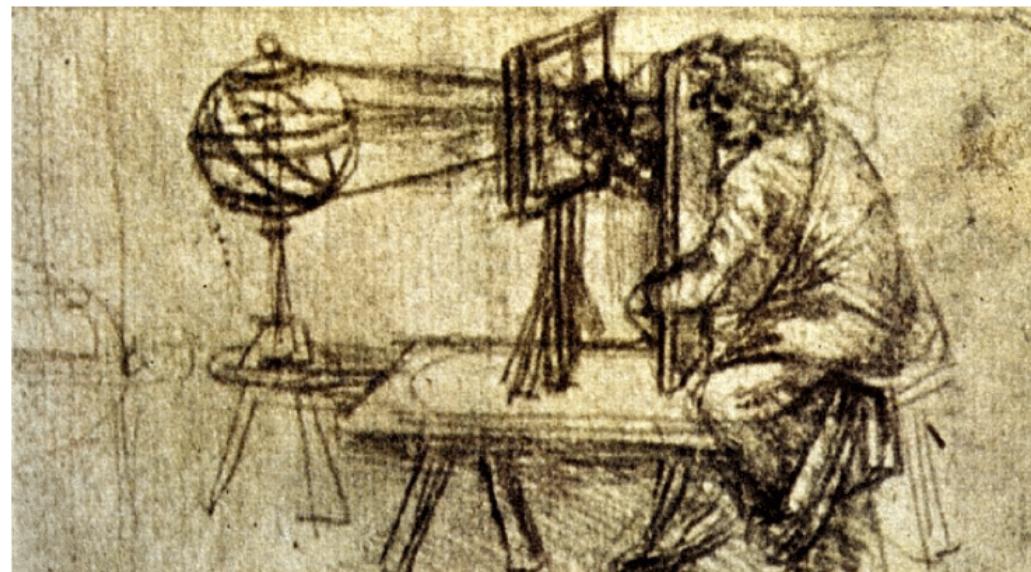
$$R_{view}^{-1} = \begin{bmatrix} x_{\hat{g} \times \hat{t}} & x_t & x_{-g} & 0 \\ y_{\hat{g} \times \hat{t}} & y_t & y_{-g} & 0 \\ z_{\hat{g} \times \hat{t}} & z_t & z_{-g} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \xrightarrow{\text{WHY?}} R_{view} = \begin{bmatrix} x_{\hat{g} \times \hat{t}} & y_{\hat{g} \times \hat{t}} & z_{\hat{g} \times \hat{t}} & 0 \\ x_t & y_t & y_t & 0 \\ x_{-g} & y_{-g} & z_{-g} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

# View / Camera Transformation

- Summary
  - Transform objects together with the camera
  - Until camera's at the origin, up at Y, look at -Z
- But why?
  - For projection transformation!

# 3D Transforms

Model-View-Projection



Leonardo da Vinci, 1515

# Projection Transformation

- Projection in 3D Visualization
  - 3D to 2D
  - Orthographic projection
  - Perspective projection

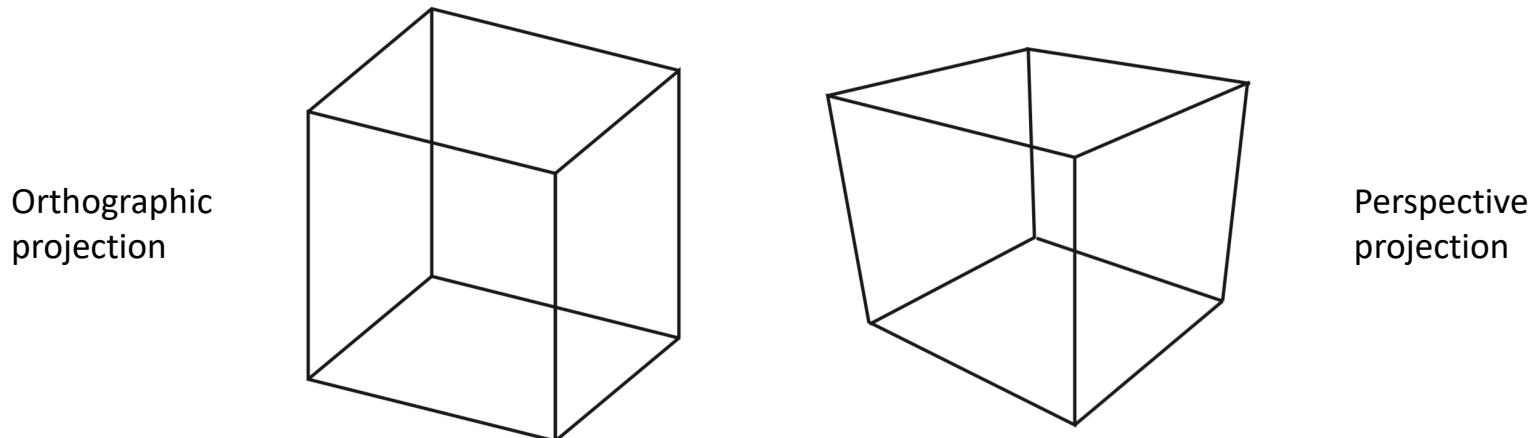
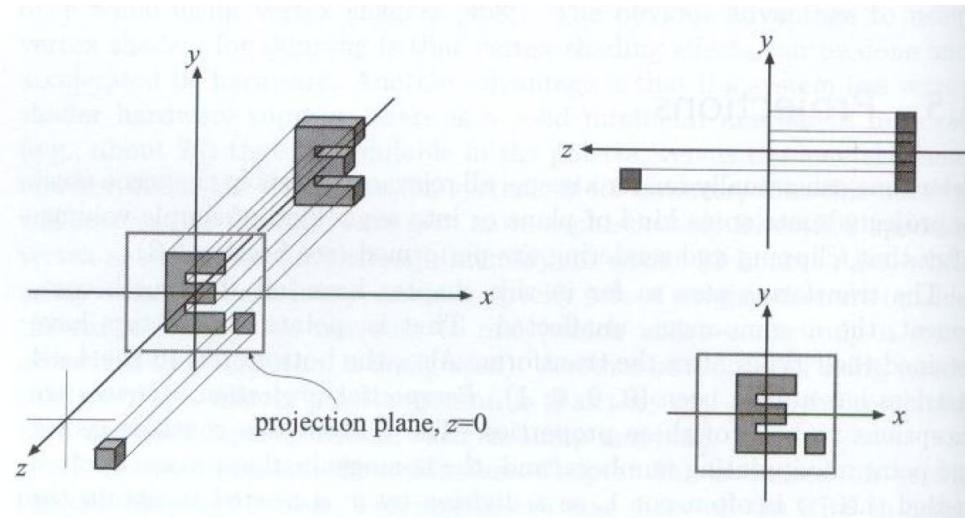


Fig. 7.1 from *Fundamentals of Computer Graphics, 4th Edition*

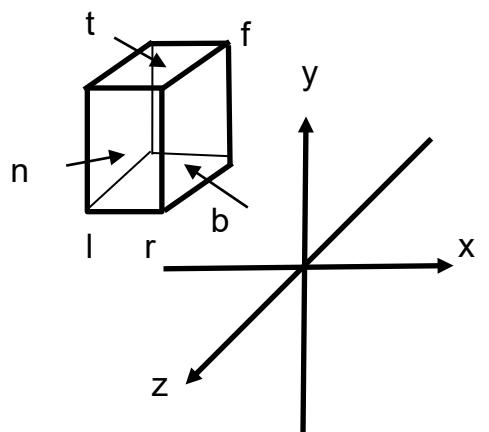
# Orthographic Projection

- A simple way of understanding
  - Camera located at origin, looking at -Z, up at Y (looks familiar?)
  - Drop Z coordinate
  - Translate and scale the resulting rectangle to “canonical” view within  $[-1, 1]^2$

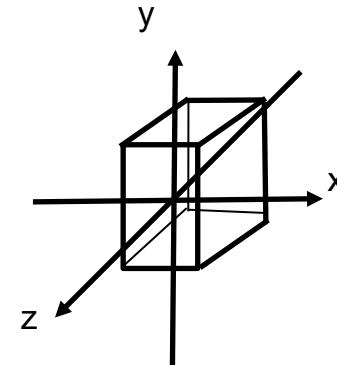


# Orthographic Projection

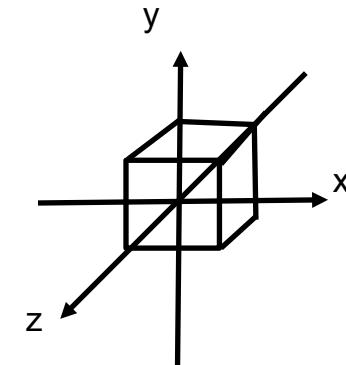
- In general
  - We want to map a cuboid  $[l, r] \times [b, t] \times [f, n]$  to the “canonical” cube  $[-1, 1]^3$



Translate

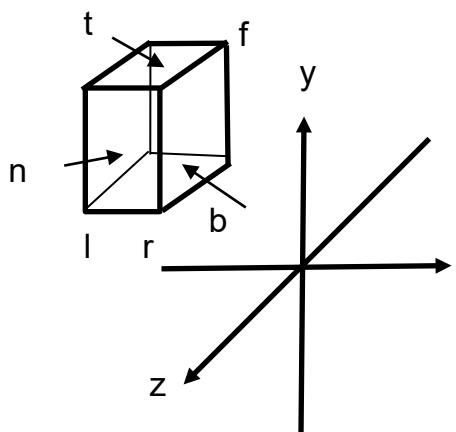


Scale

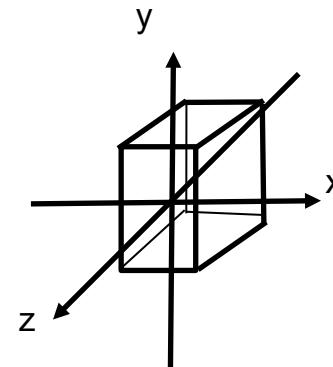


# Orthographic Projection

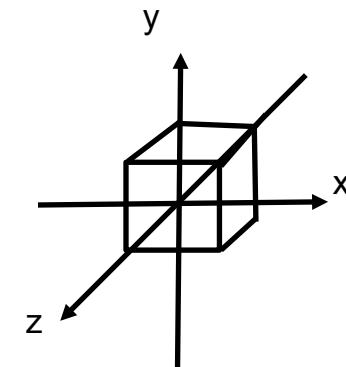
- First center cuboid by translating
- Then scale into “canonical” cube



Translate



Scale

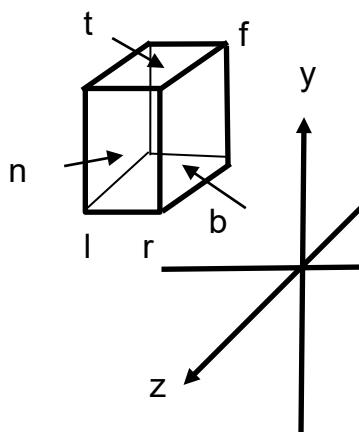


# Orthographic Projection

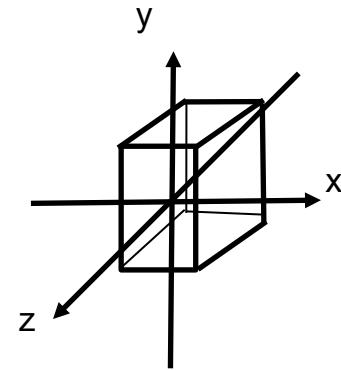
- Transformation matrix?

- Translate (**center** to origin) first, then scale (length/width/height to **2**)

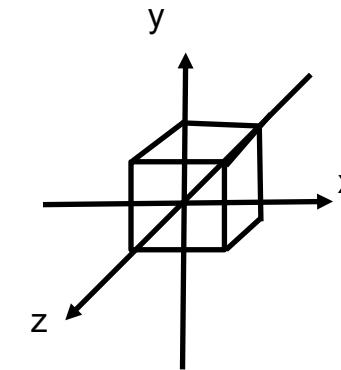
$$M_{ortho} = \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & 0 \\ 0 & \frac{2}{t-b} & 0 & 0 \\ 0 & 0 & \frac{2}{n-f} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & -\frac{r+l}{2} \\ 0 & 1 & 0 & -\frac{t+b}{2} \\ 0 & 0 & 1 & -\frac{n+f}{2} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



Translate



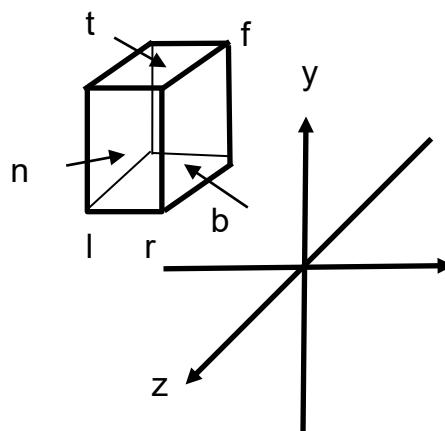
Scale



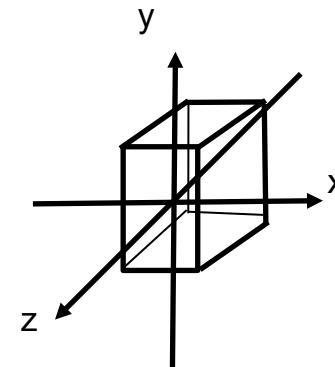
# Orthographic Projection

- Caveat

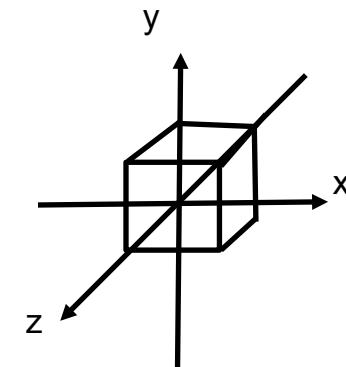
- Looking at / along -Z is making near and far not intuitive ( $n > f$ )
- FYI: that's why OpenGL (a Graphics API) uses left hand coords.



Translate

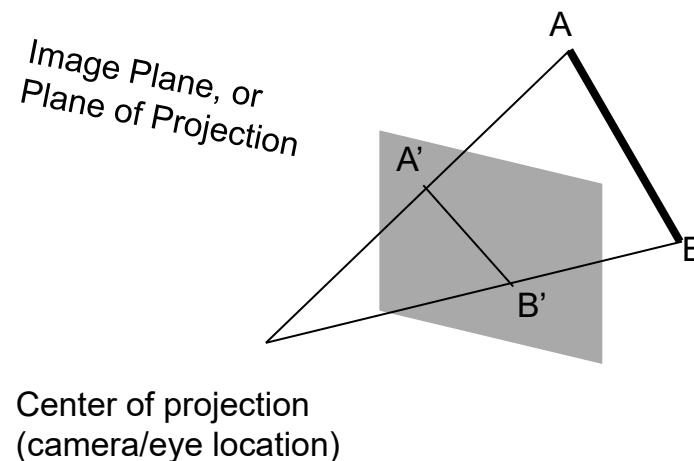


Scale



# Perspective Projection

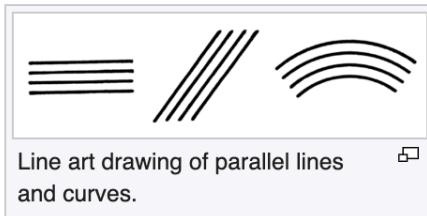
- Most common computer graphics, art, visual system
- Further objects are smaller
- Parallel lines not parallel; converge to single point



# Perspective Projection

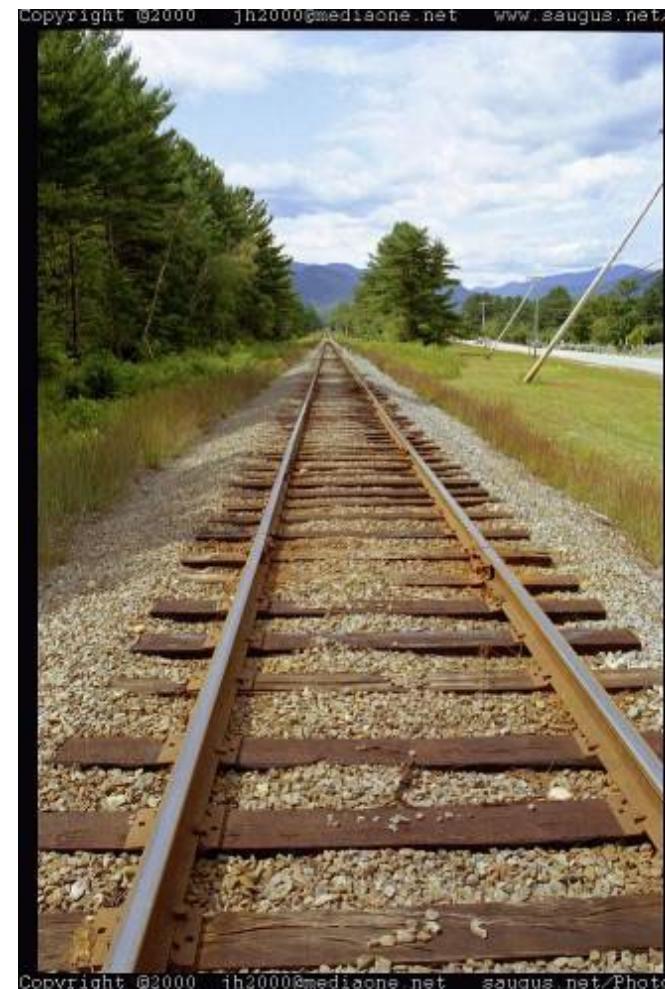
- Euclid was wrong??!!

In [geometry](#), **parallel** lines are [lines in a plane](#) which do not meet; that is, two lines in a plane that do not [intersect](#) or [touch](#) each other at any point are said to be parallel. By extension, a line and a plane, or two planes, in [three-dimensional Euclidean space](#) that do not share a point are said to be parallel. However, two lines in three-dimensional space which do not meet must be in a common plane to be considered parallel; otherwise they are called [skew lines](#). Parallel planes are planes in the same three-dimensional space that never meet.



Parallel lines are the subject of [Euclid's parallel postulate](#).<sup>[1]</sup> Parallelism is primarily a property of [affine geometries](#) and [Euclidean geometry](#) is a special instance of this type of geometry. In some other geometries, such as [hyperbolic geometry](#), lines can have analogous properties that are referred to as parallelism.

[https://en.wikipedia.org/wiki/Parallel\\_\(geometry\)](https://en.wikipedia.org/wiki/Parallel_(geometry))



# Final Project Q&A

- Clarification/concerns on the proposal: we will reach out to individuals if any. Otherwise, go ahead with your plan
- “How many figures/buttons do I need?”
  - Algorithm Projects (i.e., reproducing research articles): whatever the paper does to solve a problem
  - Applications Projects: as an “application”/system, they are **commonly** interactive. Also think about using all the techniques we covered. Consider the interaction to bridge spatial/temporal/network data, etc.
  - Capstone/Research projects: dependss

# 3D Visualization

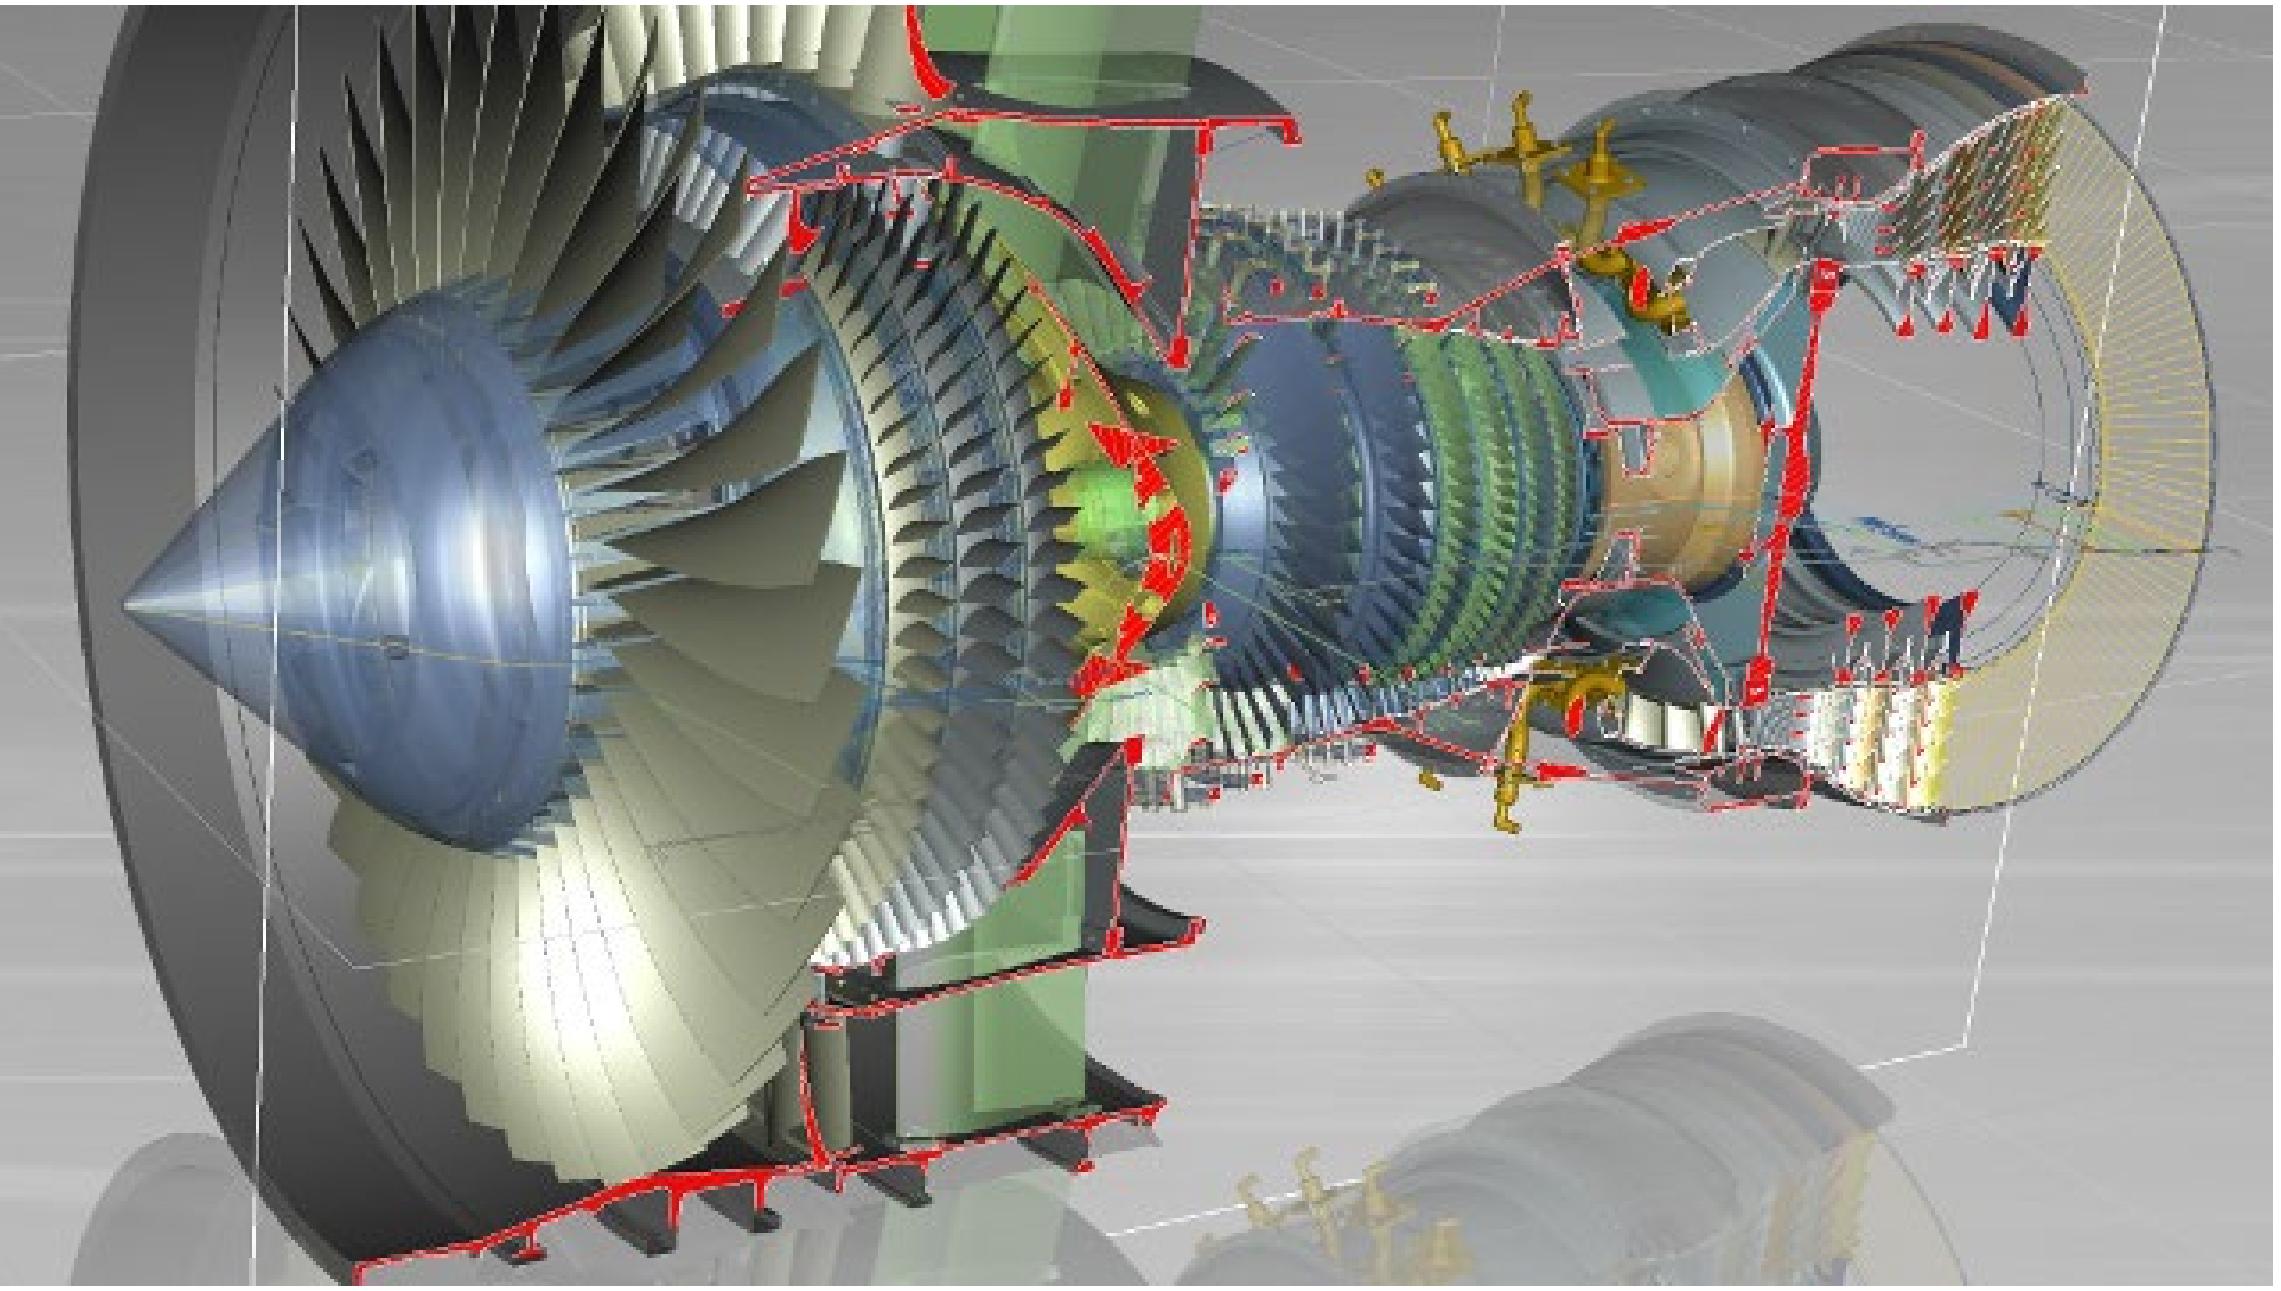
- 3D Data (Scientific Visualization)
- Non-3D Data

---

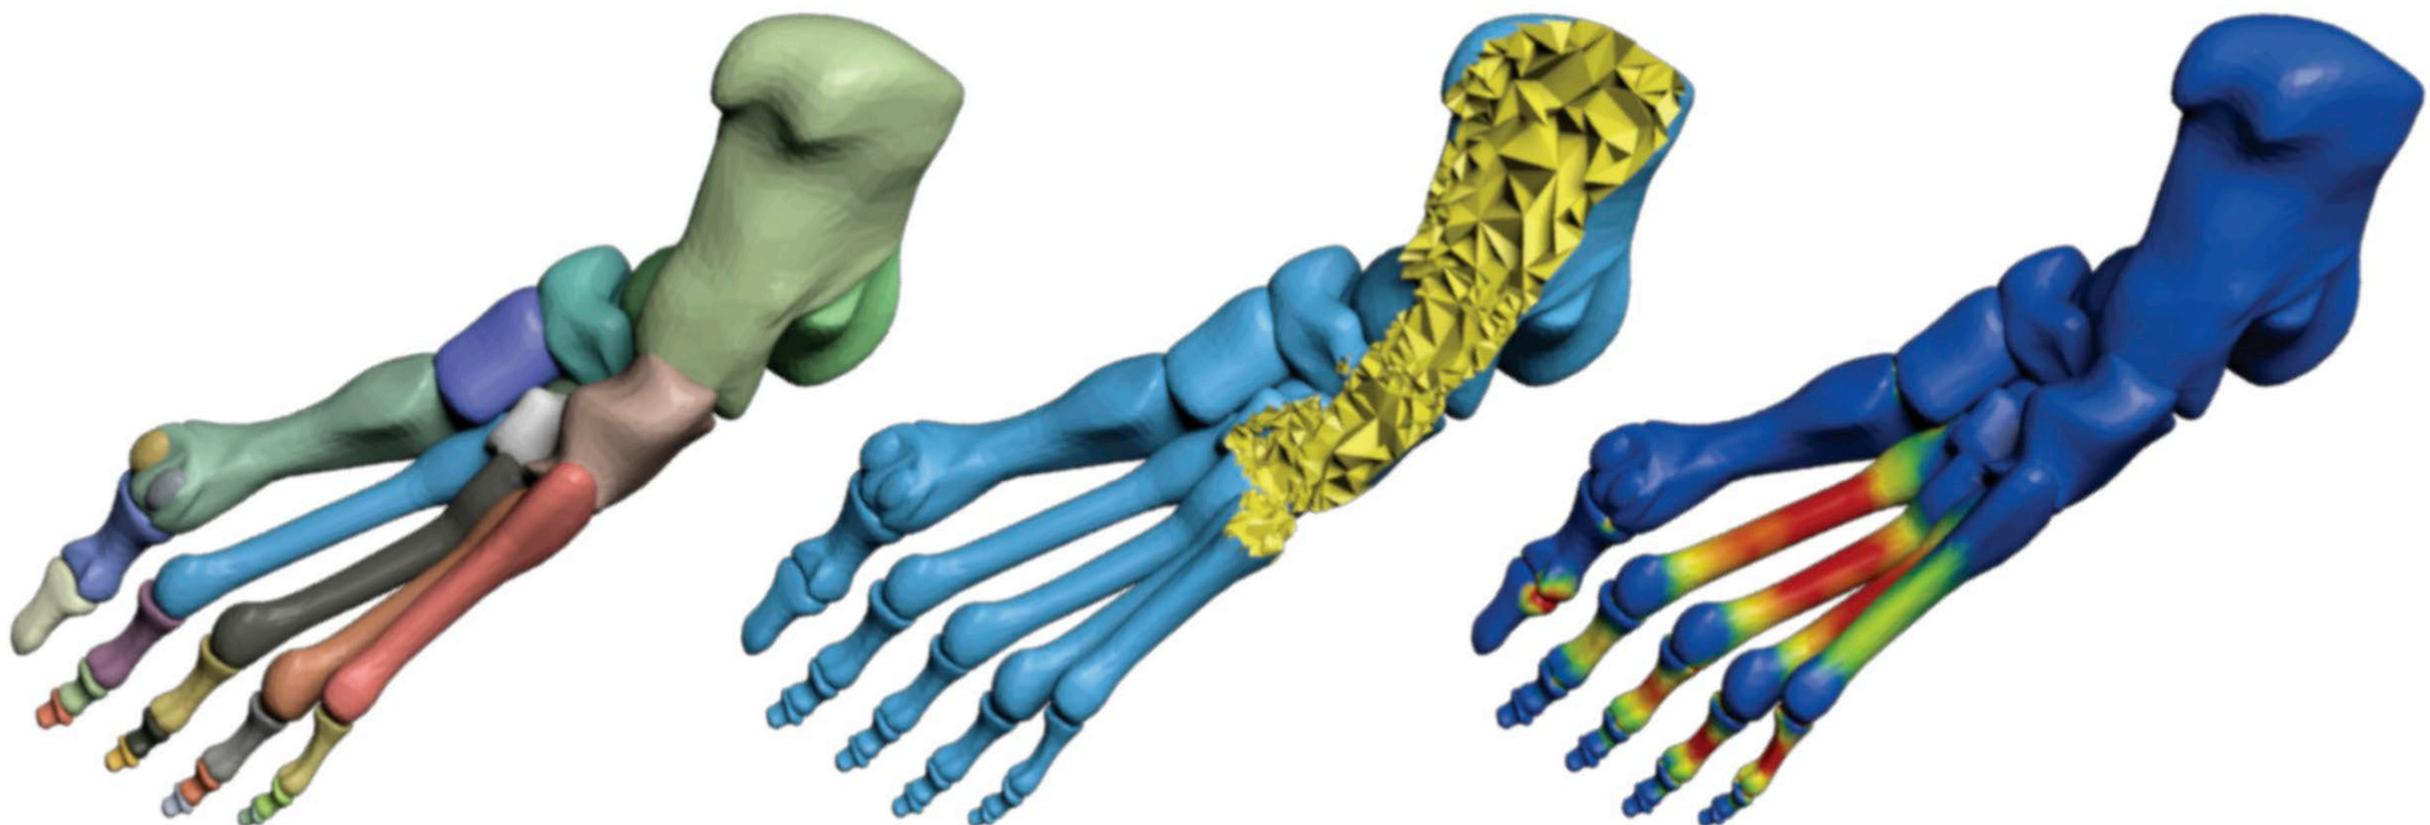
# 3D Visualization for 3D Data

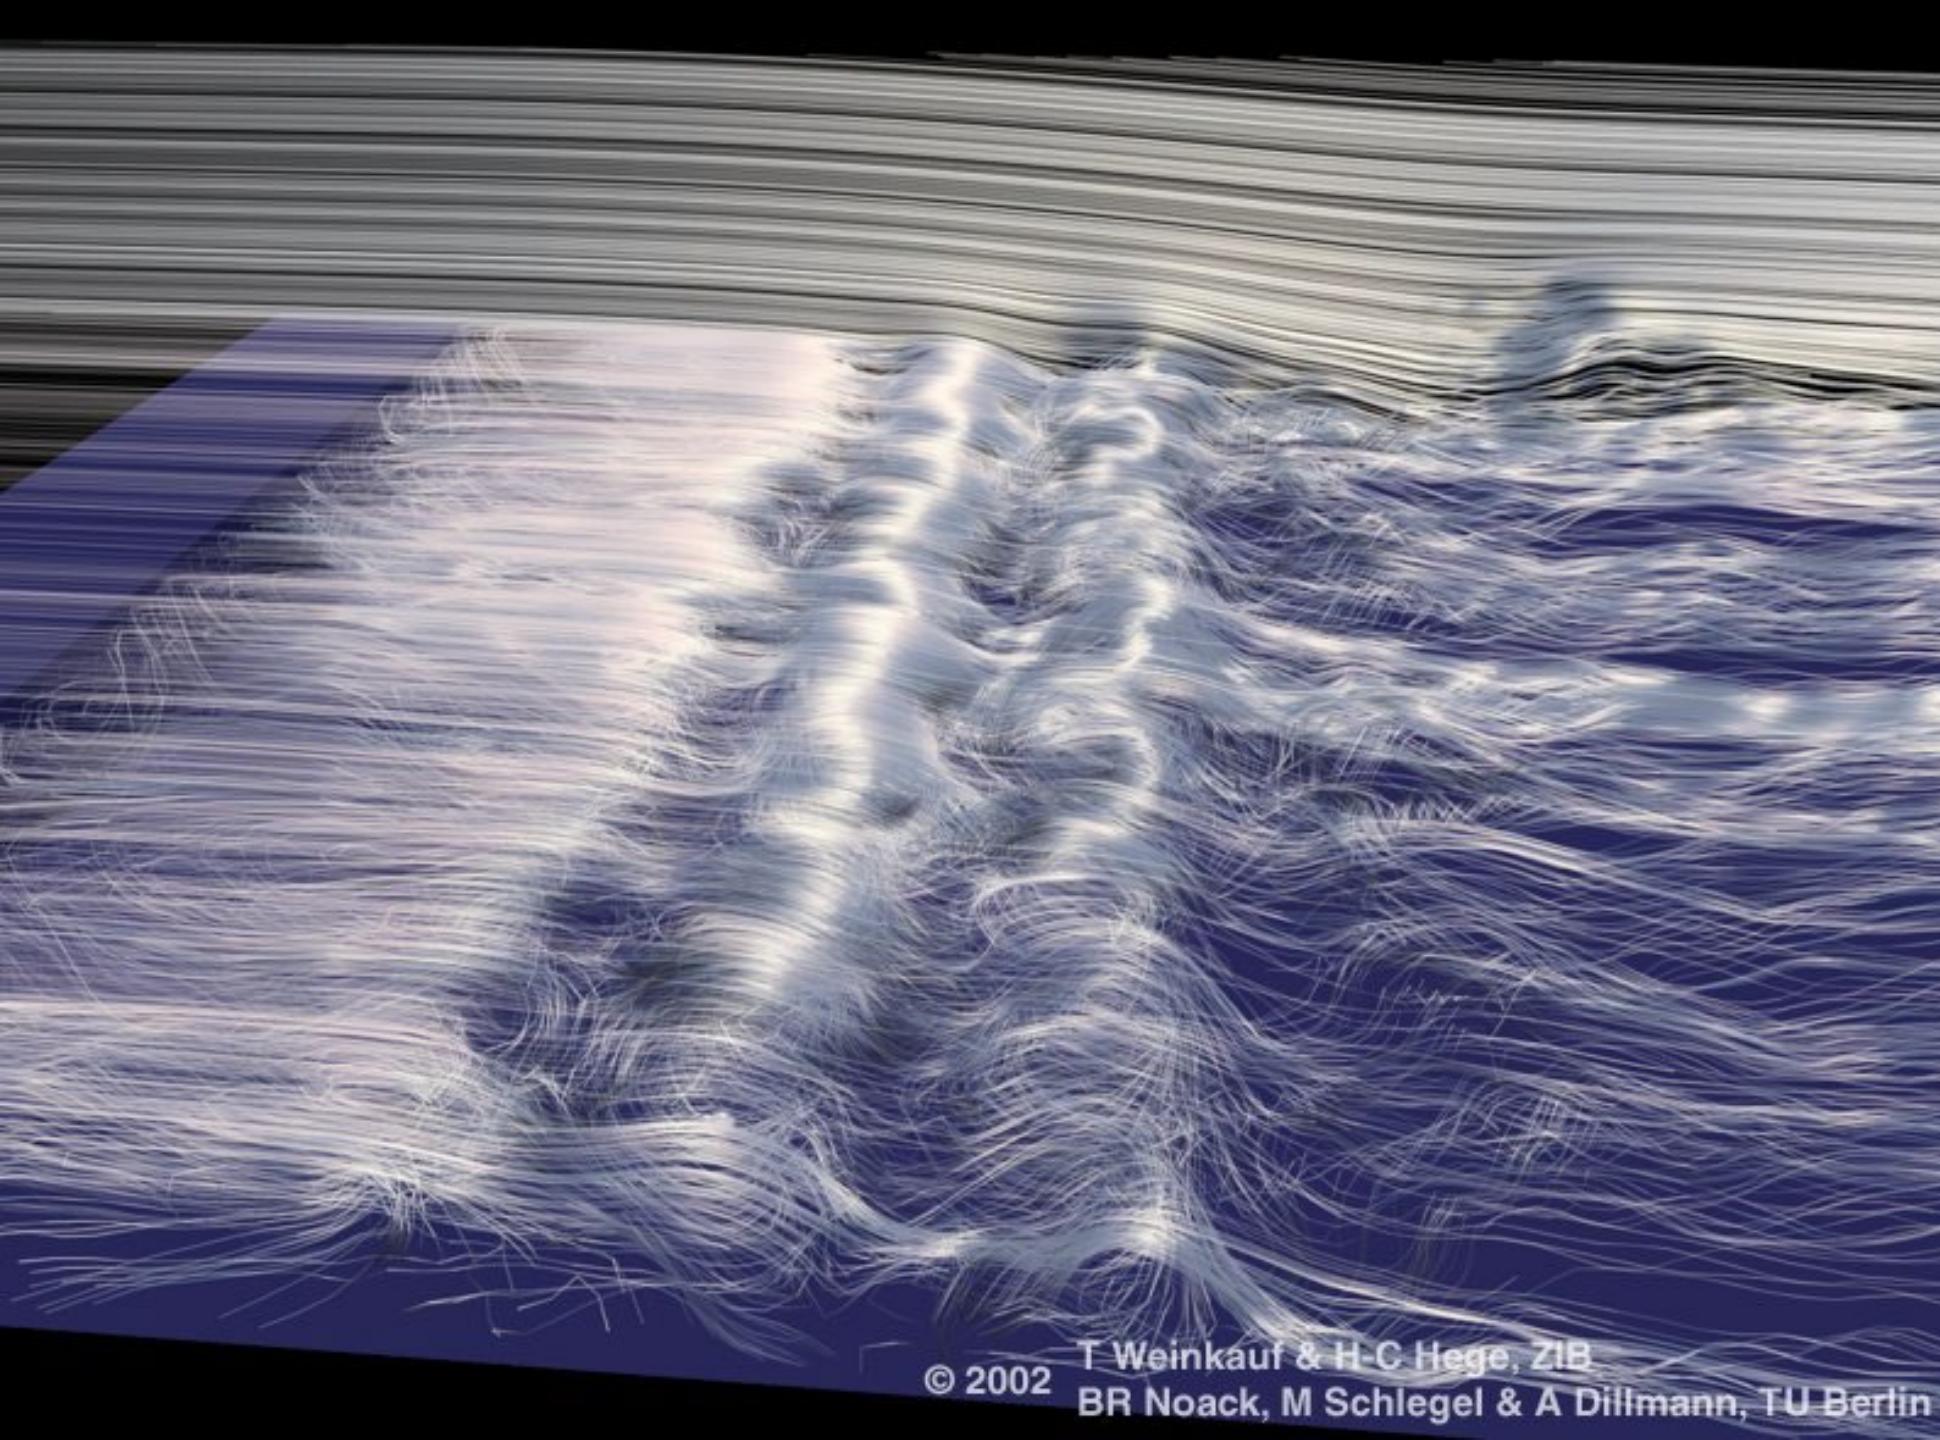
# 3D Data

- Data that are intrinsically 3D-related.
  - Every observation is associated with position (x, y, z).
- Two sub-categories:
  - 3D geometry is the data.
  - Data is associated with 3D geometry.



Source: <https://www.plm.automation.siemens.com/global/en/products/collaboration/digital-mockup.html>

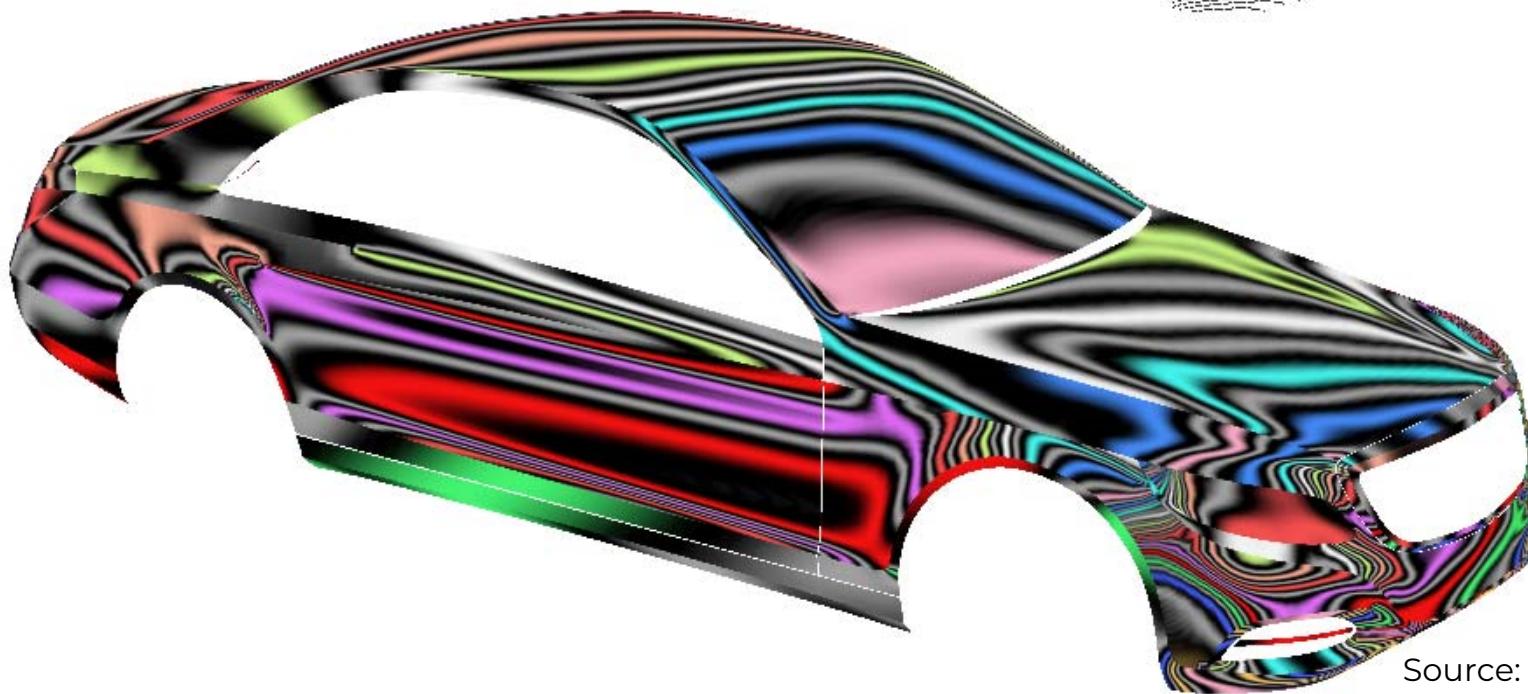
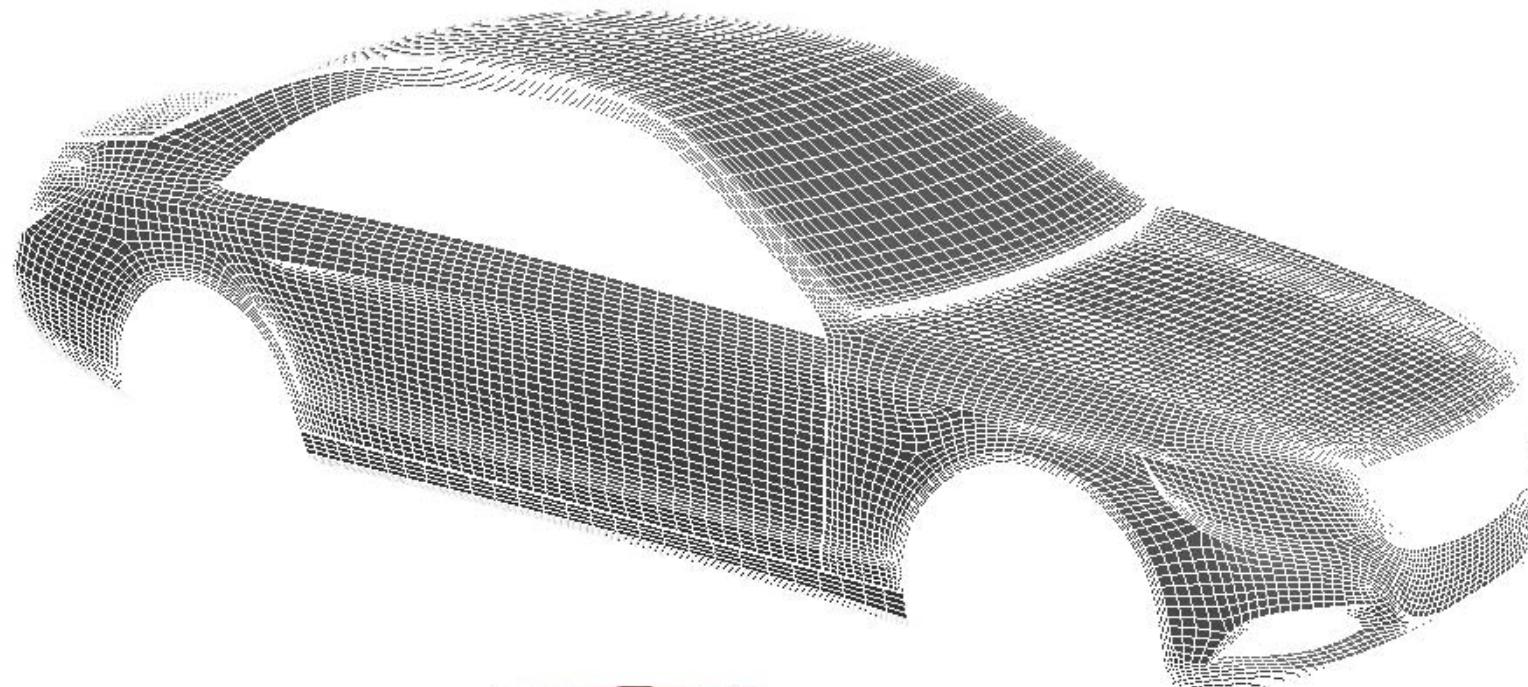




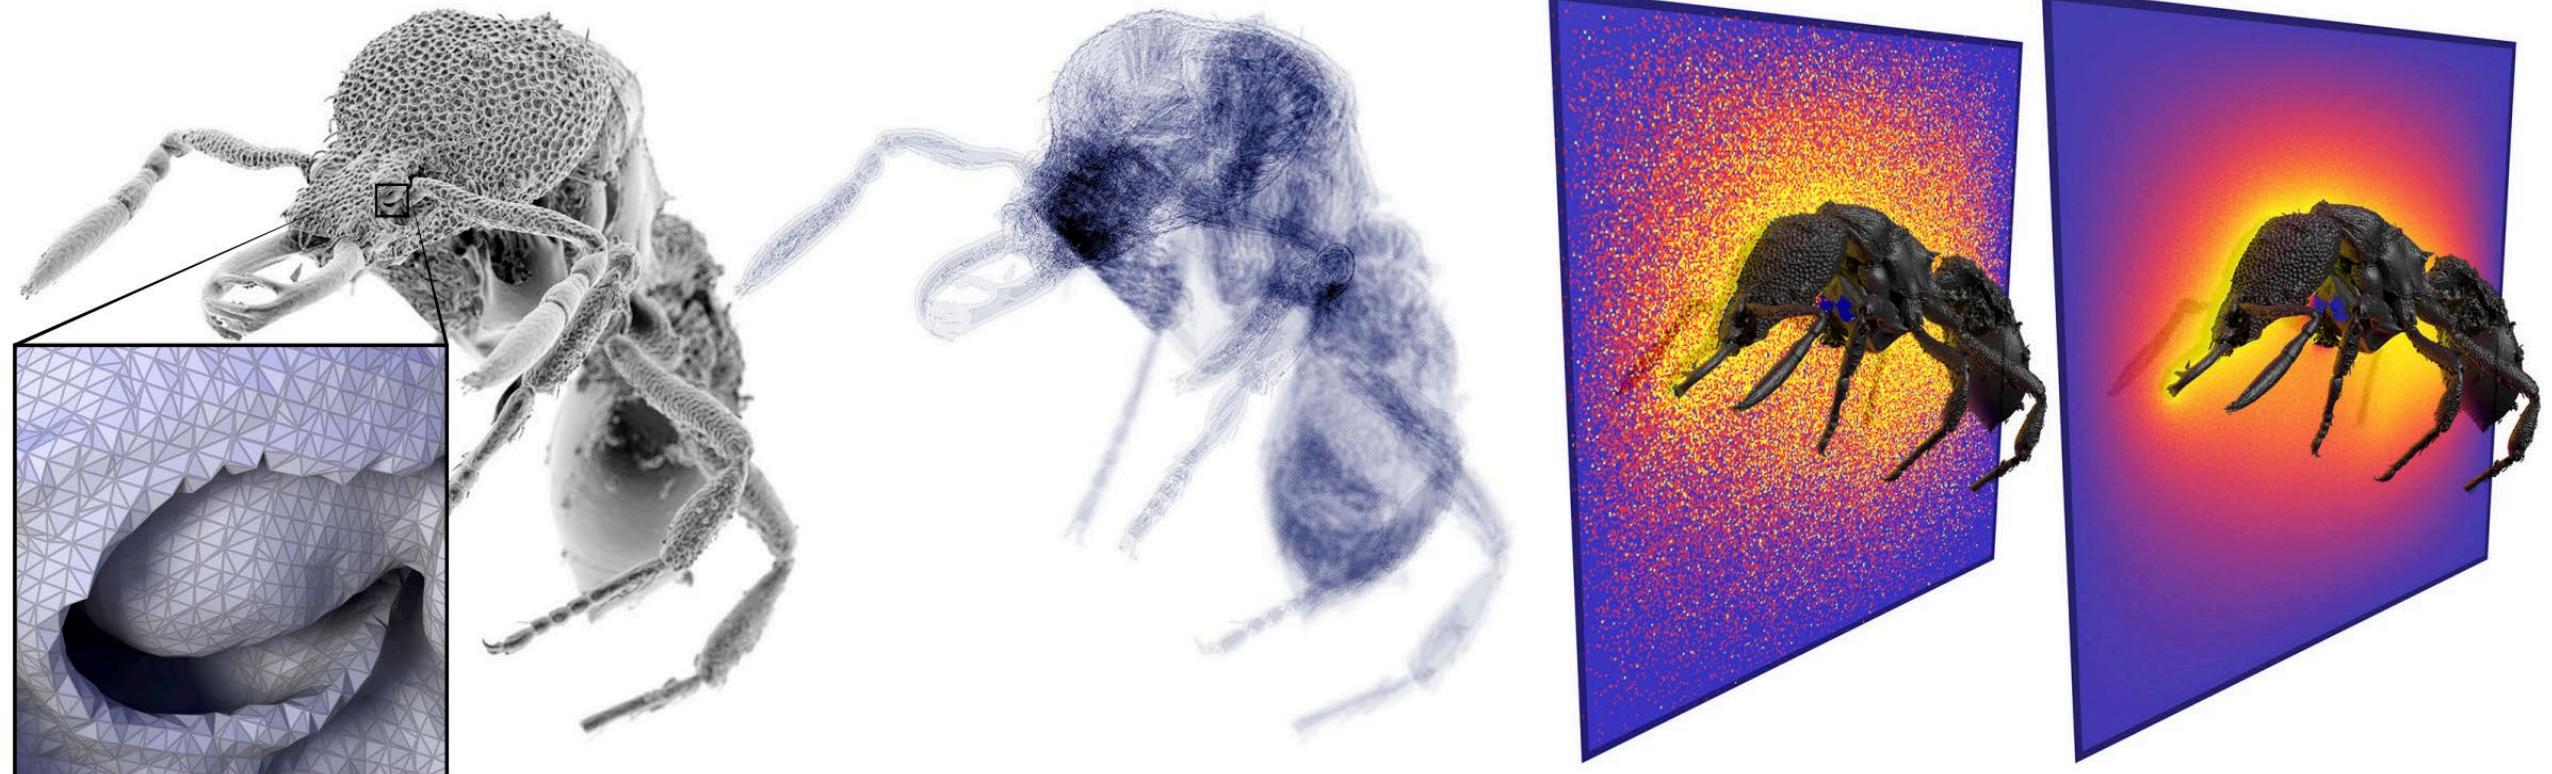
## Stream Line Visualization of the Flow around a Backward-Facing Step

Source:

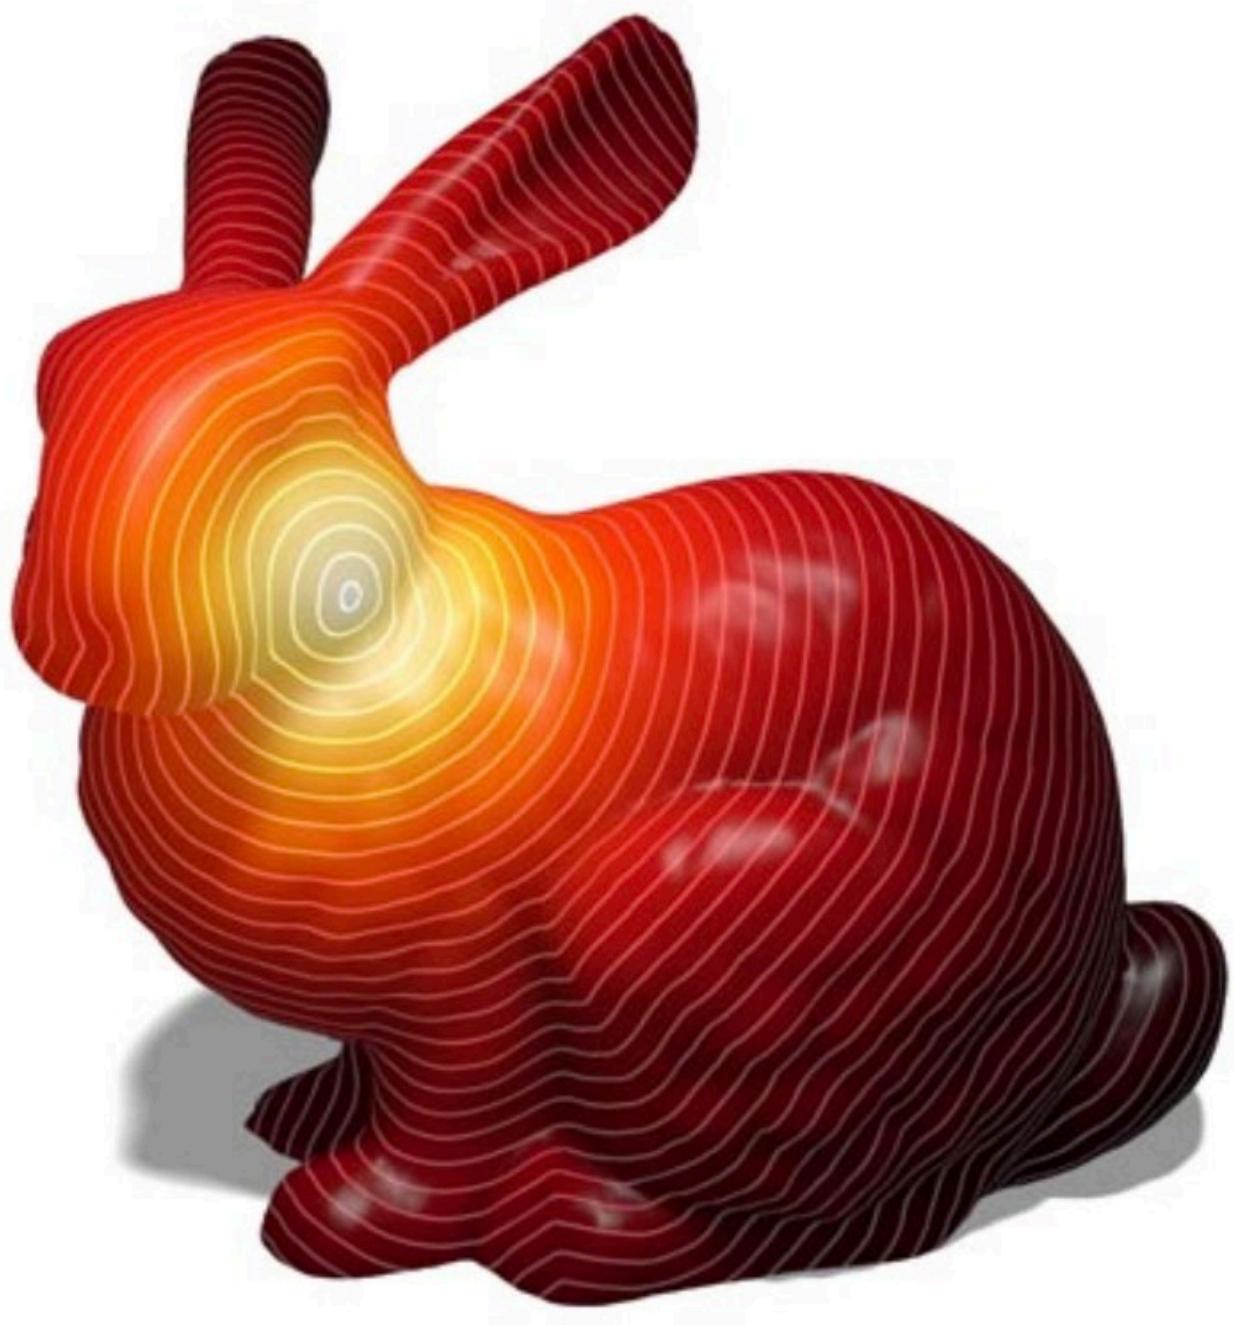
<https://www.csc.kth.se/~weinkauf/notes/step.html>



Source: <https://www.ebalstudios.com/blog/modeling-cars-polygons>

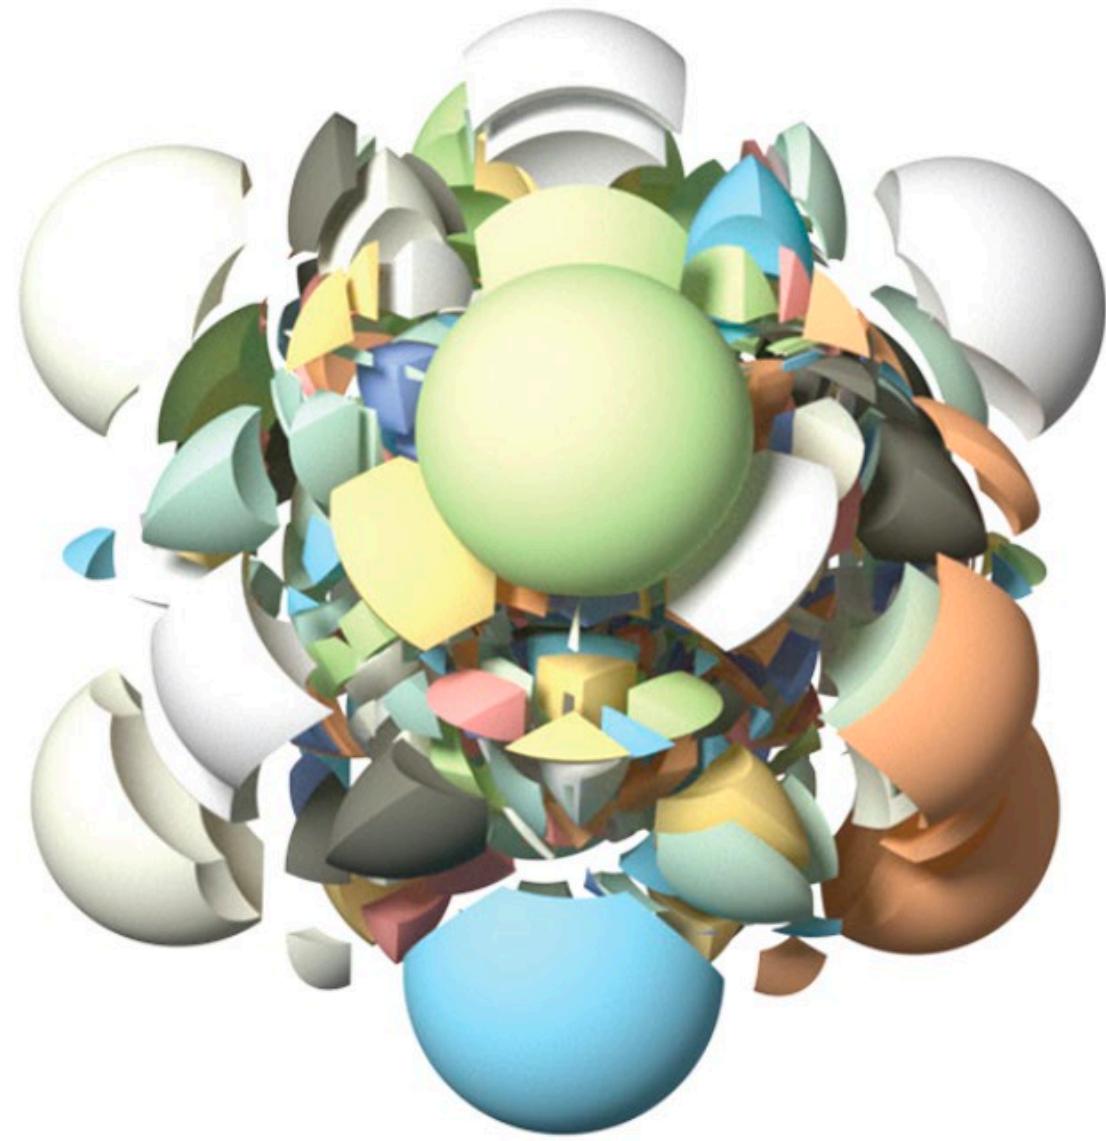
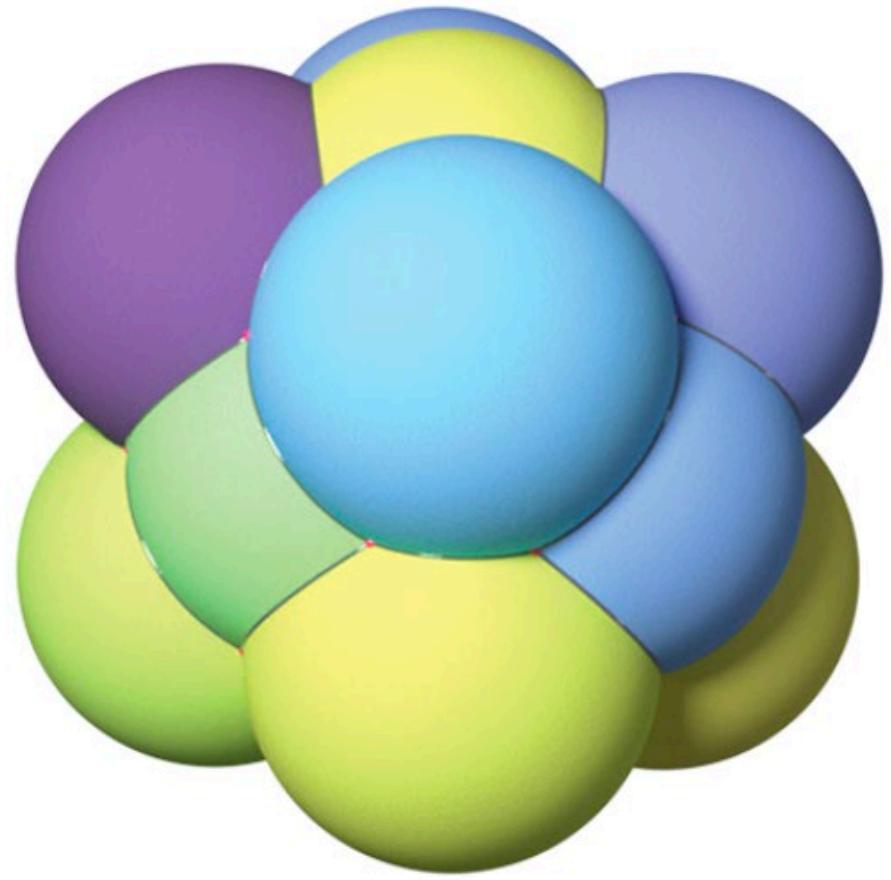


Source: <https://www.cs.cmu.edu/~kmcrane/Projects/MonteCarloGeometryProcessing/paper.pdf>



Source: <https://arxiv.org/pdf/1204.6216.pdf>







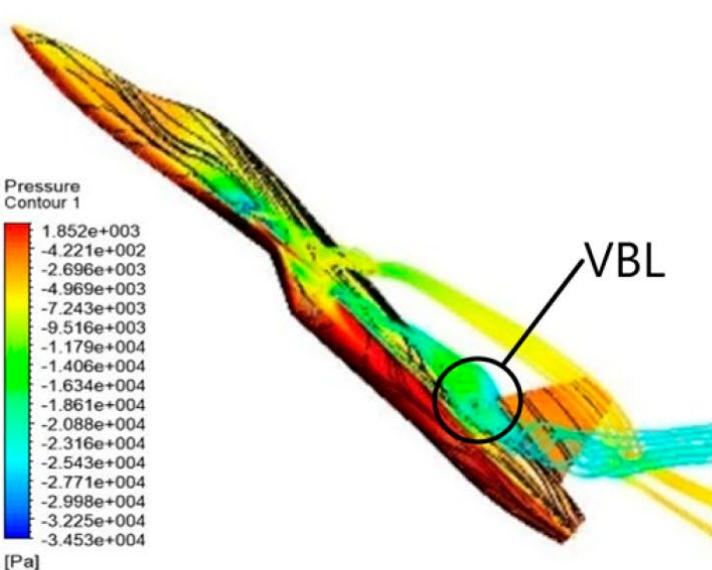
## A NASA study on wingtip vortices

Source:

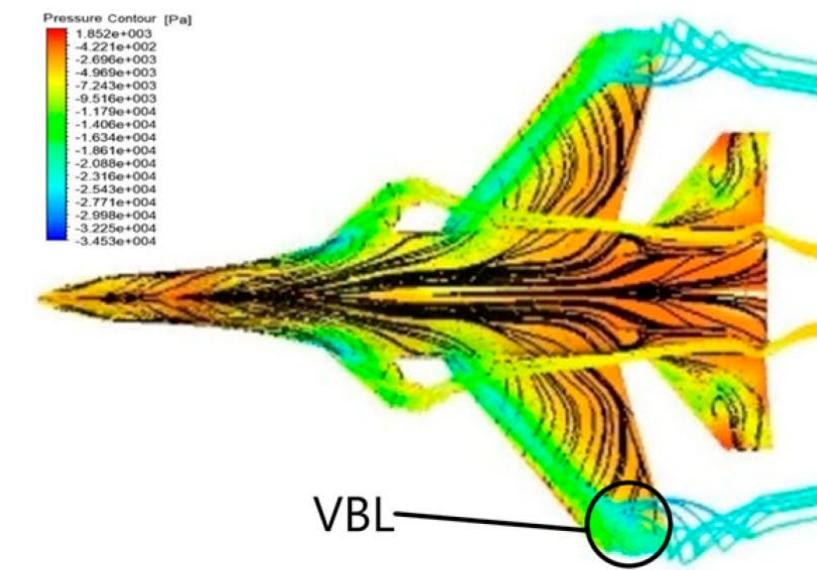
[https://en.wikipedia.org/wiki/Wingtip\\_vortices](https://en.wikipedia.org/wiki/Wingtip_vortices)



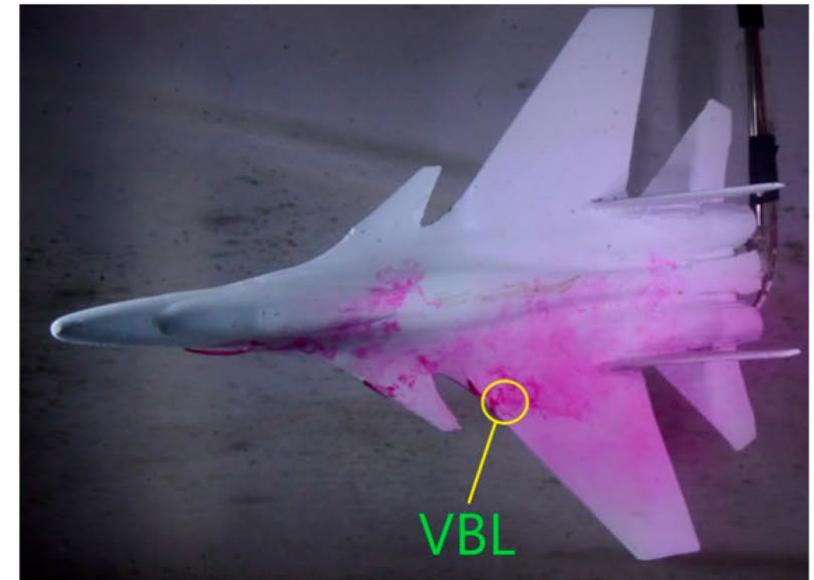
(a) WaTu AoA 30°



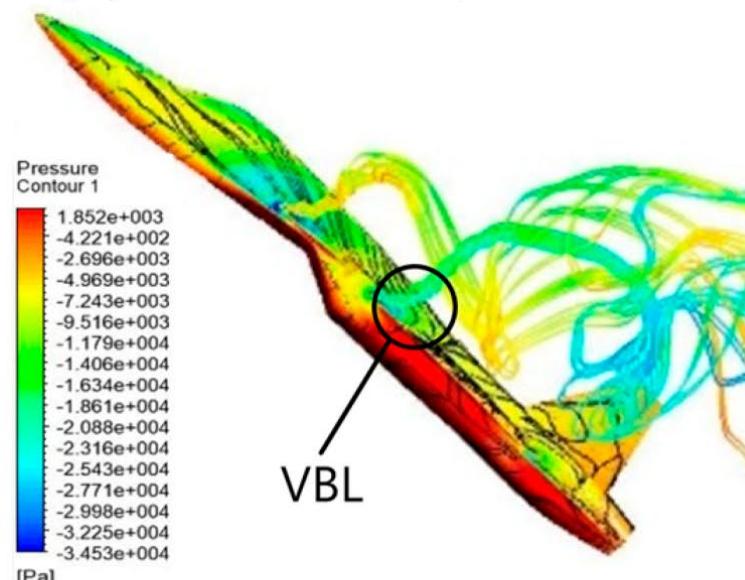
(b) CFD AoA 30°; side view



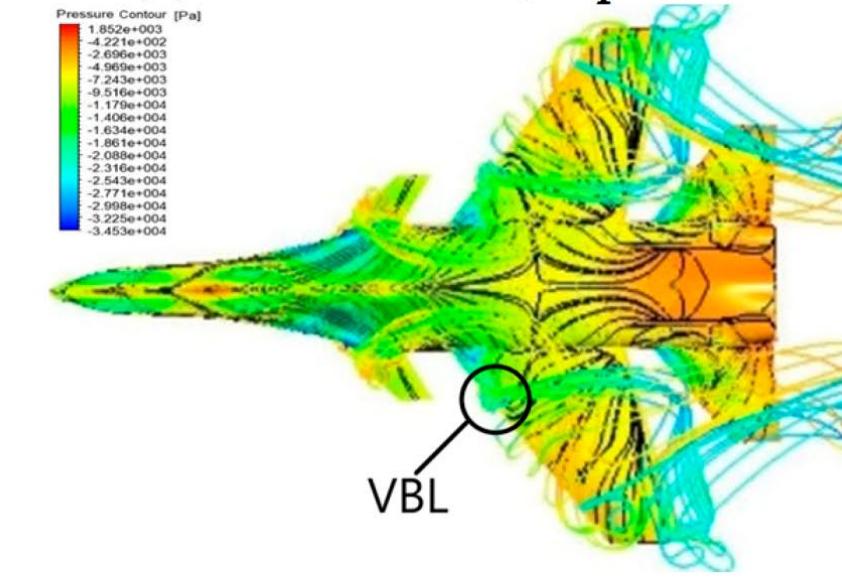
(c) CFD AoA 30°; top view



(d) WaTu AoA 50°

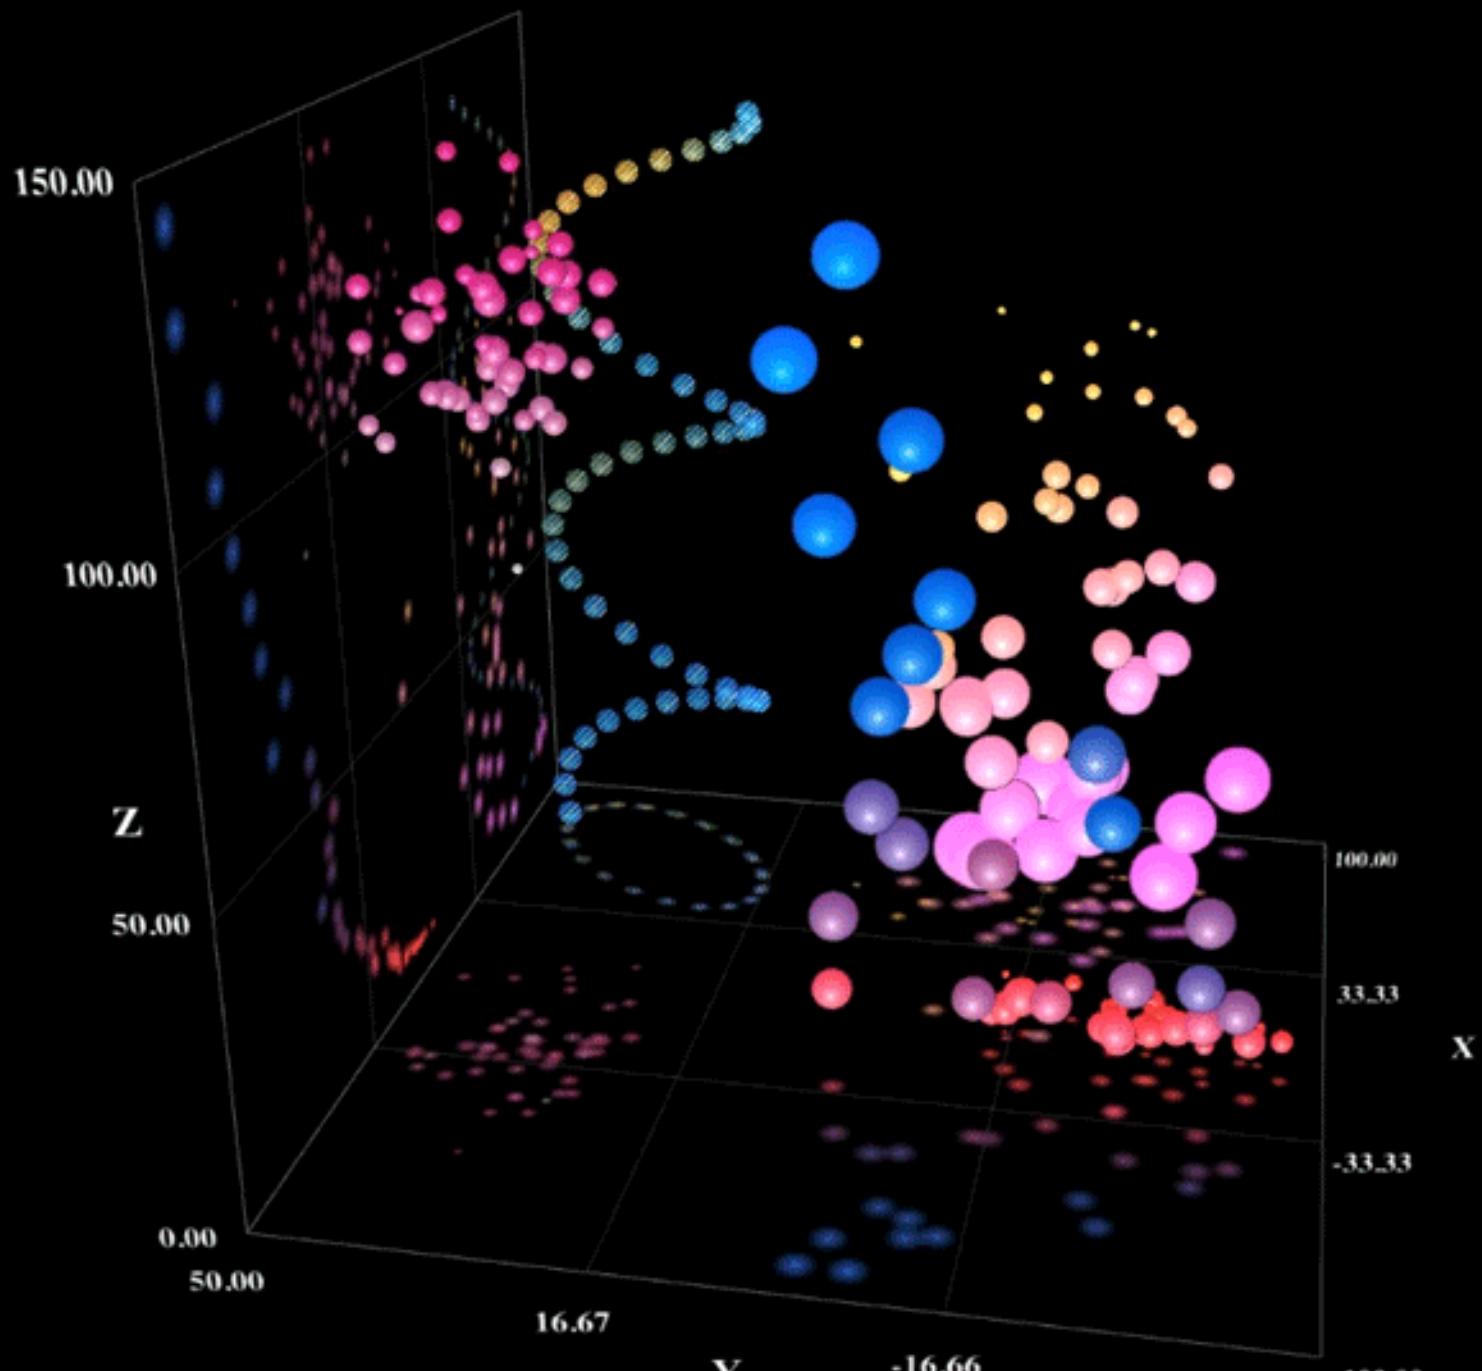


(e) CFD AoA 50°; side view

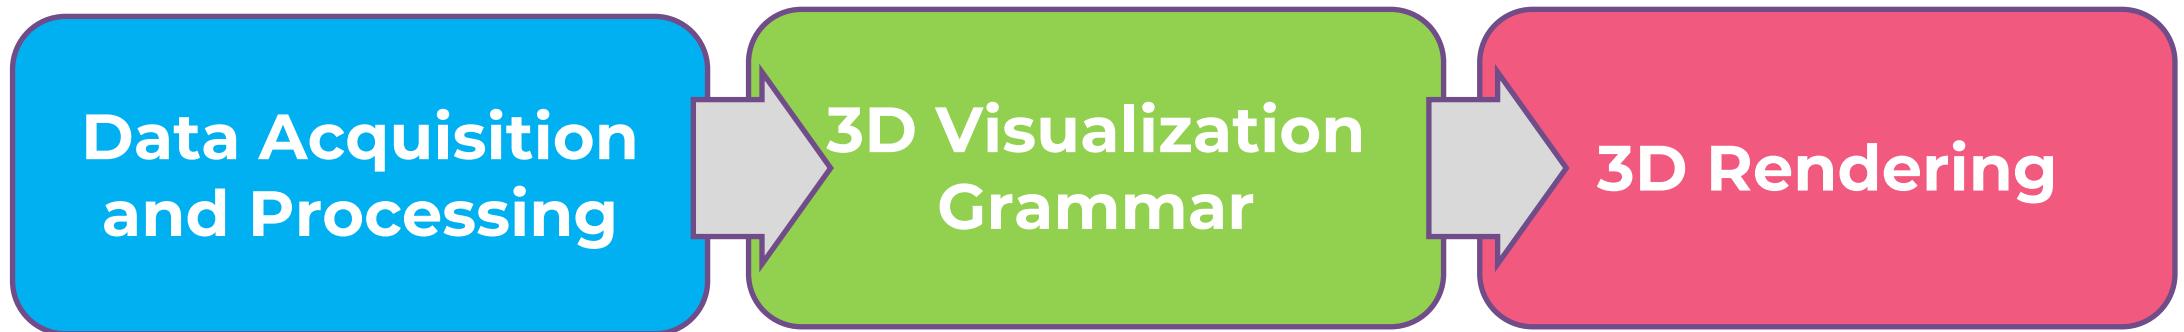


(f) CFD AoA 50°; top view

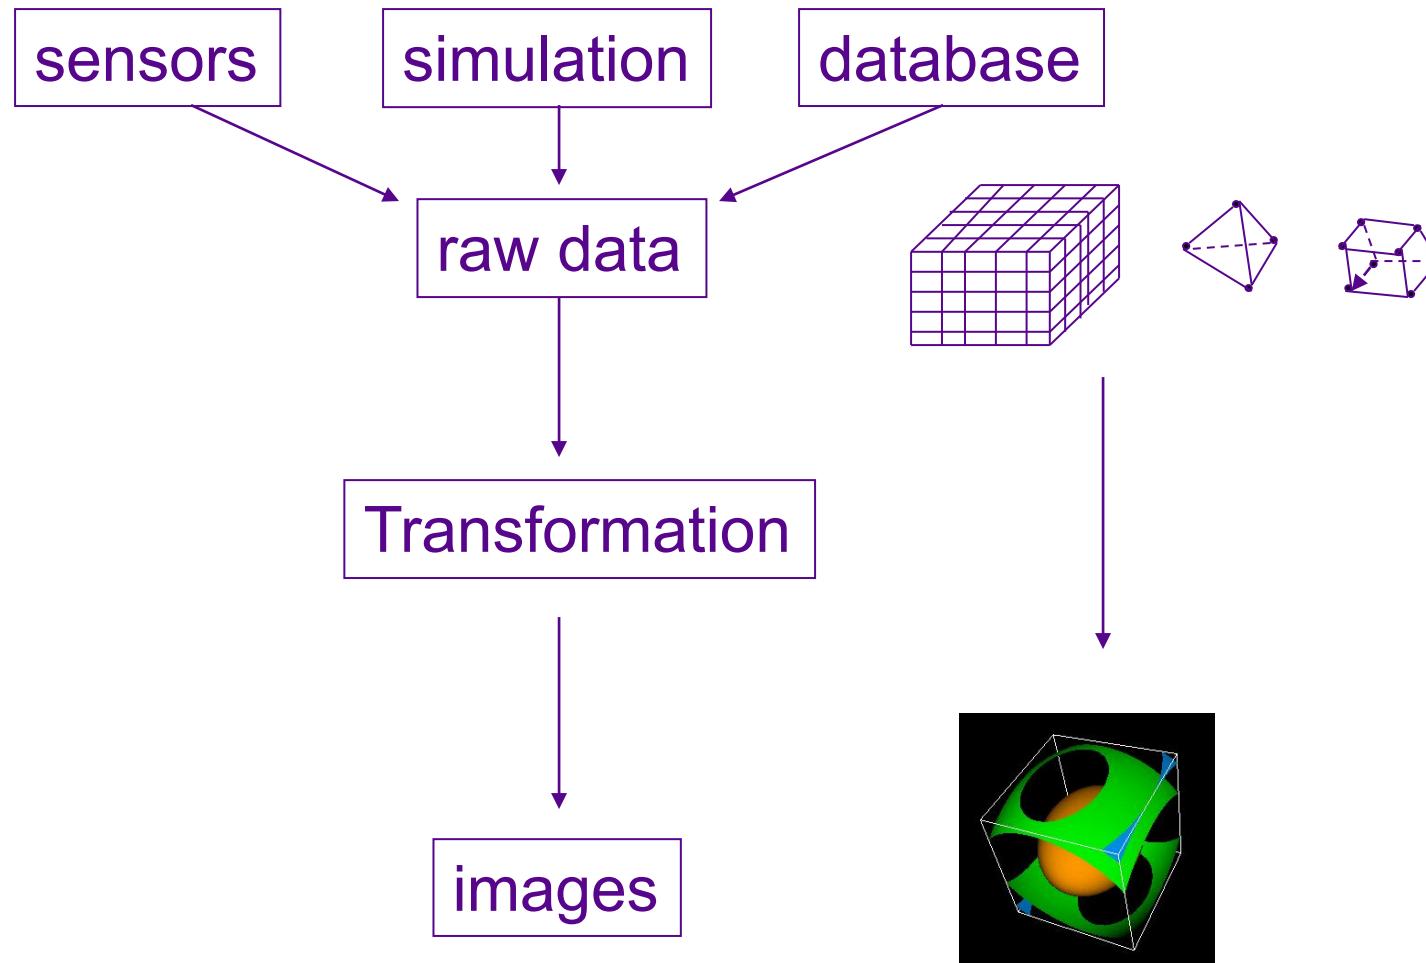




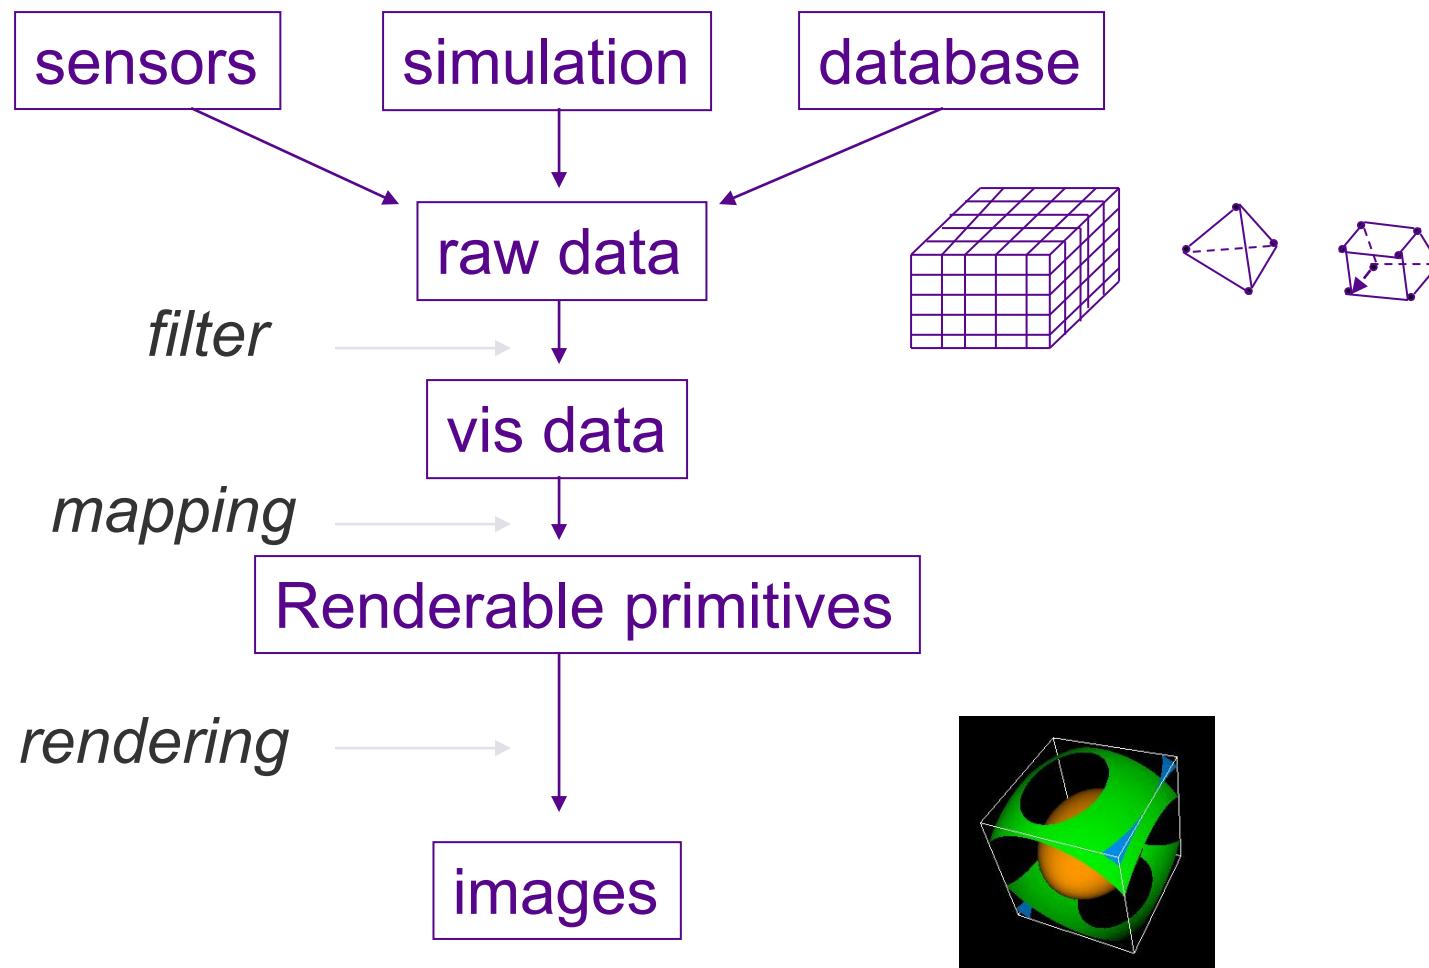
# 3D Visualization Pipeline



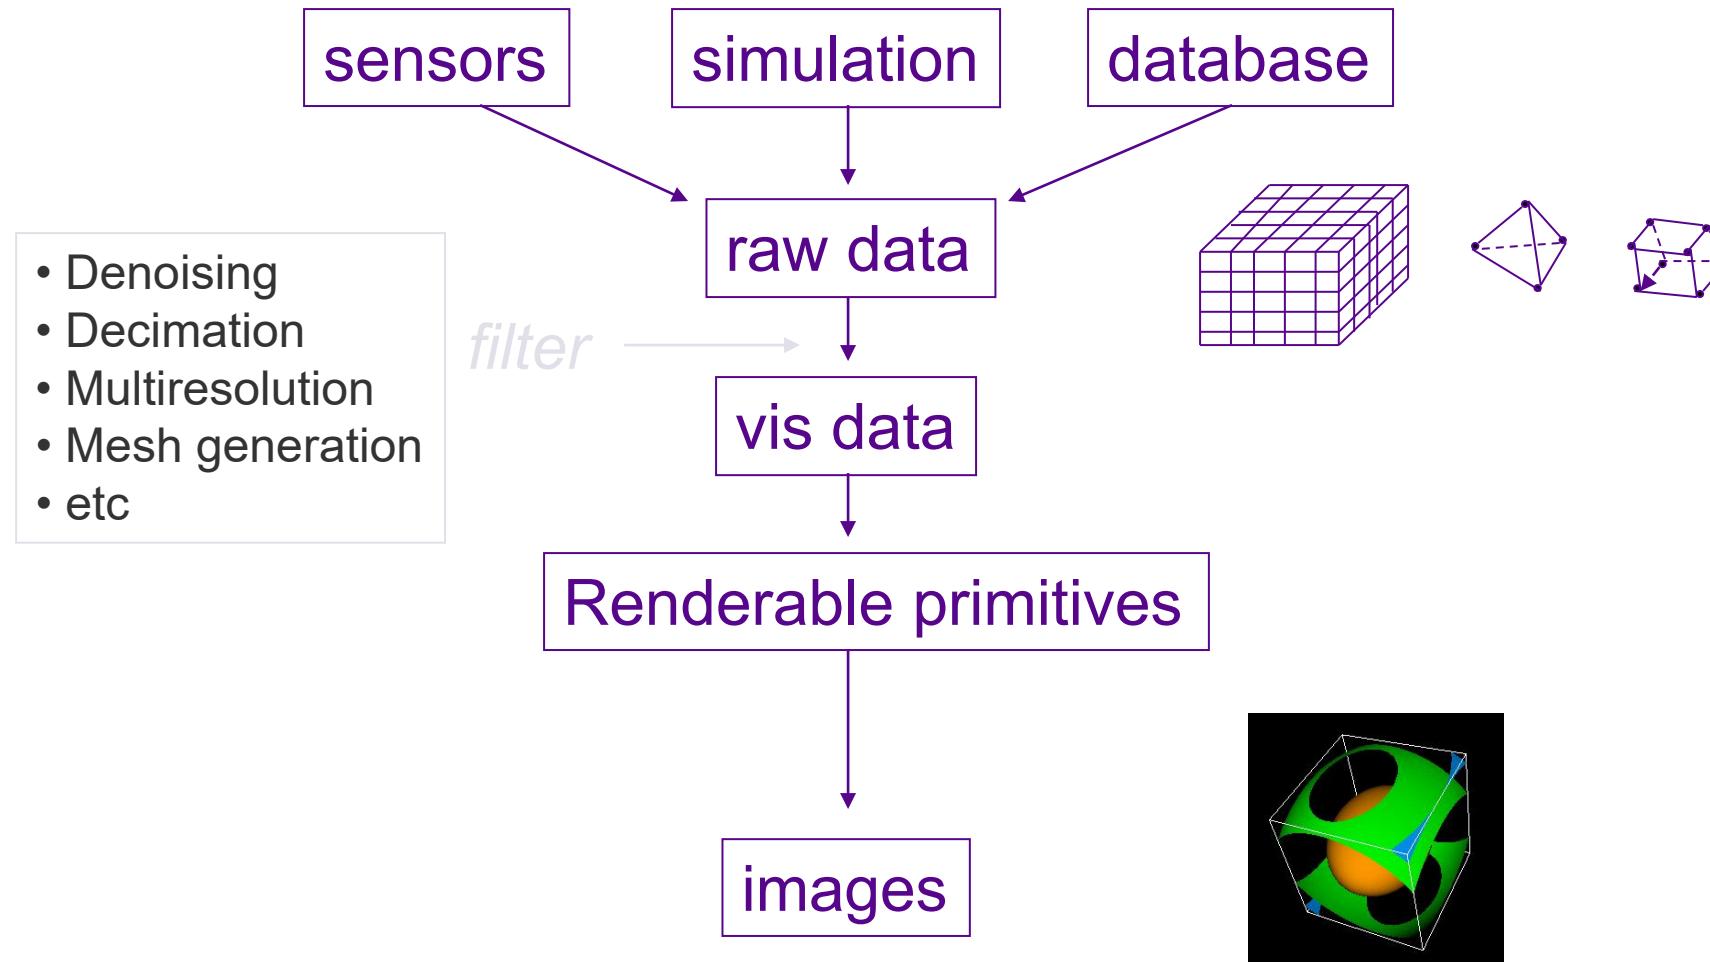
# Visualization Process



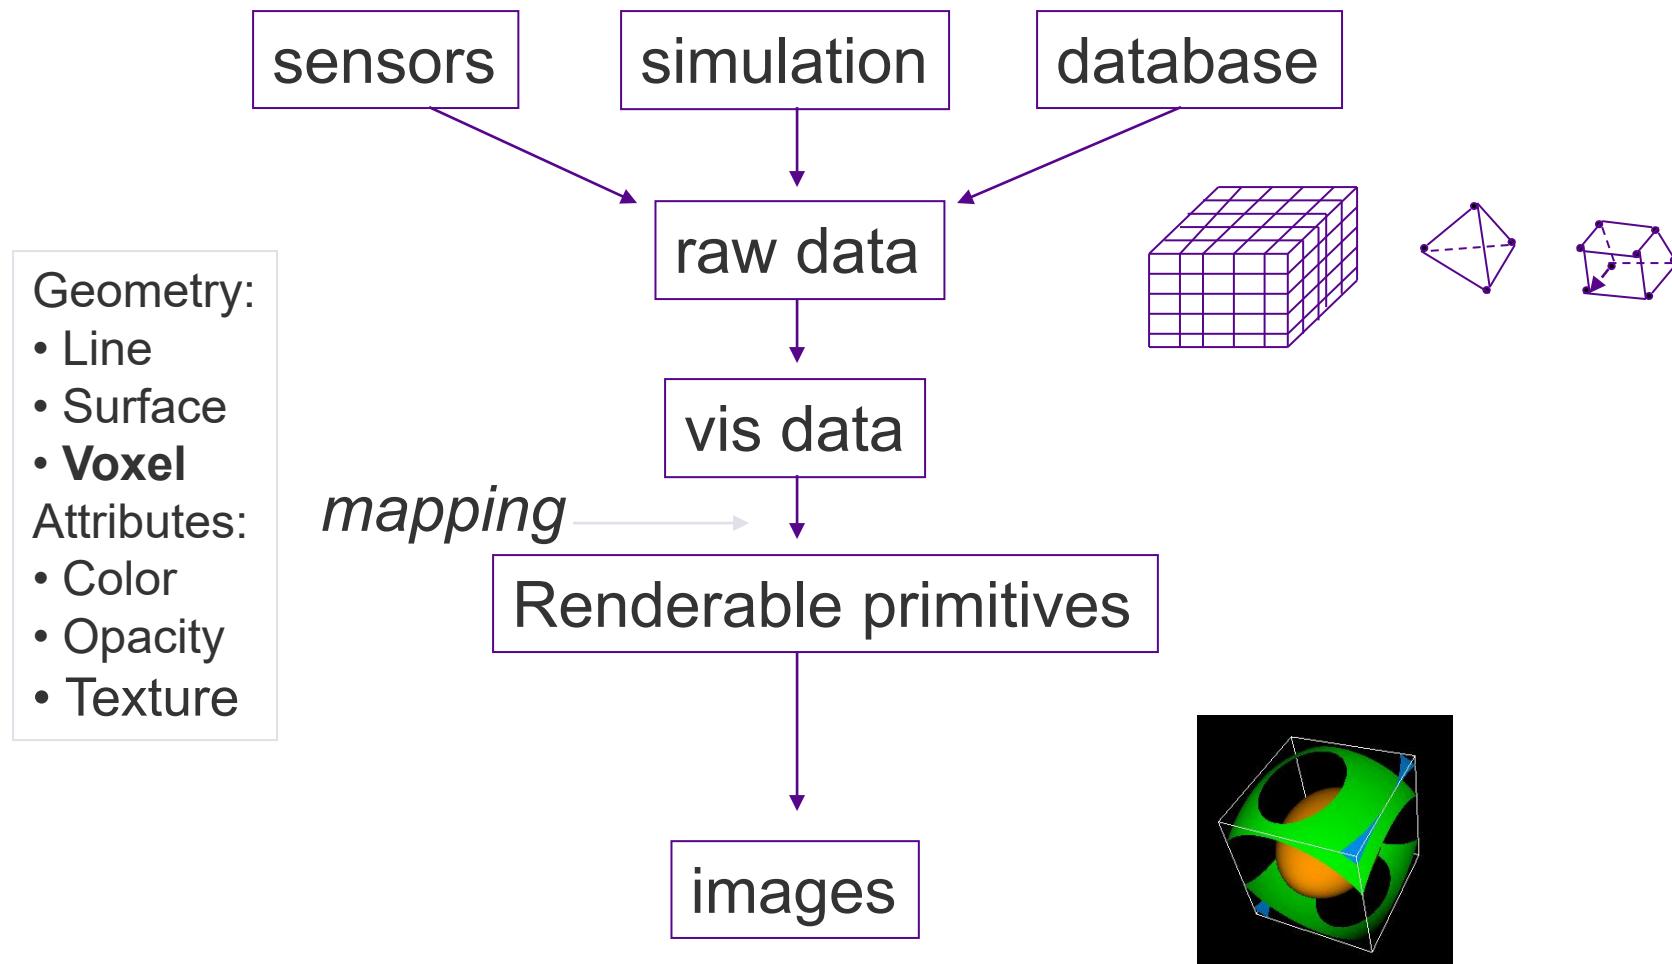
# Visualization Pipeline



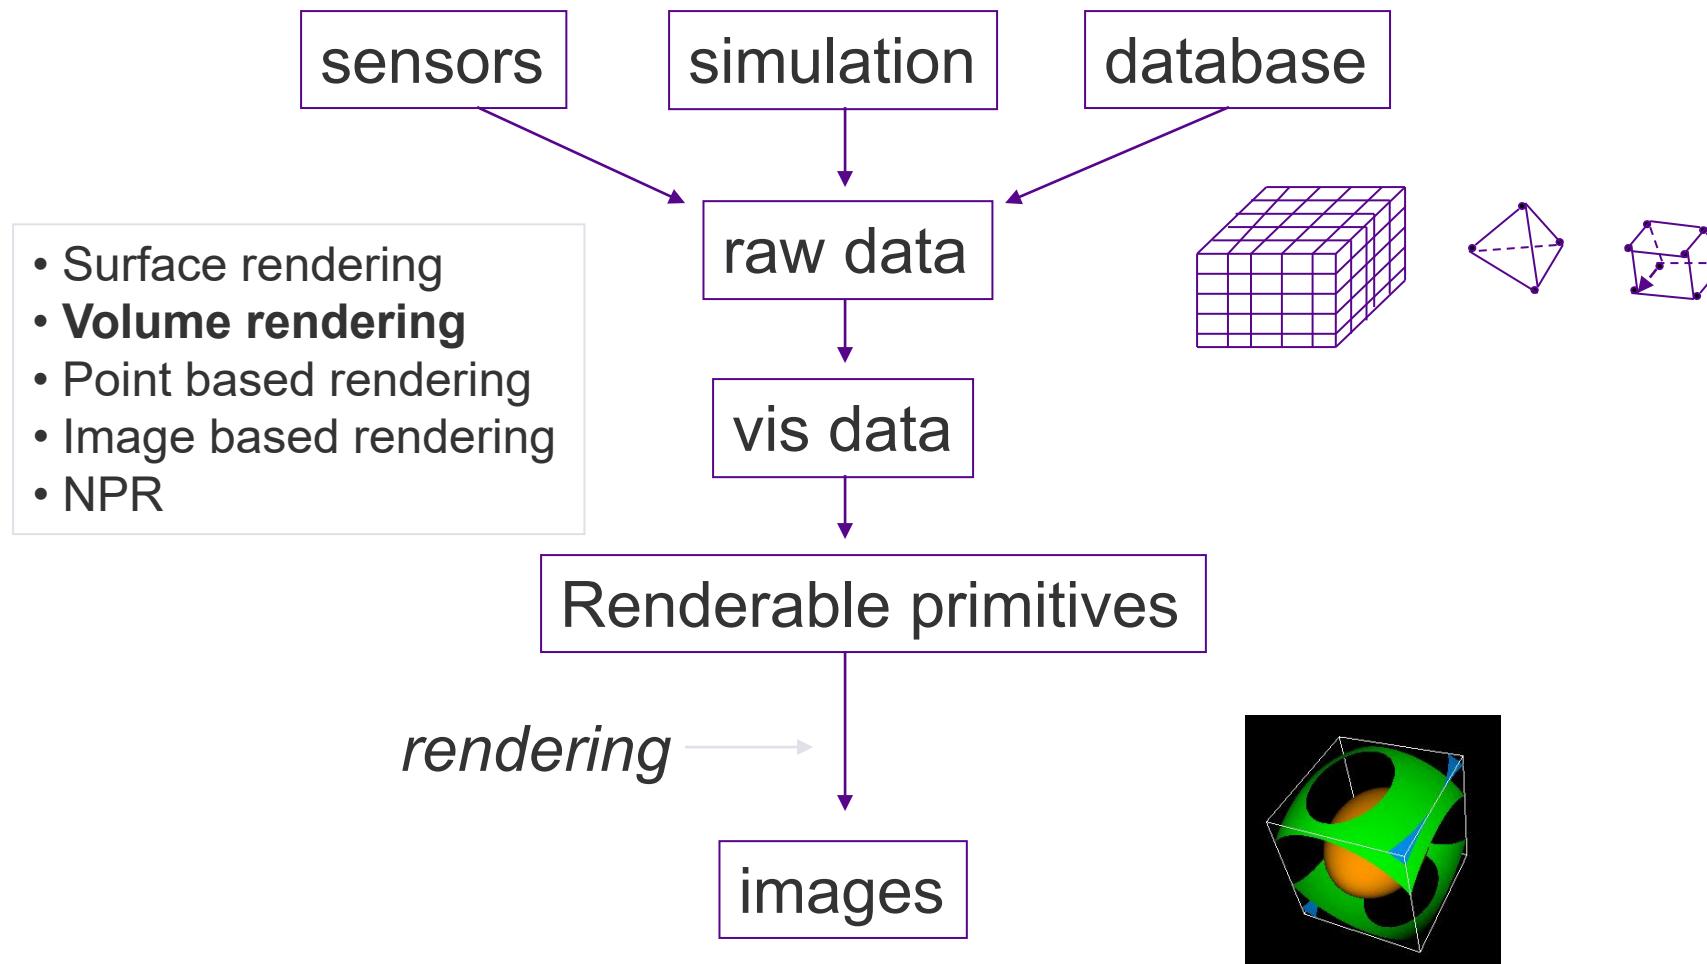
# Visualization Pipeline



# Visualization Pipeline

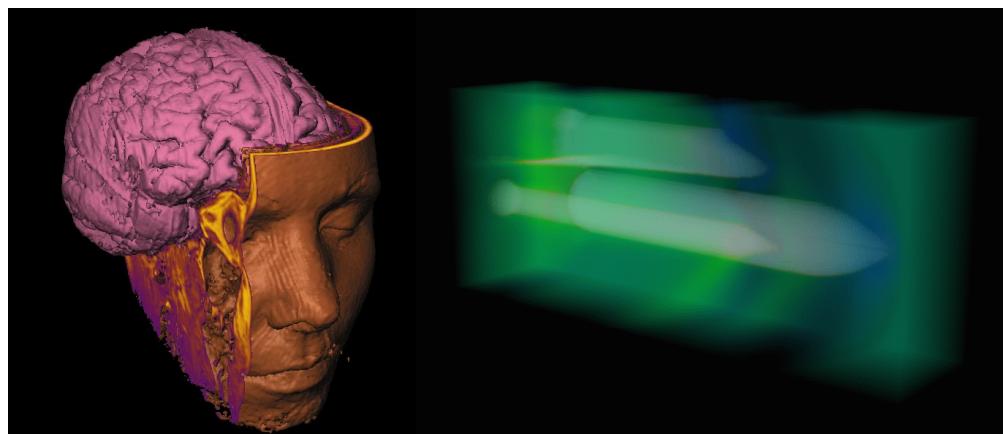


# Visualization Pipeline

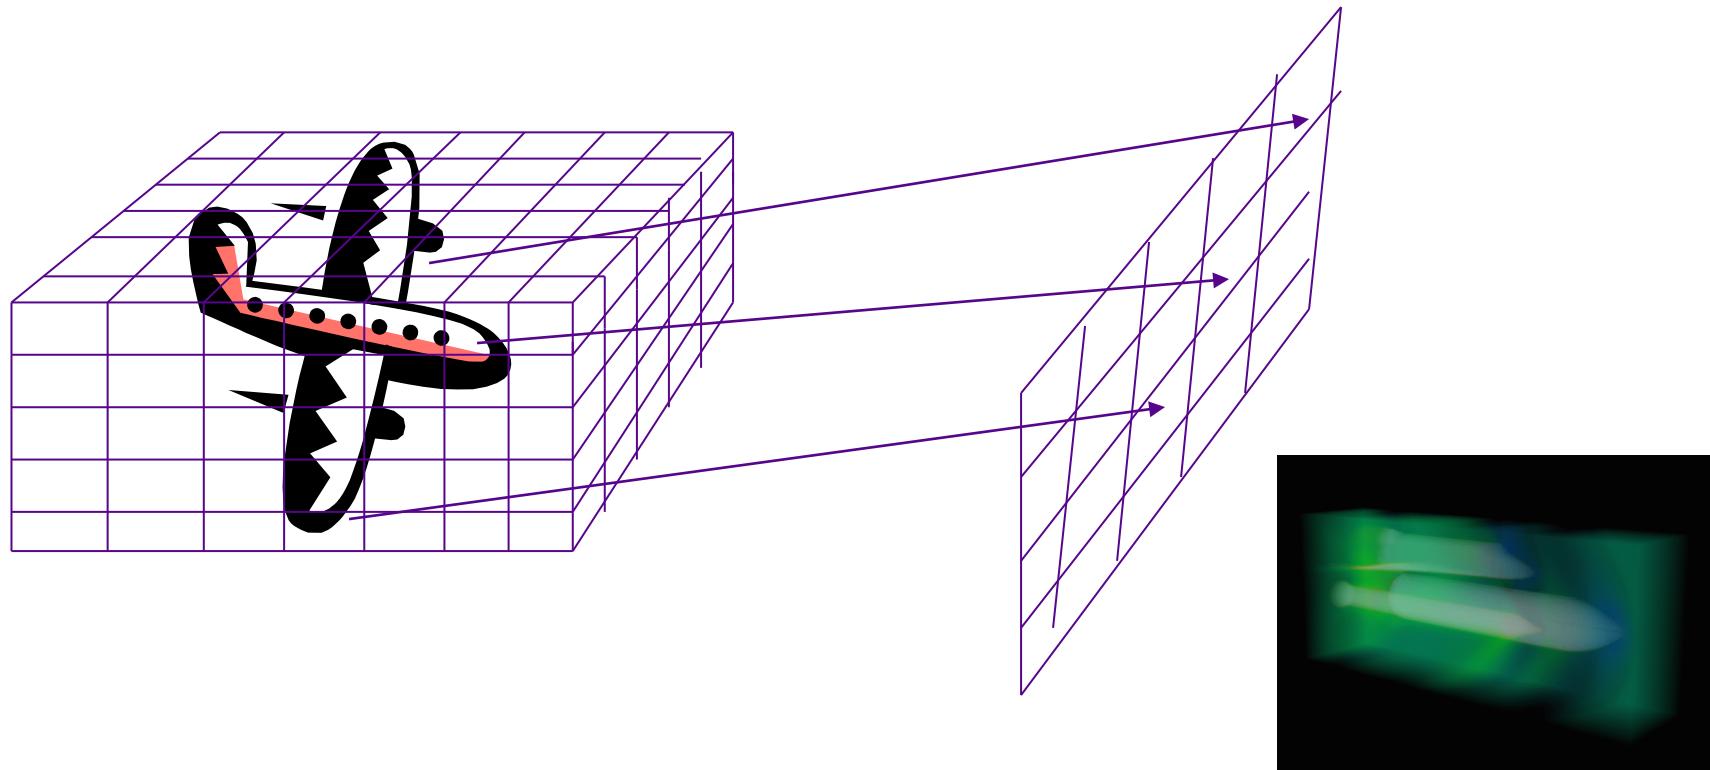


# Volume Rendering

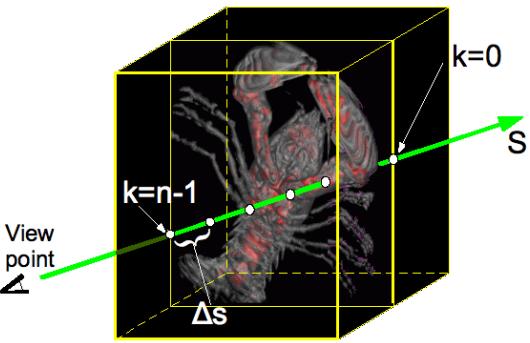
- Goal: visualize three-dimensional functions
  - Measurements (medical imaging)
  - Numerical simulation output
  - Analytic functions



# Volume Rendering

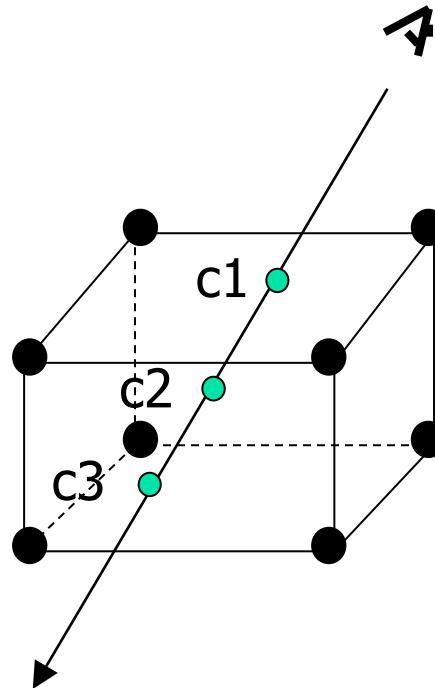
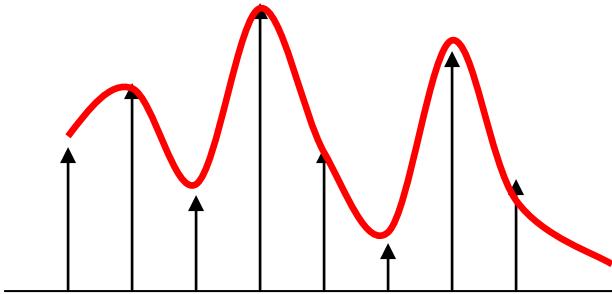


# Volume Data Visualization



# Ray Casting

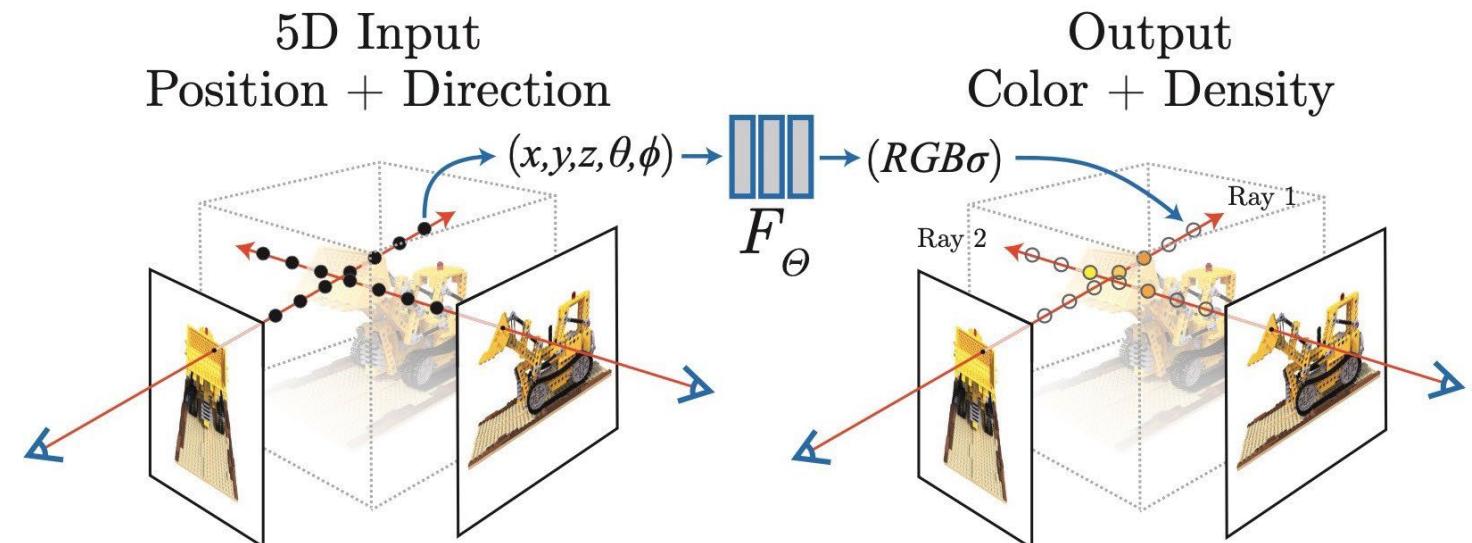
Recover the original continuous function from discrete samples



# View Synthesis from Volume Rendering: Neural Radiance Field

- ❖ Given a dataset containing RGB images of a static scene, their corresponding camera poses, and intrinsic parameters,
- ❖ Predict the color and volume density for every viewing location and direction

Inputs:	
❖ $x, y, z$	Target position of the camera
❖ $\theta, \phi$	Target orientation of the camera
Outputs:	
❖ $c = (R, G, B)$	Color
❖ $\sigma$	Volume density



# NeRF: Volume Rendering

- ❖ Generating a view from NeRF requires rendering all rays that pass through each pixel of the desired virtual camera

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt, \text{ where } T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

Expected color of a camera ray      Predicted Volume Density      Predicted Color      Probability that nothing has blocked the ray up to this point

- ❖ Numerically estimated using quadrature and stratified sampling

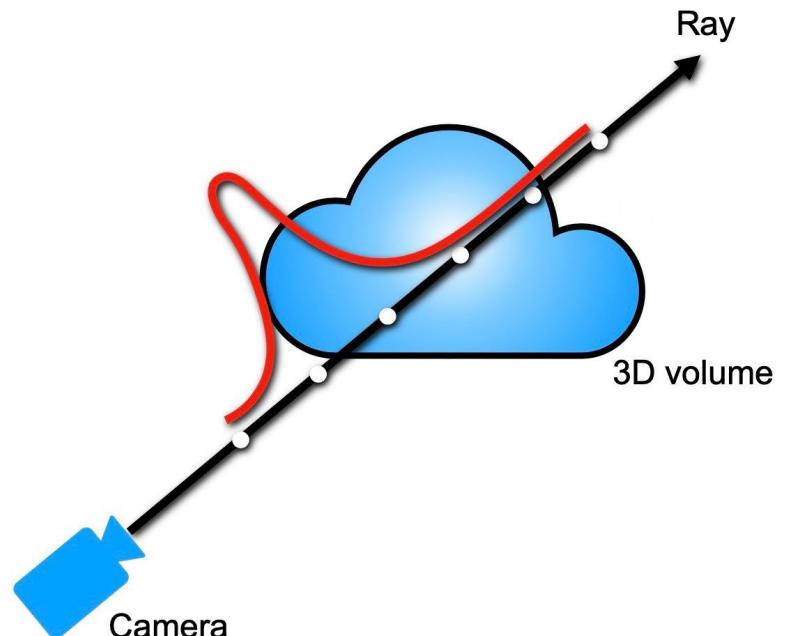
$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i, \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

The relative contribution of this segment      Predicted Color      Probability that nothing has blocked the ray up to this point

- ❖ Differentiable: allows optimization using gradient descent

# NeRF: Hierarchical Sampling

- ❖ Problem: It is inefficient to integrate over empty and occluded spaces in a scene
- ❖ Solution: Allocate samples proportionally to their expected effect on the final rendering.
- ❖ Evaluate a “coarse” network on a set of  $N_c$  locations along a ray to produce a PDF along the ray
- ❖ Evaluate a “fine” network on  $N_c$  and a second a set of  $N_f$  locations sampled from the PDF



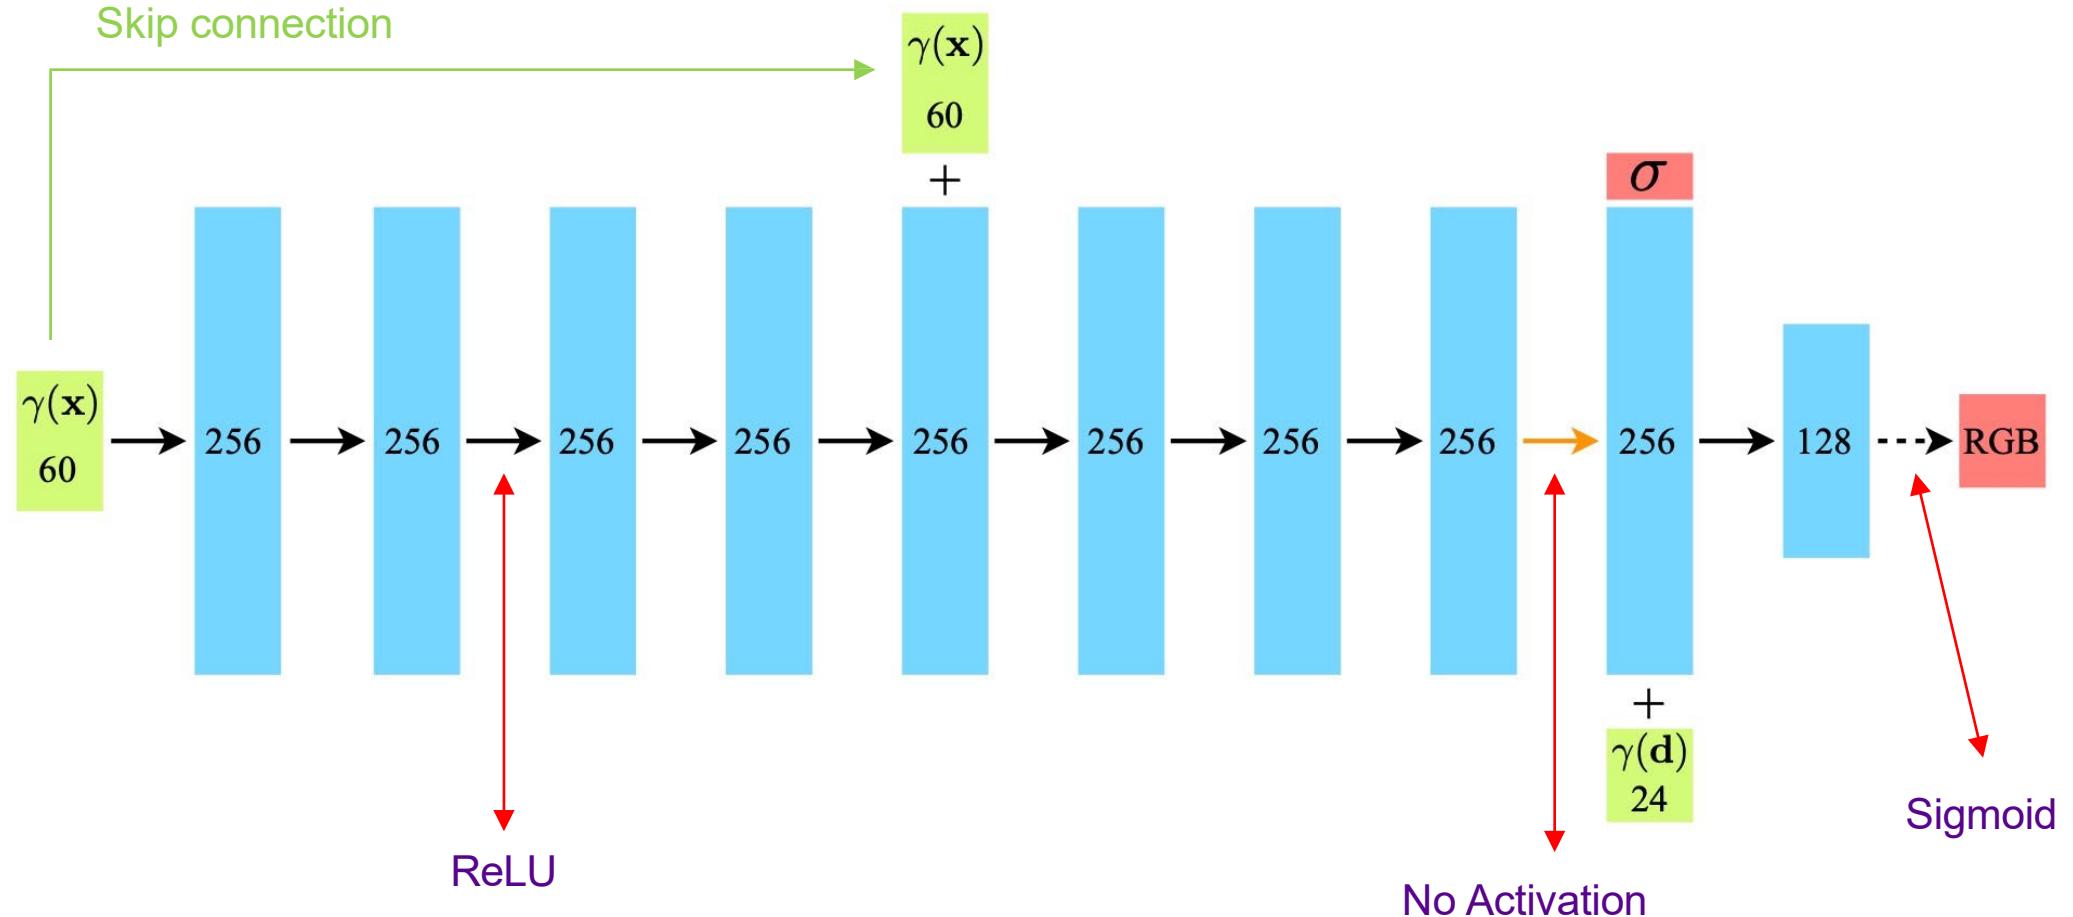
# NeRF: Positional Encoding

- ❖ Map individual components of position and direction vectors to a higher dimensional space

$$\gamma(p) = (\sin(2^0 \pi p), \cos(2^0 \pi p), \dots, \sin(2^{L-1} \pi p), \cos(2^{L-1} \pi p))$$

- ❖ Similar concept as positional encoding in Transformer Architectures
- ❖ Empirically improves the preservation of high-frequency geometry and texture
- ❖ Surprising result, explored further in a follow-up work by the same authors
  - ❖ “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains” (Tancik, Srinivasan, Mildenhall, et al., 2020)

# Network Architecture



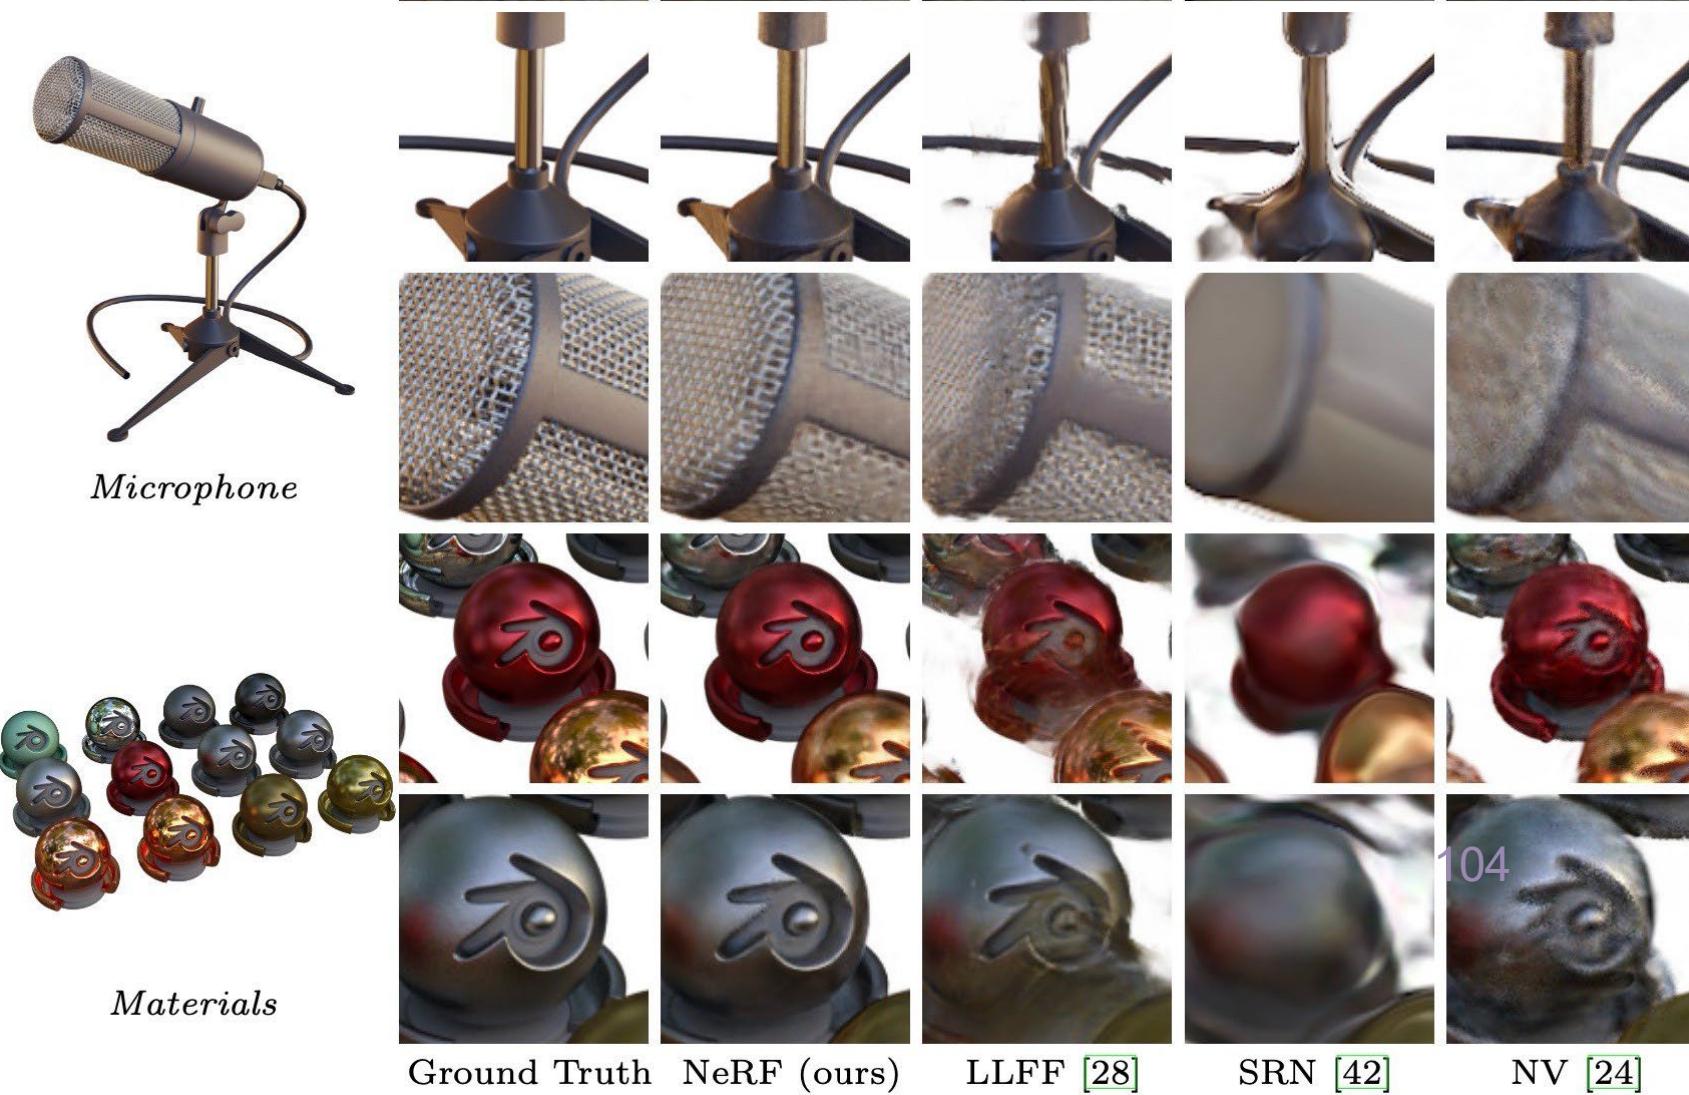
# Training Summary

- ❖ Sample a batch of camera rays from the dataset (bs=4096)
- ❖ Use hierarchical sampling to query coarse and fine points
- ❖ Use the volume rendering equation to calculate the color of the ray
- ❖ Compute the squared error between rendered and true pixel colors

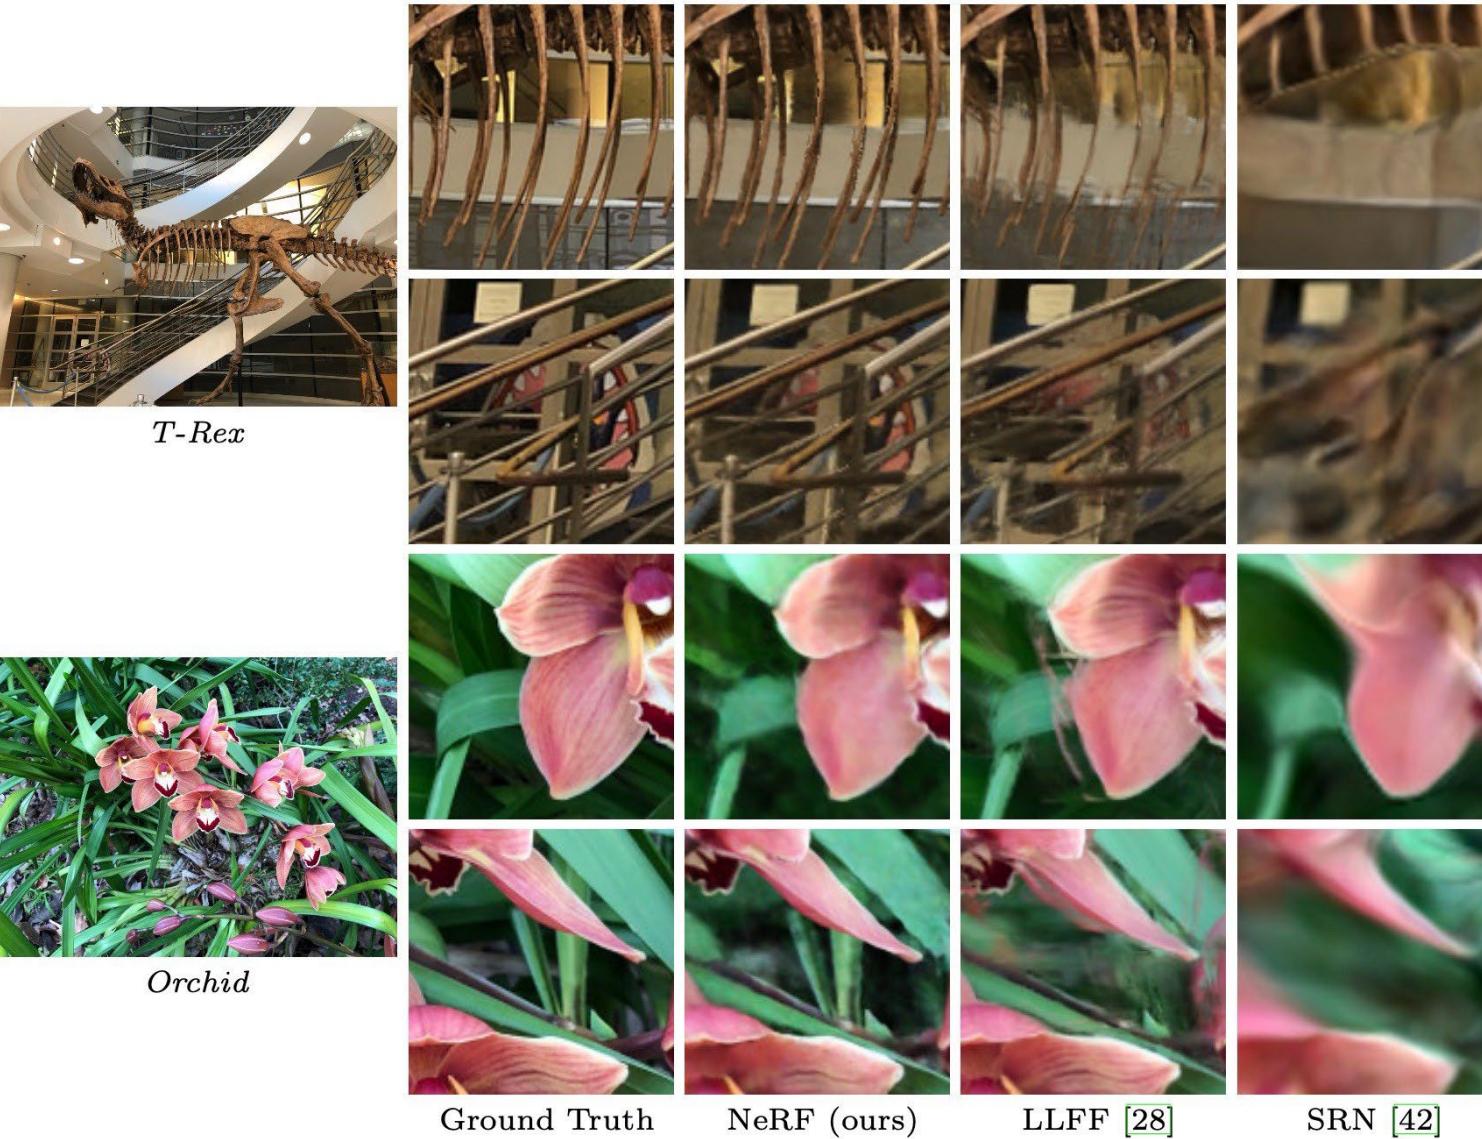
$$\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \left[ \left\| \hat{C}_c(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 + \left\| \hat{C}_f(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 \right]$$

- ❖ Optimize network parameters using Adam

# Qualitative Results: Realistic Synthetic Objects

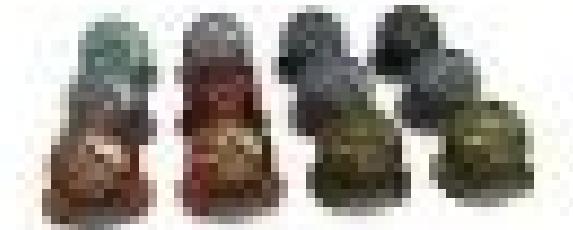


# Qualitative Results: Real-world scenes

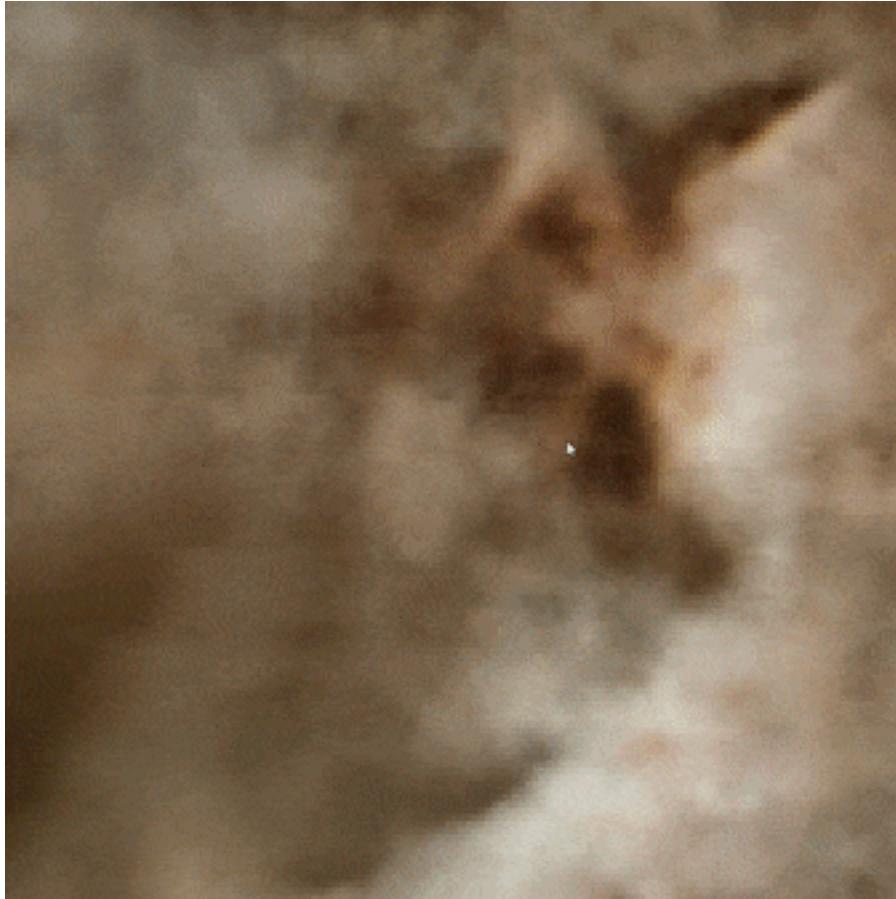


# Animated Results

---



# Instant Neural Graphics Primitives



<https://github.com/NVlabs/instant-ngp>

# Applications



Real-time rendering

# Readings

- ❖ Neural Volumes (NV), 2019
  - ❖ Deep 3D convolutional network architecture
  - ❖ Predicts a fixed-size discretized voxel grid
  - ❖ Limitation: discrete voxel grids do not scale well and lose fine detail at high resolutions
  - ❖ Limitation: requires a bounded volume and knowledge of the background
- ❖ Scene Representation Networks (SRN), 2019
  - ❖ Uses a recurrent neural network to model a rendering function
  - ❖ Limited to simple shapes with low geometric complexity
- ❖ Local Light Field Fusion (LLFF), 2019
  - ❖ 3D convolutional network architecture
  - ❖ Predicts multiplane images and fuses them to create new views
  - ❖ Fast to train (<10 minutes) at the cost of large storage requirements (~GB for each scene)

---

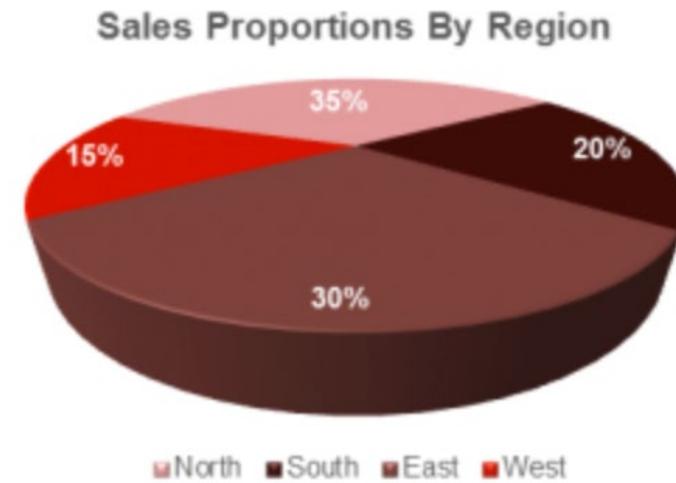
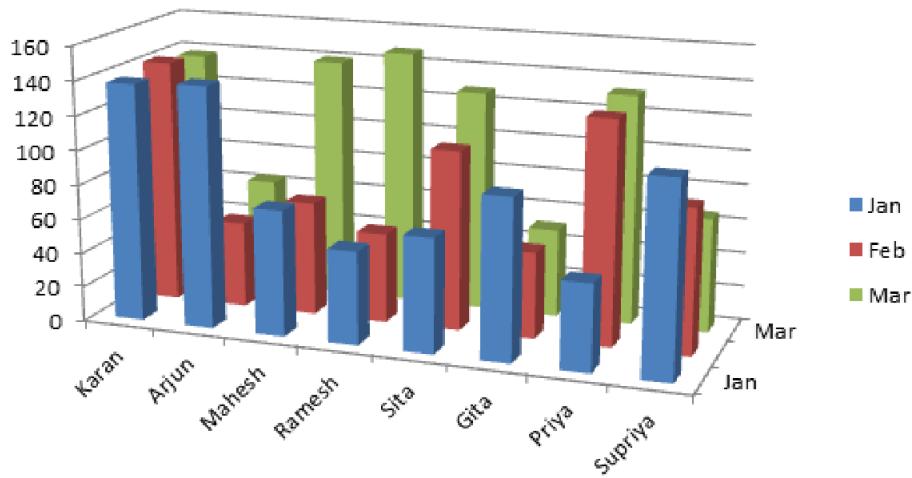
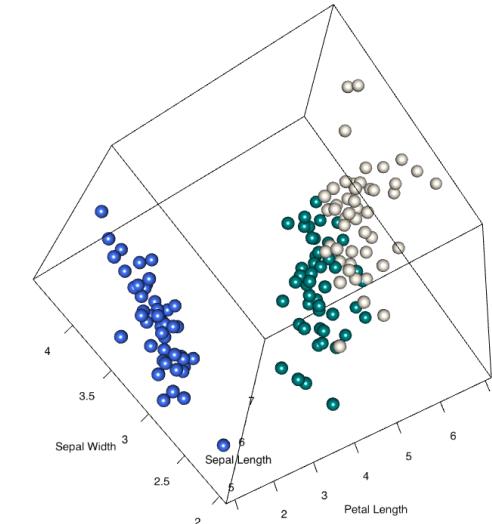
# 3D Visualization for non-3D Data

A wooden relief sculpture of a mountain range, featuring several peaks and valleys carved into a light-colored wooden base. The sculpture is set against a plain, light-colored wall.

To 3D or not to 3D?

## When not to 3D?

- 2D visualization is sufficient.
- 3rd dimension encode no information.
- Size/lengths are foreshortened.
- Occlusion.

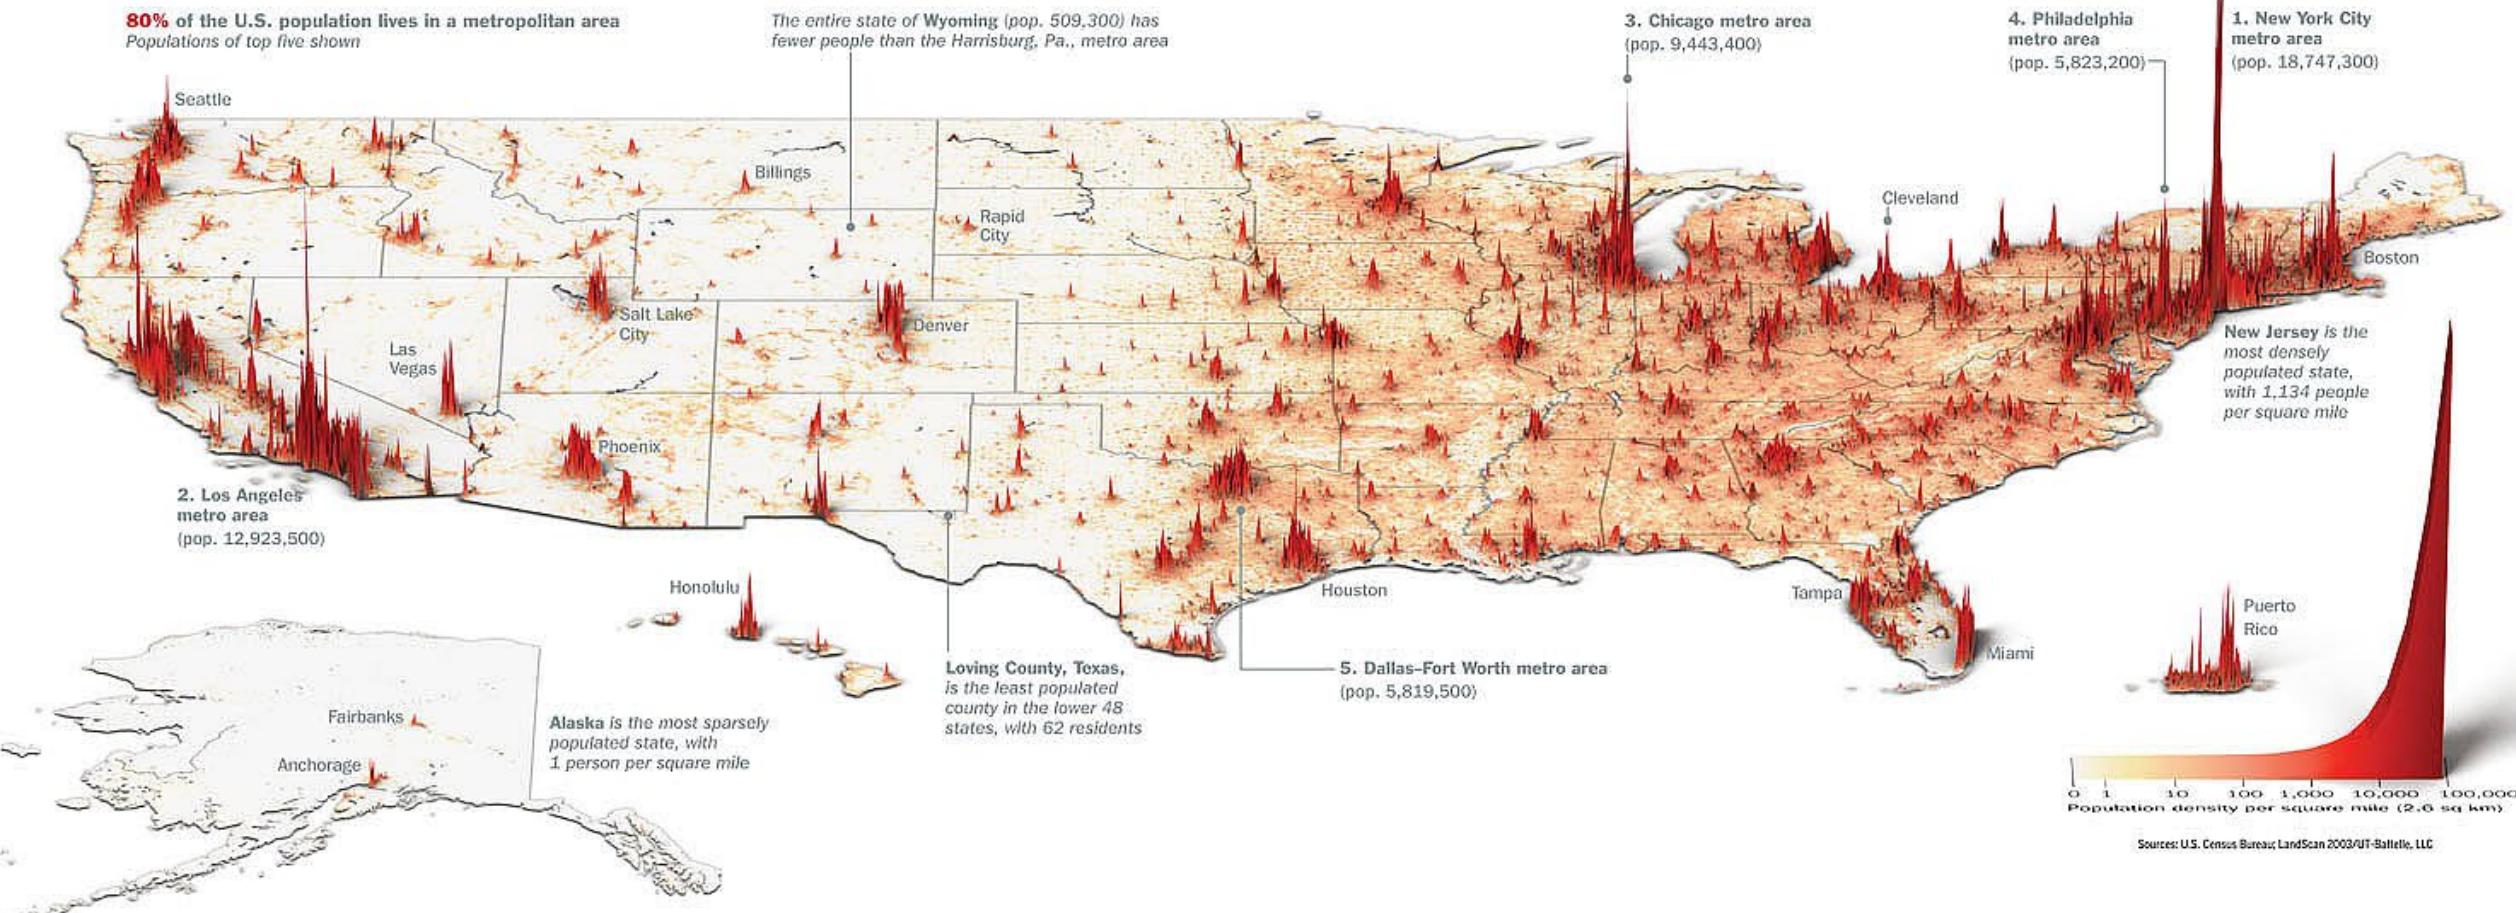


# Non-3D Data

- Data that are not intrinsically 3D-related.
  - No variables that are directly related to x, y and z.
  - E.g. stock price, house price, votes, population, etc.

# Where We Live...

Unlike many developed countries, the U.S. keeps growing. We are also moving south and west. But compared with China or India, the nation is a vast prairie



Our families are getting smaller—with one vital exception. Compared with those of Europe and Japan, the U.S. population is younger and more colorful because of the continued arrival of immigrants and their higher-than-average birthrates. Of the 100 million Americans who will join us in the next 37 years, half will be immigrants or their children. In the next few decades, 97% of the world's population growth will occur in the developing world; the U.S. is the largest developed country in the world that is still growing at a healthy clip. That matters, strategically, economical-

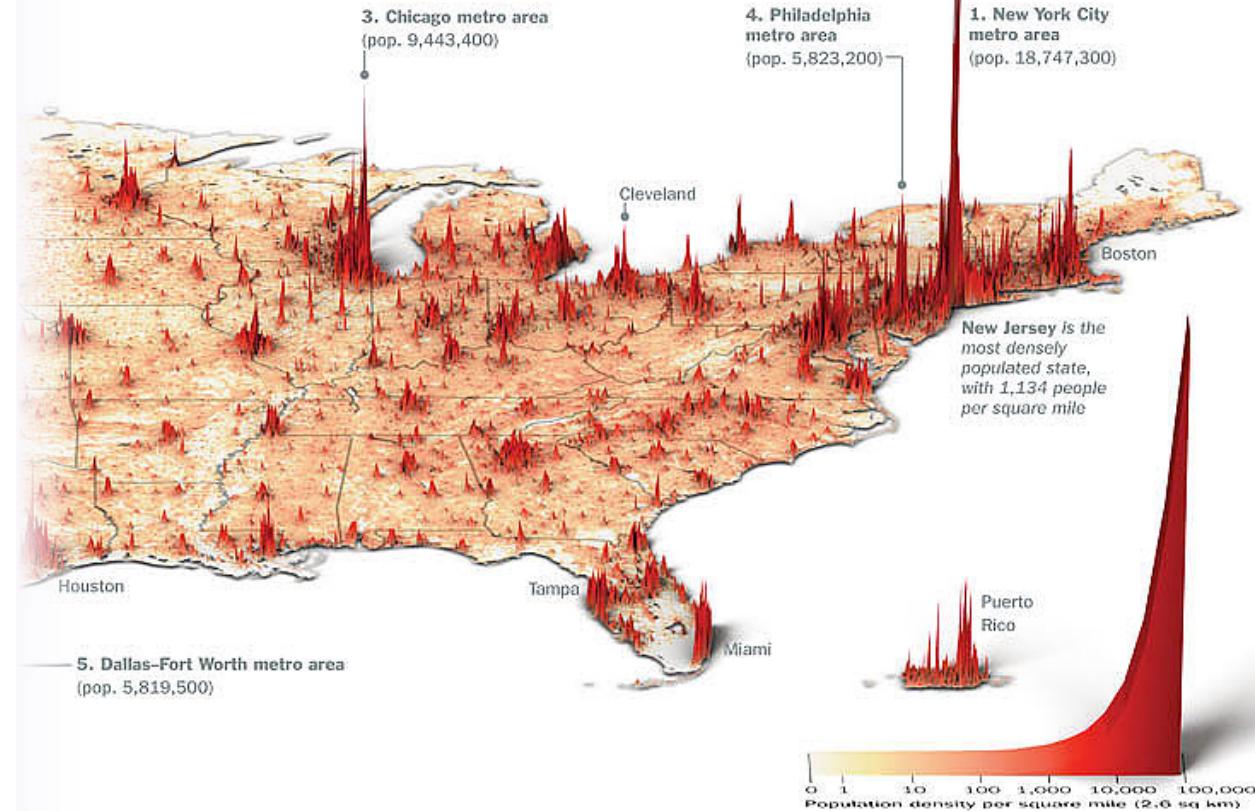
Ala.; Possum Trot, Ky.; or Lonelyville, N.Y. But they are all probably close to someone's idea of paradise. —By Nancy Gibbs

# What works?

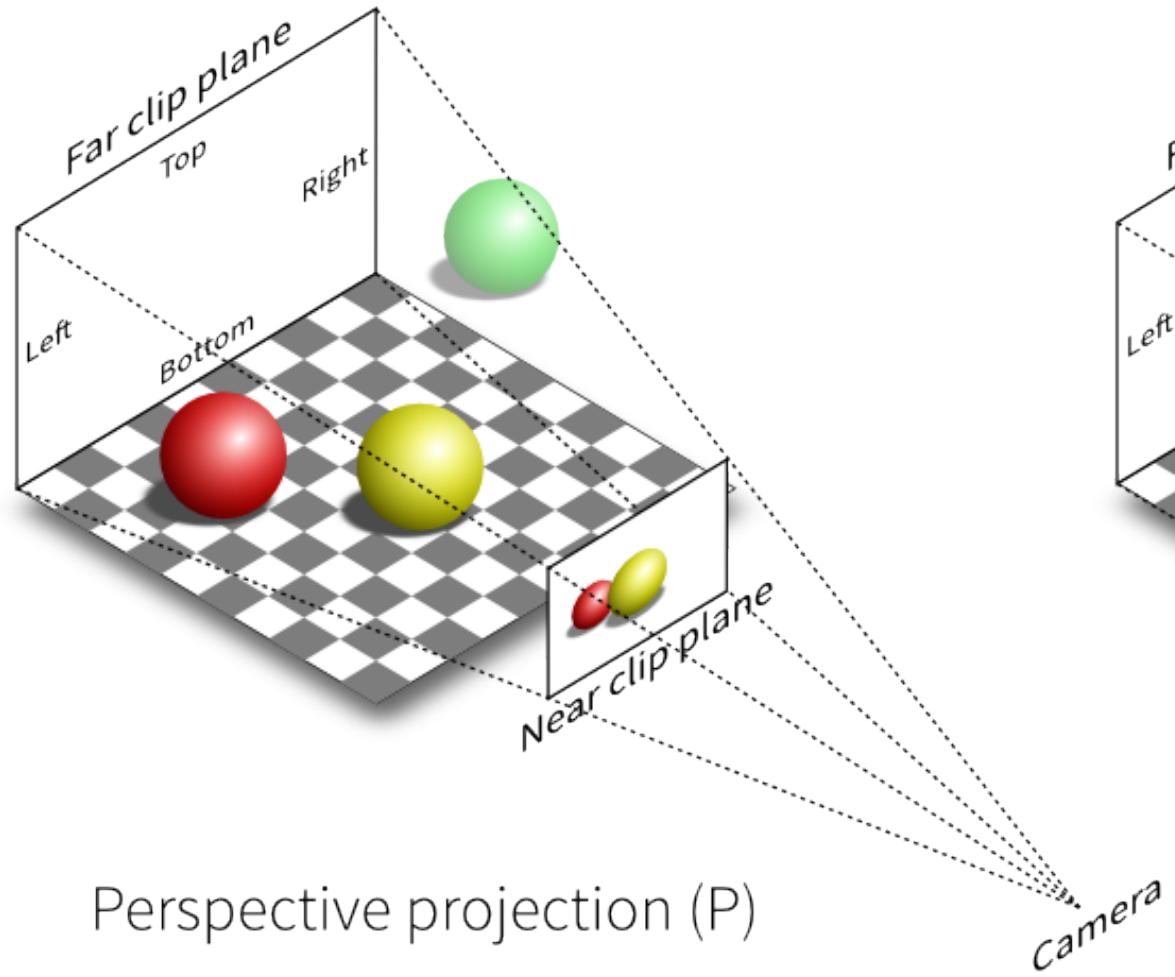
- Orthographic projection.
- Extensive labeling.
- Great rendering.
- Occlusion is not severe due to the usage of log scale.

Our families are getting smaller—with one vital exception. Compared with those of Europe and Japan, the U.S. population is younger and more colorful because of the continued arrival of immigrants and their higher-than-average birthrates. Of the 100 million Americans who will join us in the next 37 years, half will be immigrants or their children. In the next few decades, 97% of the world's population growth will occur in the developing world; the U.S. is the largest developed country in the world that is still growing at a healthy clip. That matters, strategically, economical-

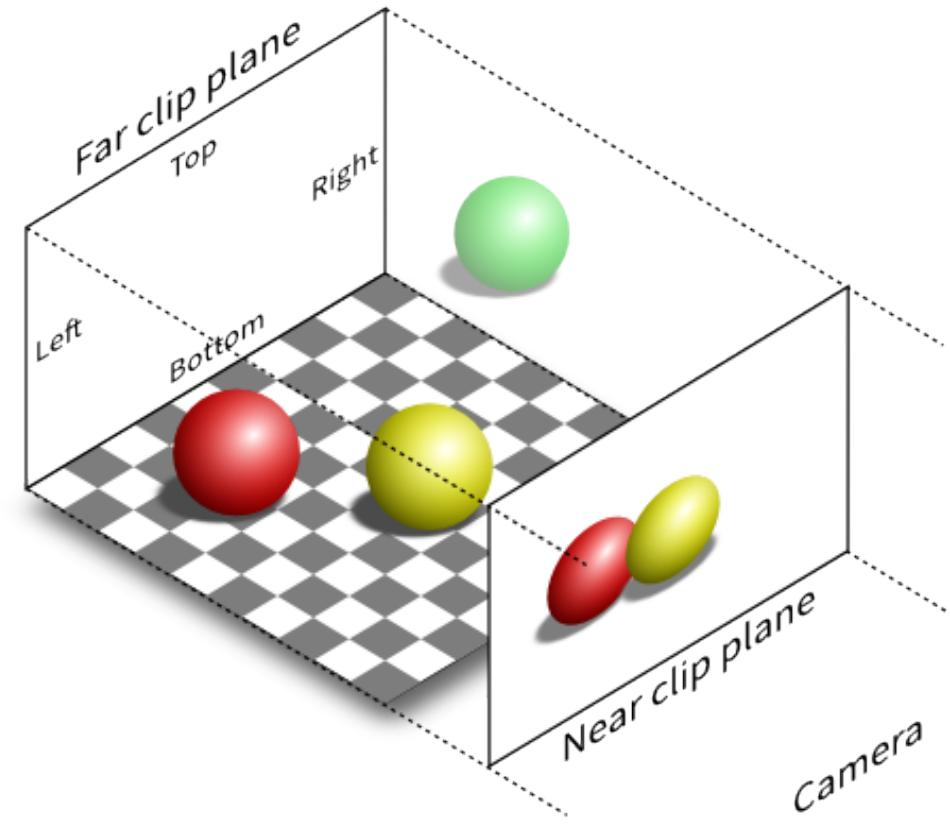
Ala.; Possum Trot, Ky.; or Lonelyville, N.Y. But they are all probably close to someone's idea of paradise. —By Nancy Gibbs



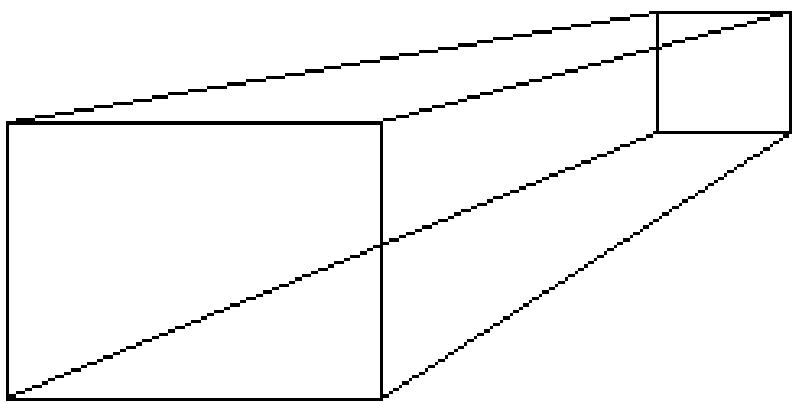
Source: [https://johnstonarchitects.files.wordpress.com/2012/04/pop\\_lg.jpg](https://johnstonarchitects.files.wordpress.com/2012/04/pop_lg.jpg)



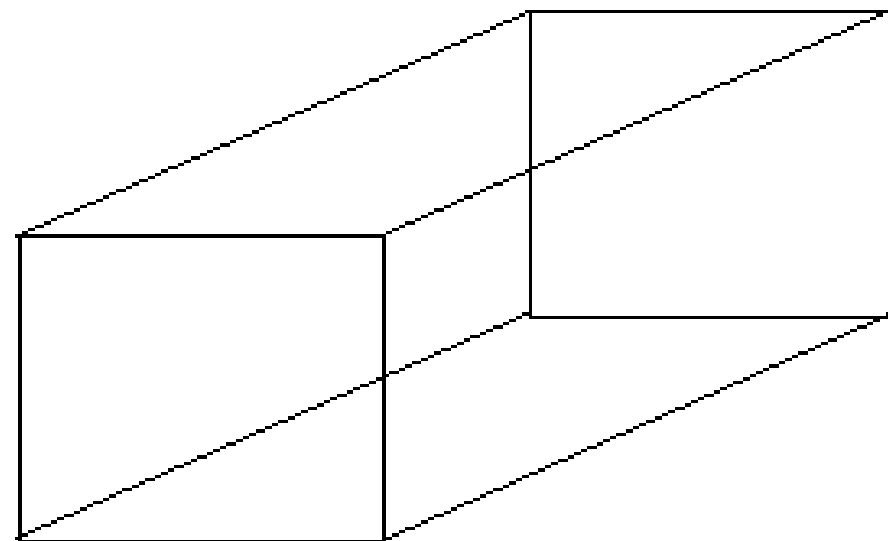
Perspective projection (P)



Orthographic projection (O)



Perspective projection



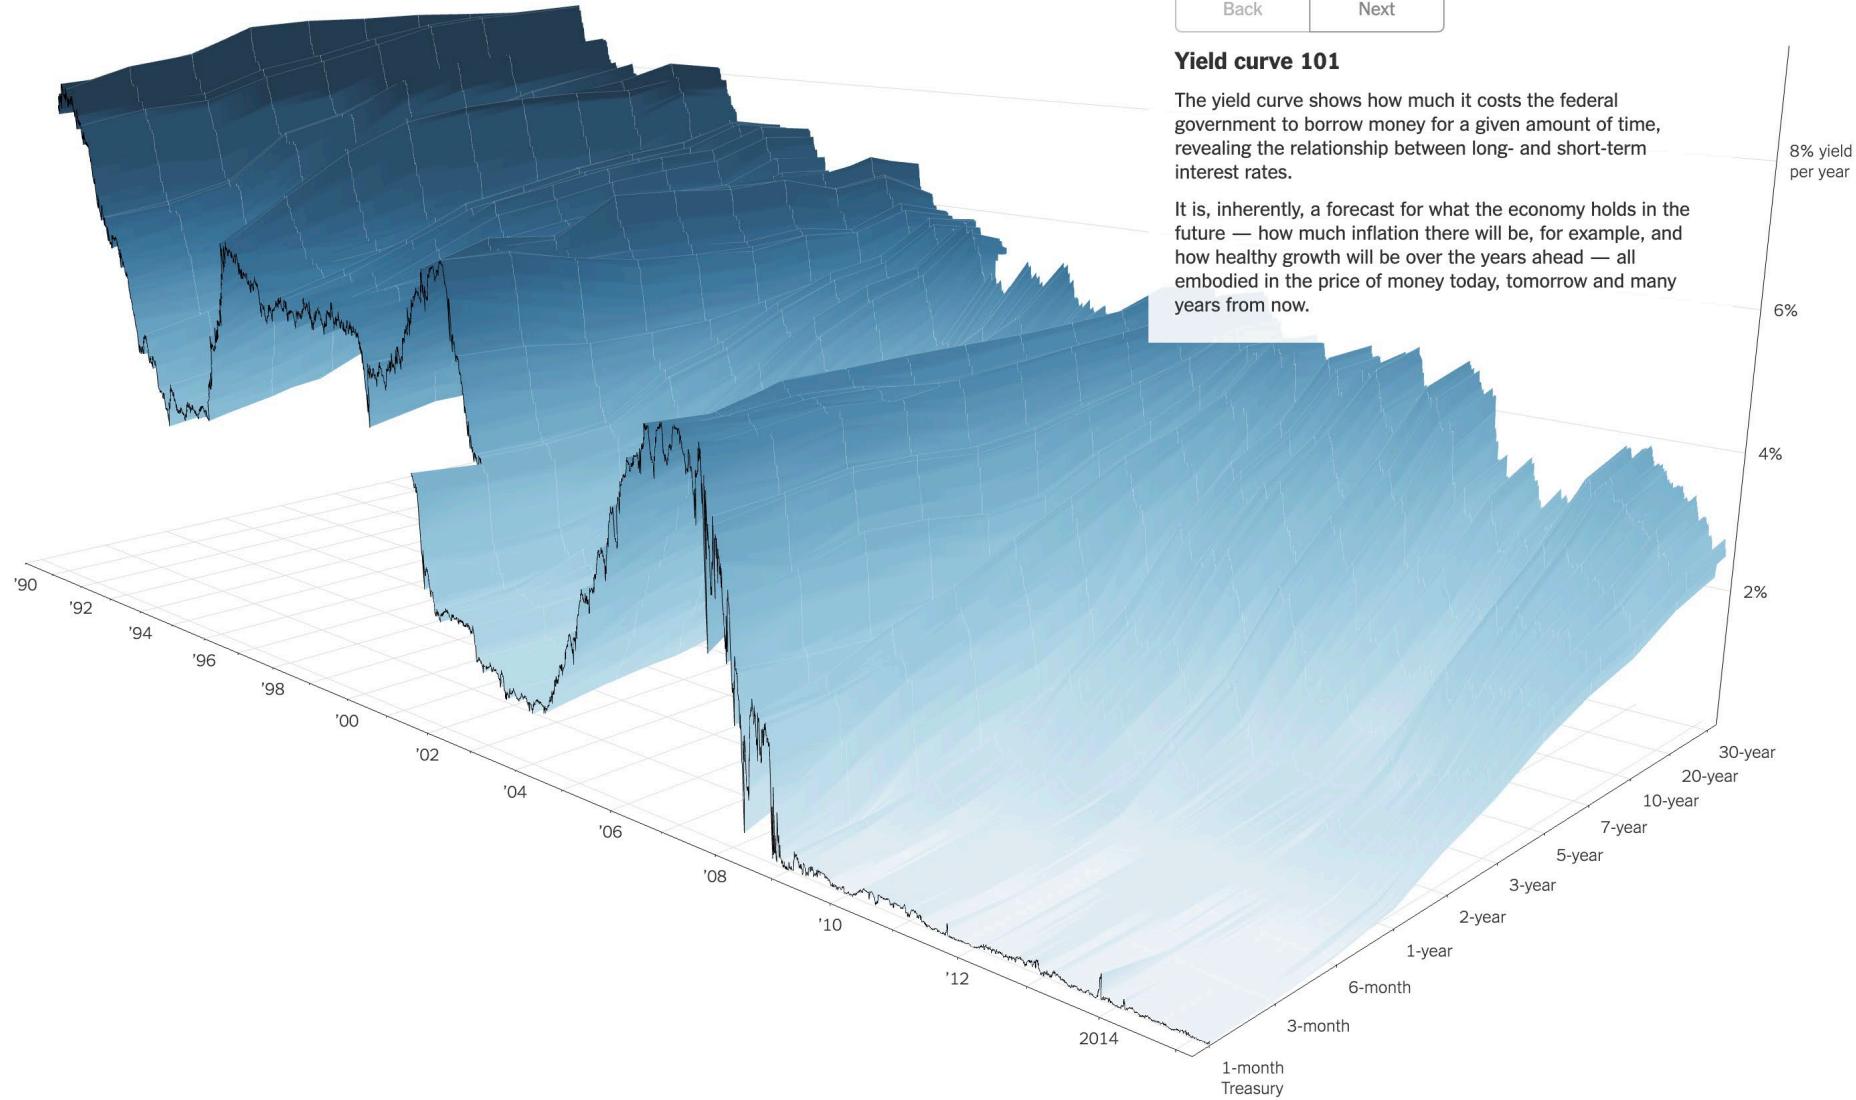
Orthographic projection

[Back](#)[Next](#)

### Yield curve 101

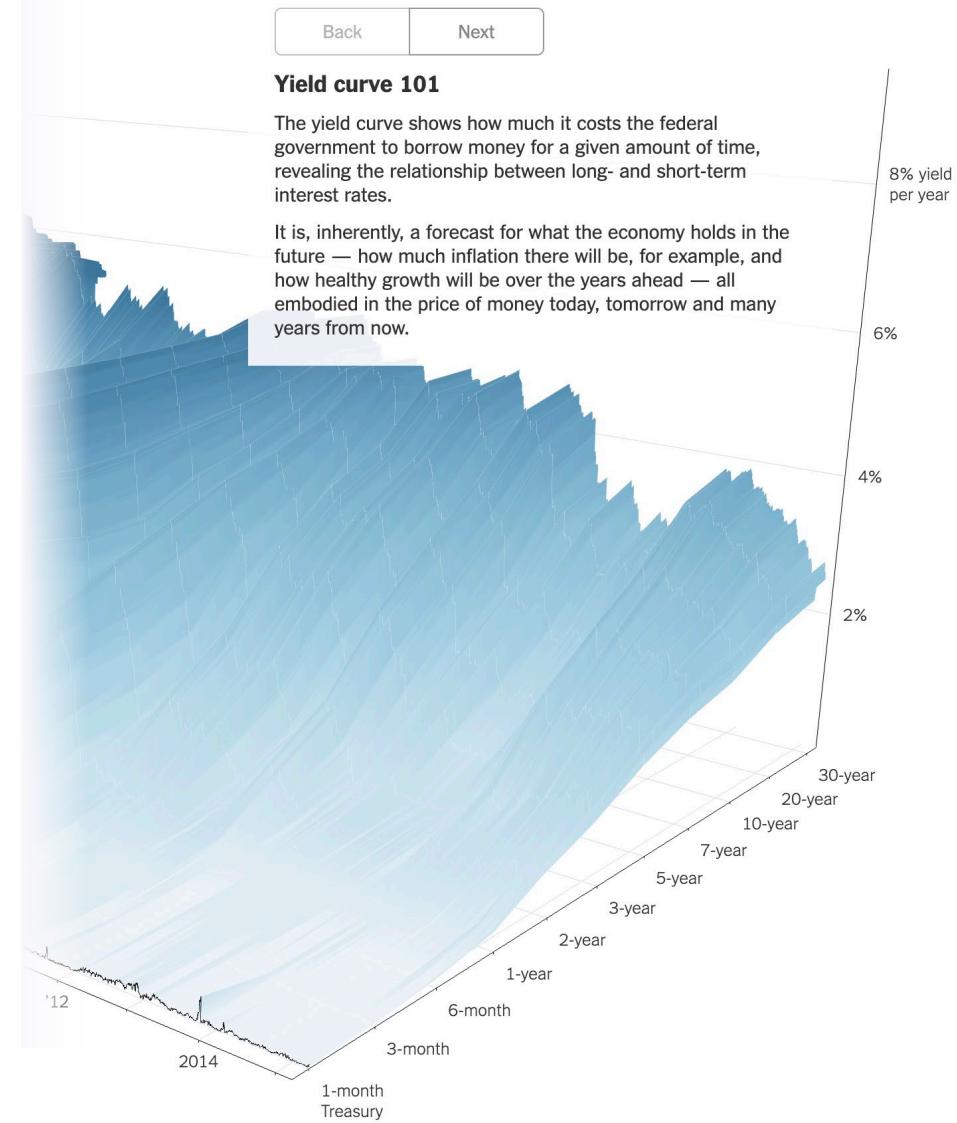
The yield curve shows how much it costs the federal government to borrow money for a given amount of time, revealing the relationship between long- and short-term interest rates.

It is, inherently, a forecast for what the economy holds in the future — how much inflation there will be, for example, and how healthy growth will be over the years ahead — all embodied in the price of money today, tomorrow and many years from now.



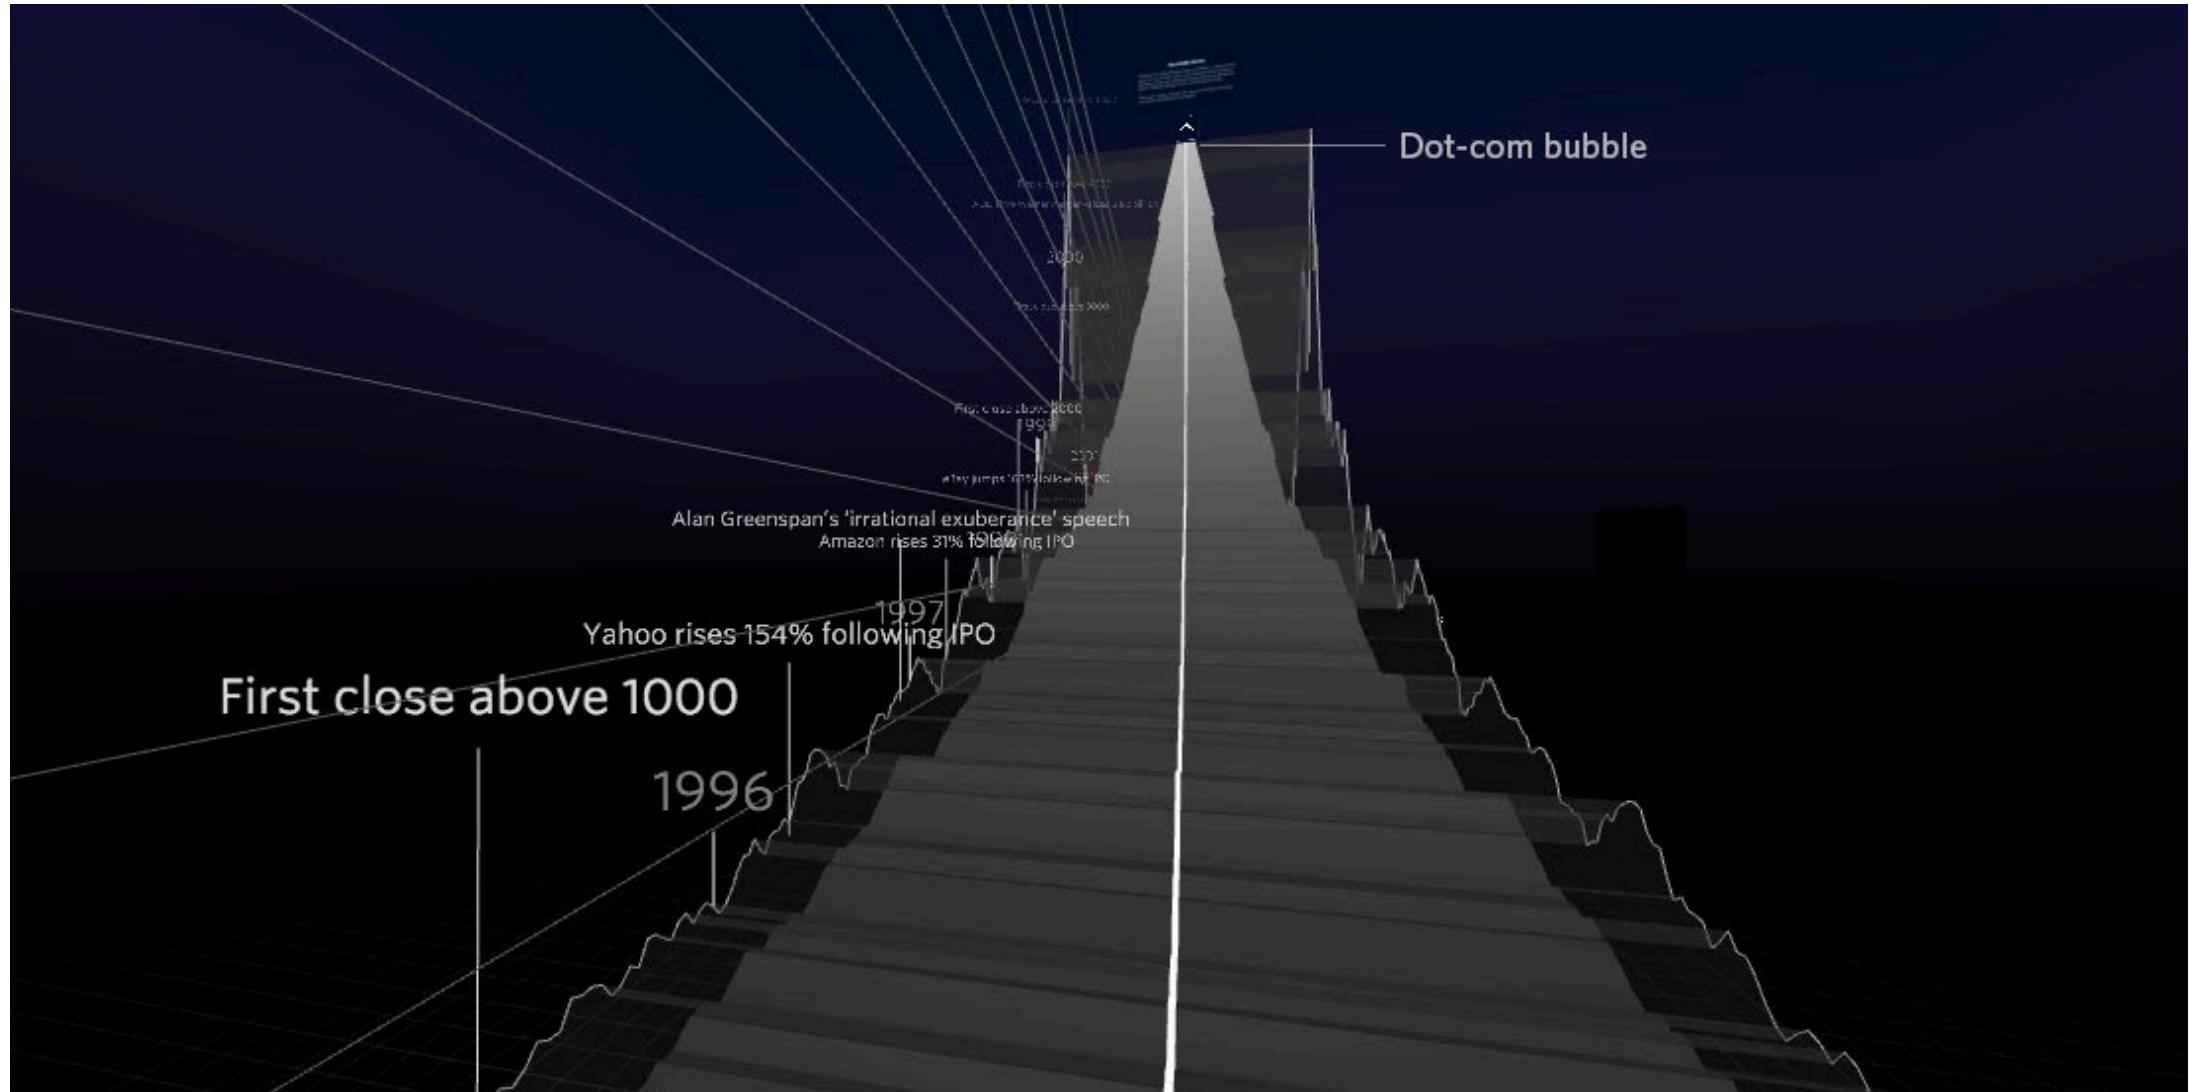
# What Works?

- Interactivity.
- Multiple selected views.
- In-chart explanation.
- Grid lines.



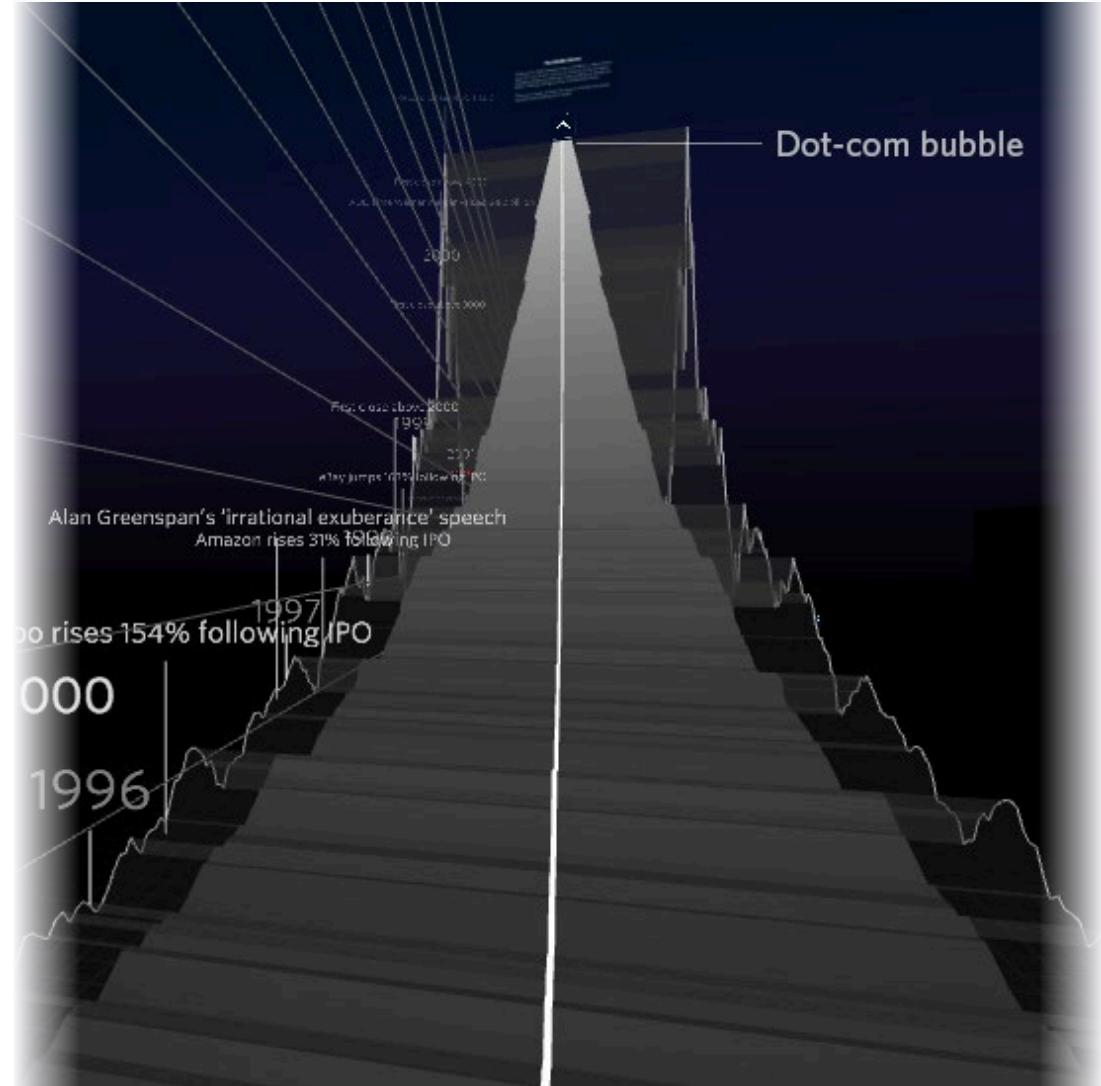
128

Source: <https://www.nytimes.com/interactive/2015/03/19/upshot/3d-yield-curve-economic-growth.html>

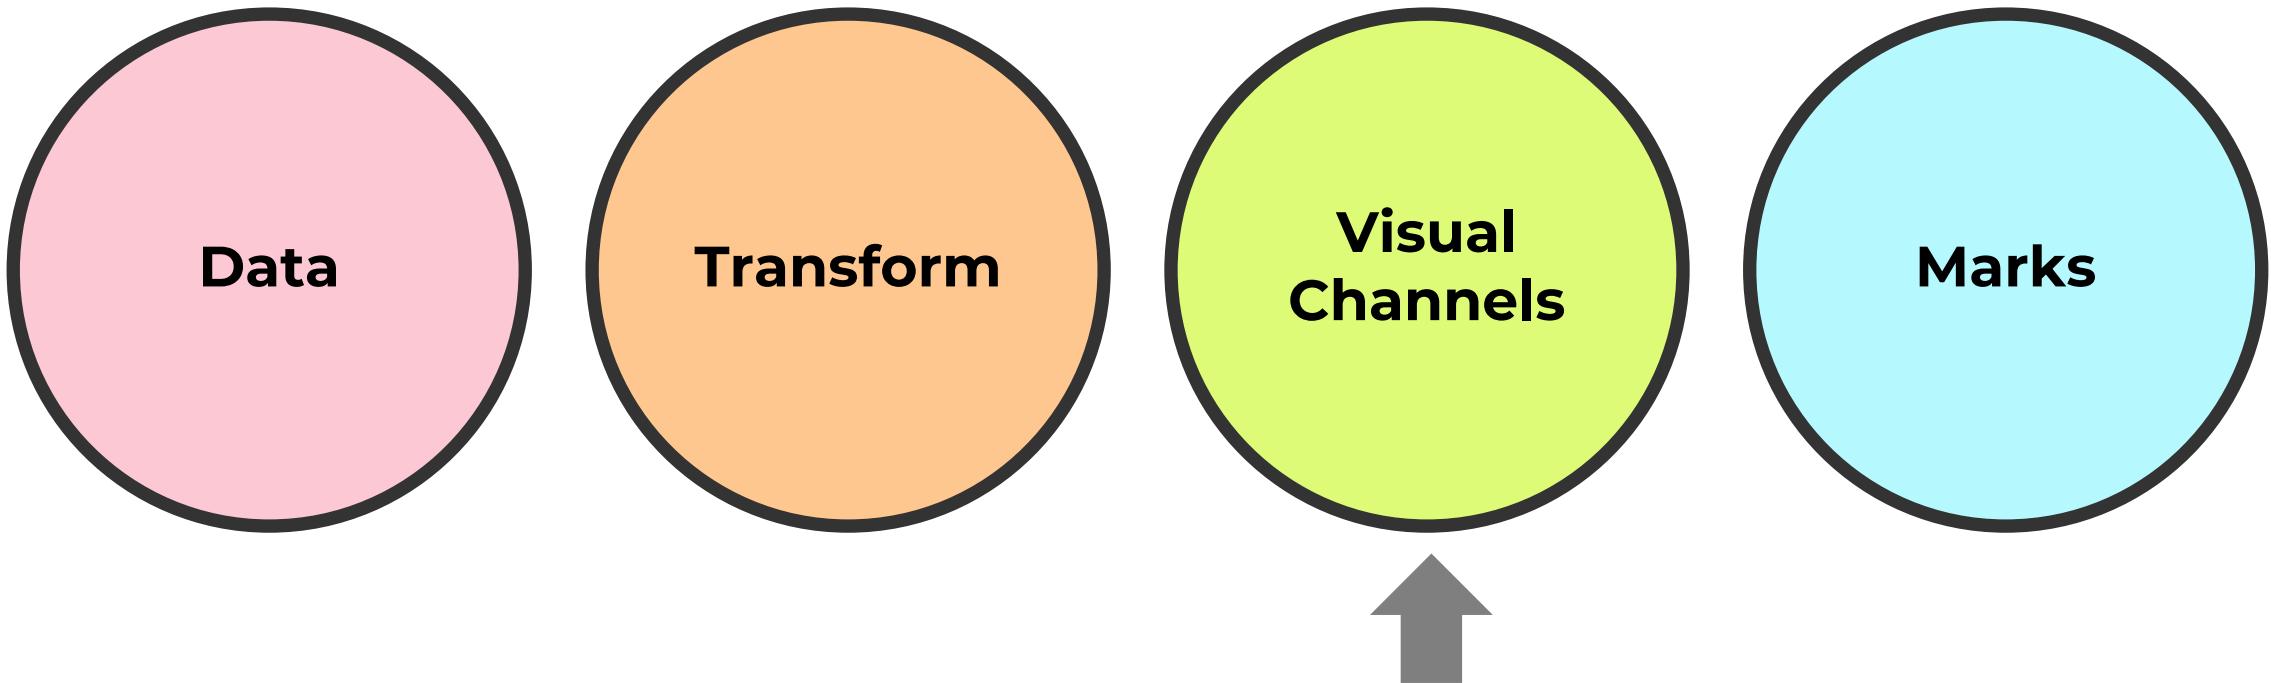


# What Works?

- Novel viewing medium (VR).
- Draw from familiar experience.
- Well annotated.



# 3D Visualization Grammar



# Visual Channels

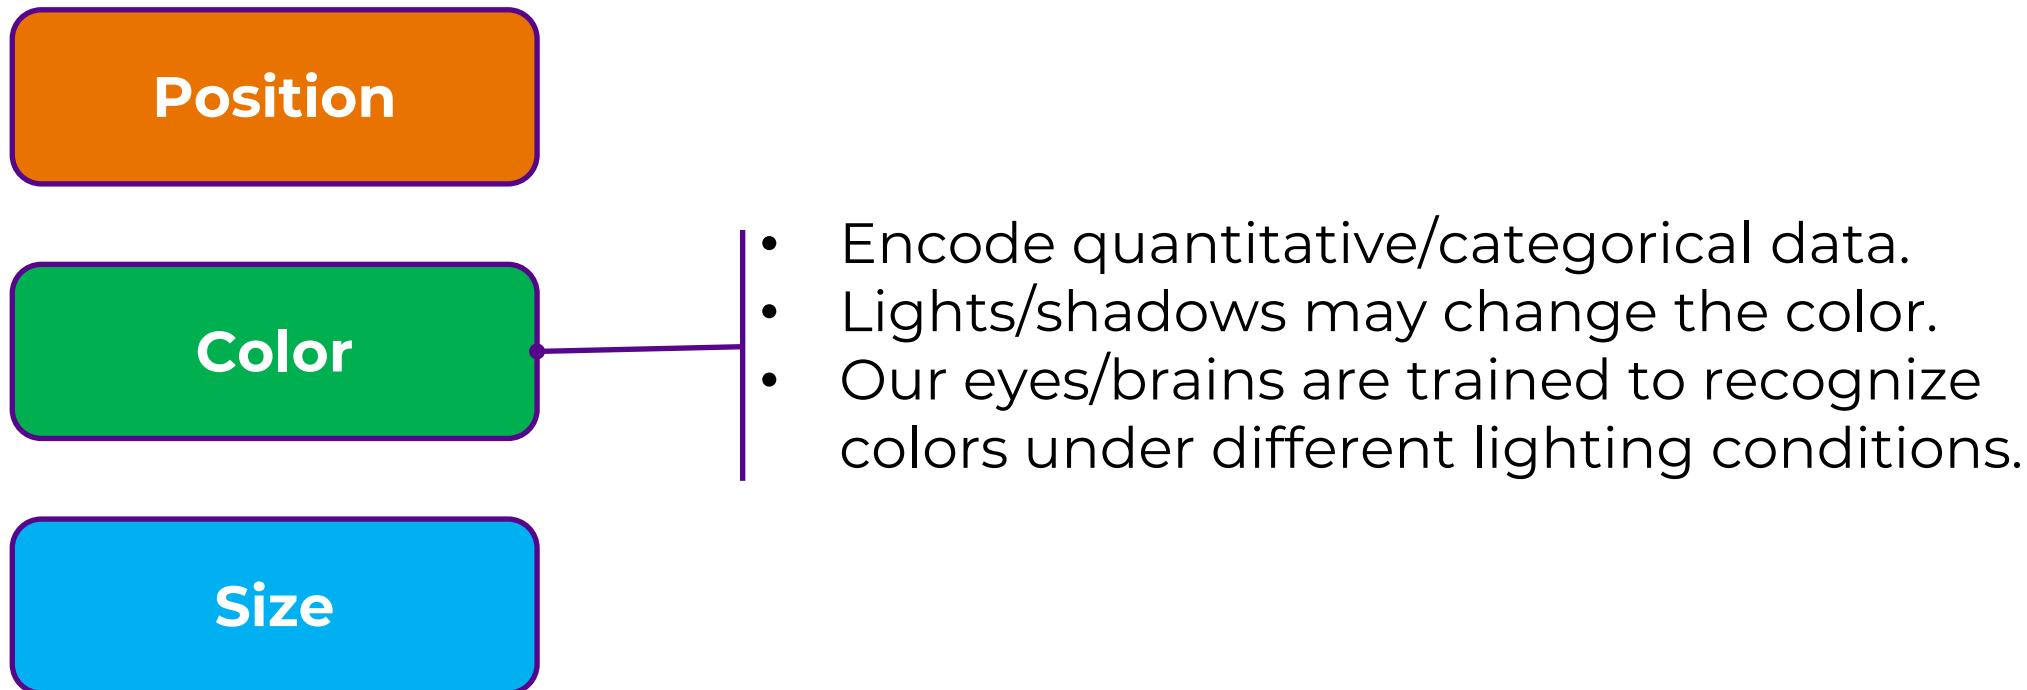
**Position**

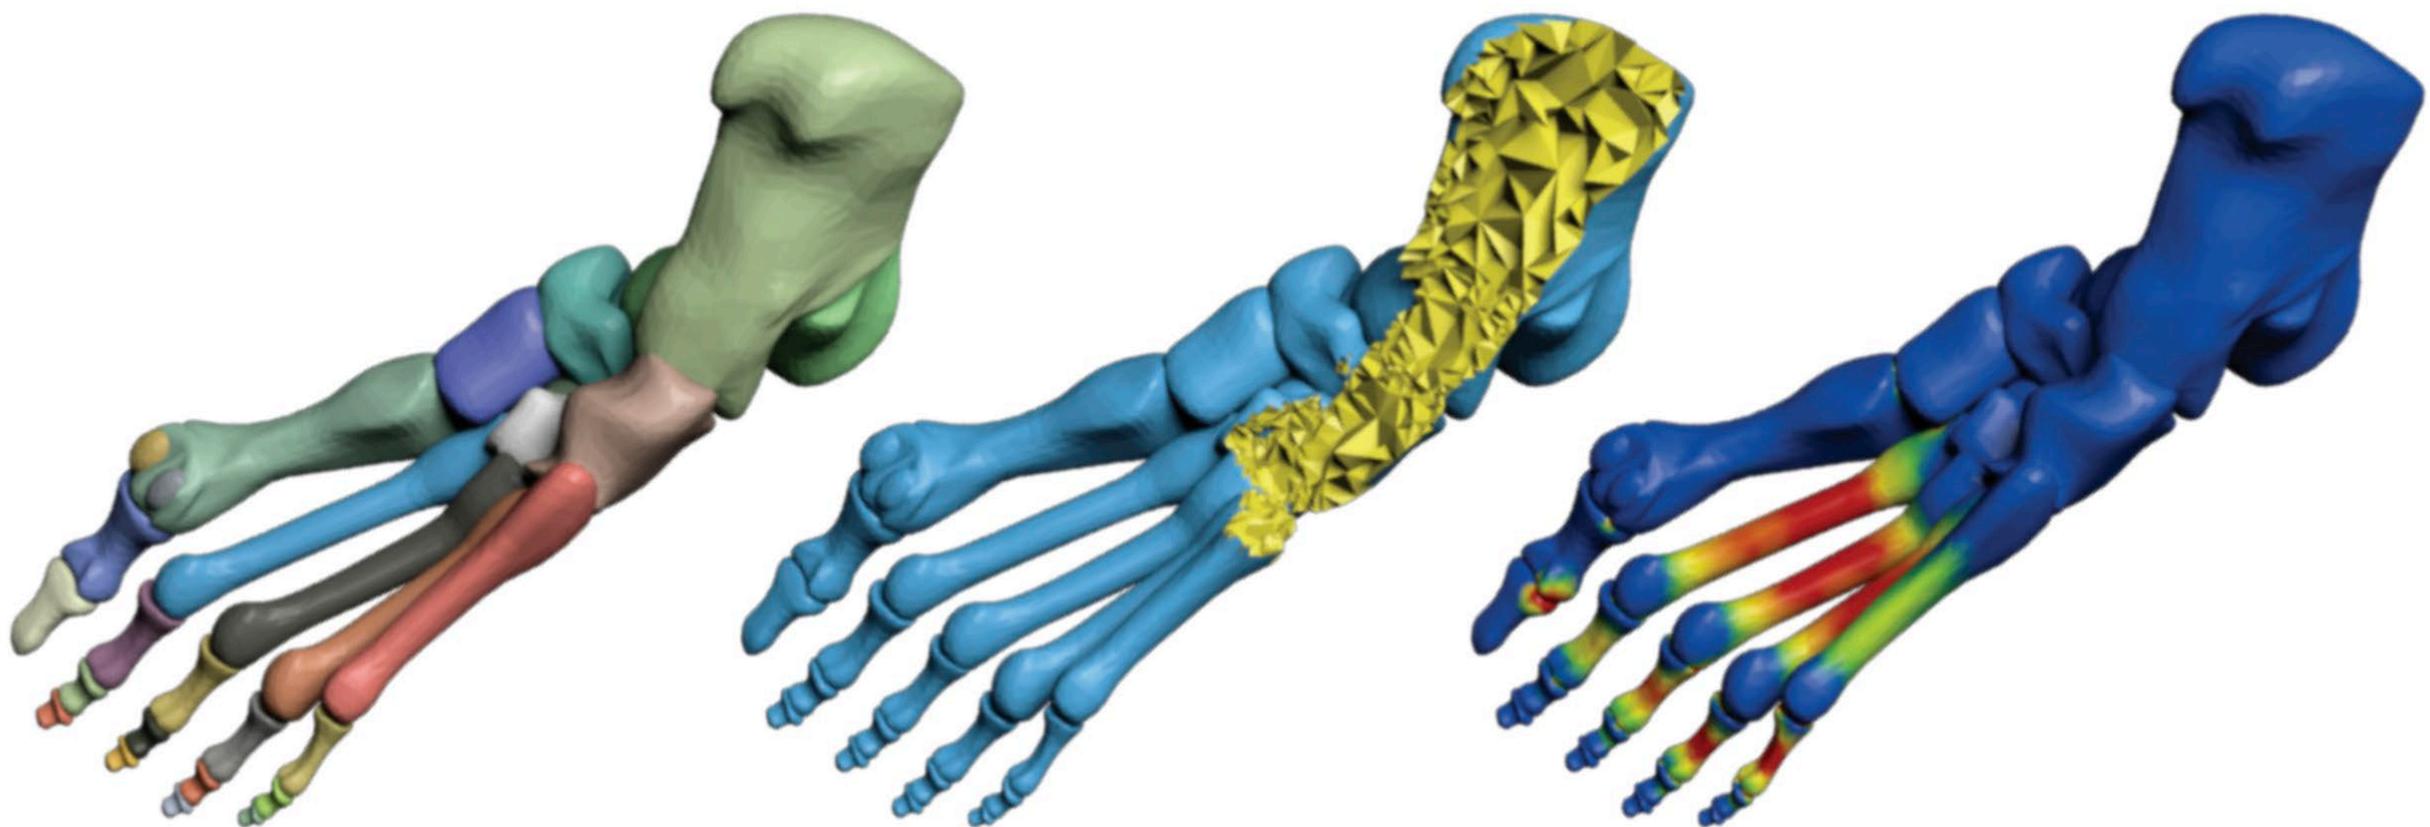
Used for encoding the geometry.

**Color**

**Size**

# Visual Channels





# Visual Channels

**Position**

**Color**

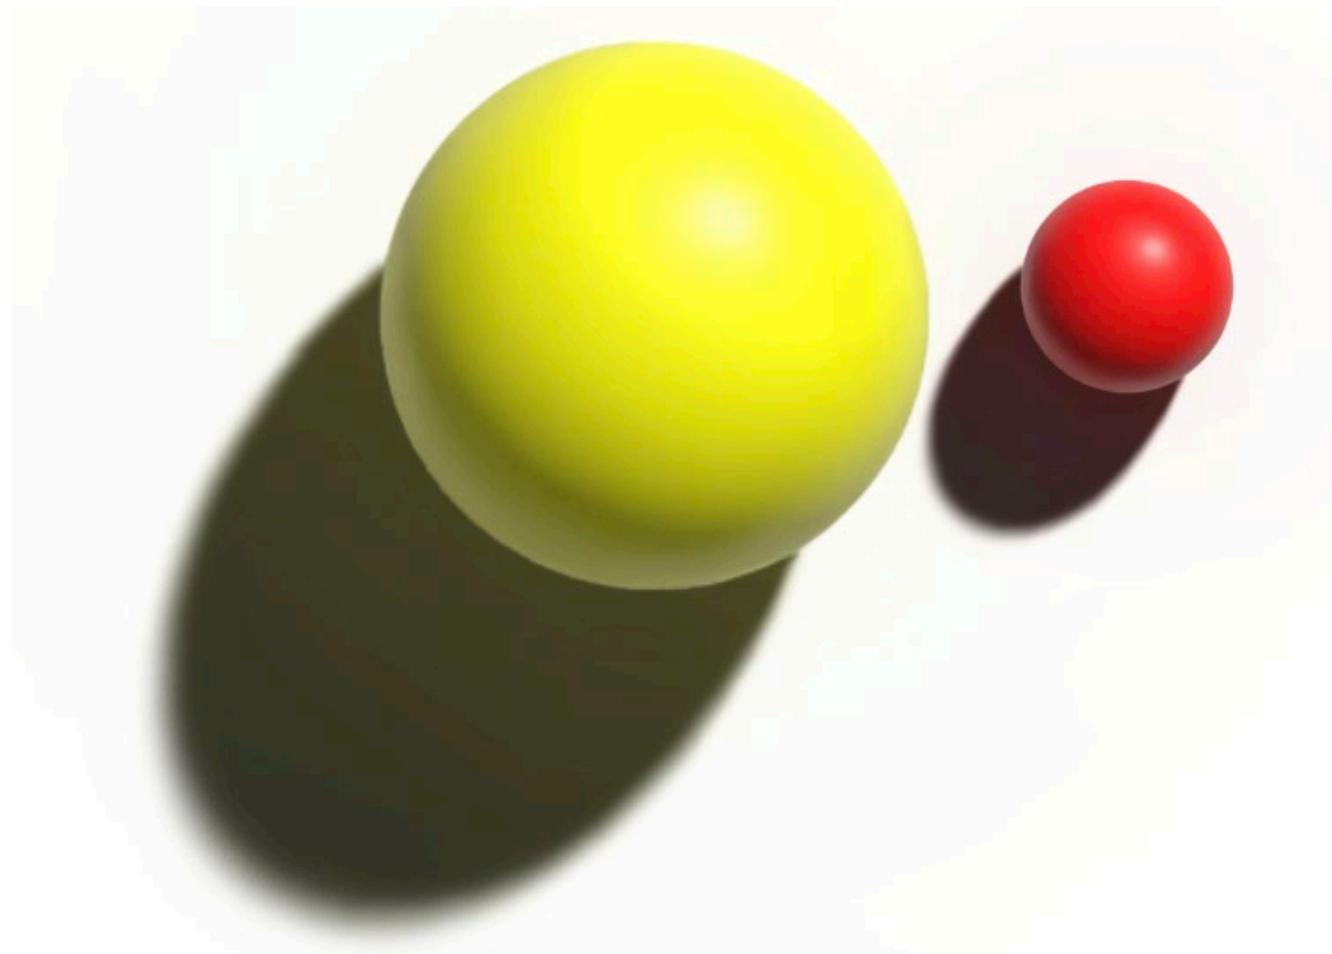
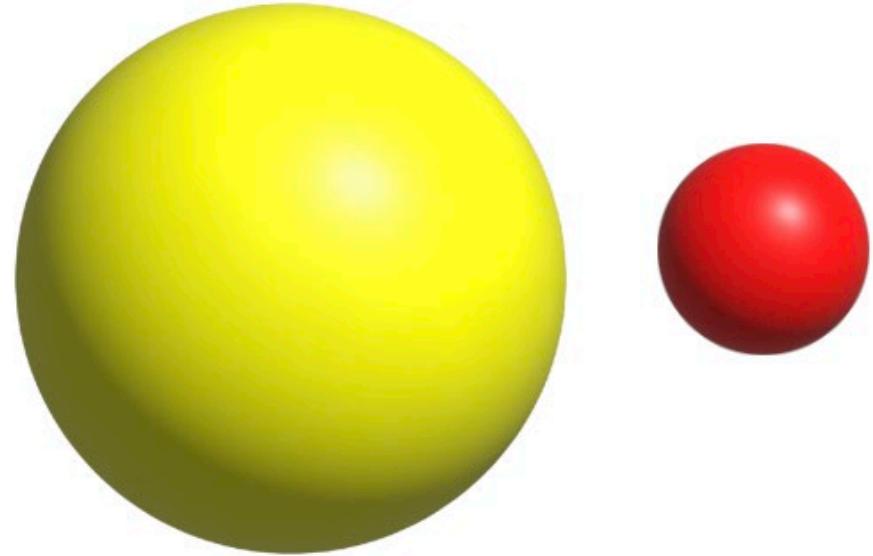
**Size**

Not used as often because foreshortening makes it difficult to distinguish different sizes.

Is the red ball **smaller** than the yellow ball?

*Or...*

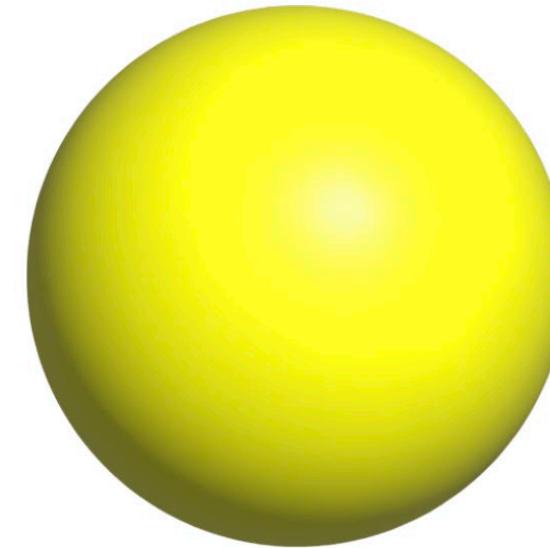
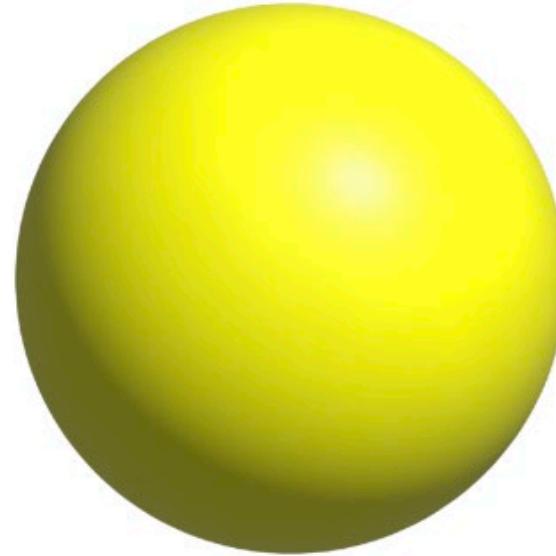
Is the red ball **further away** than the yellow ball?



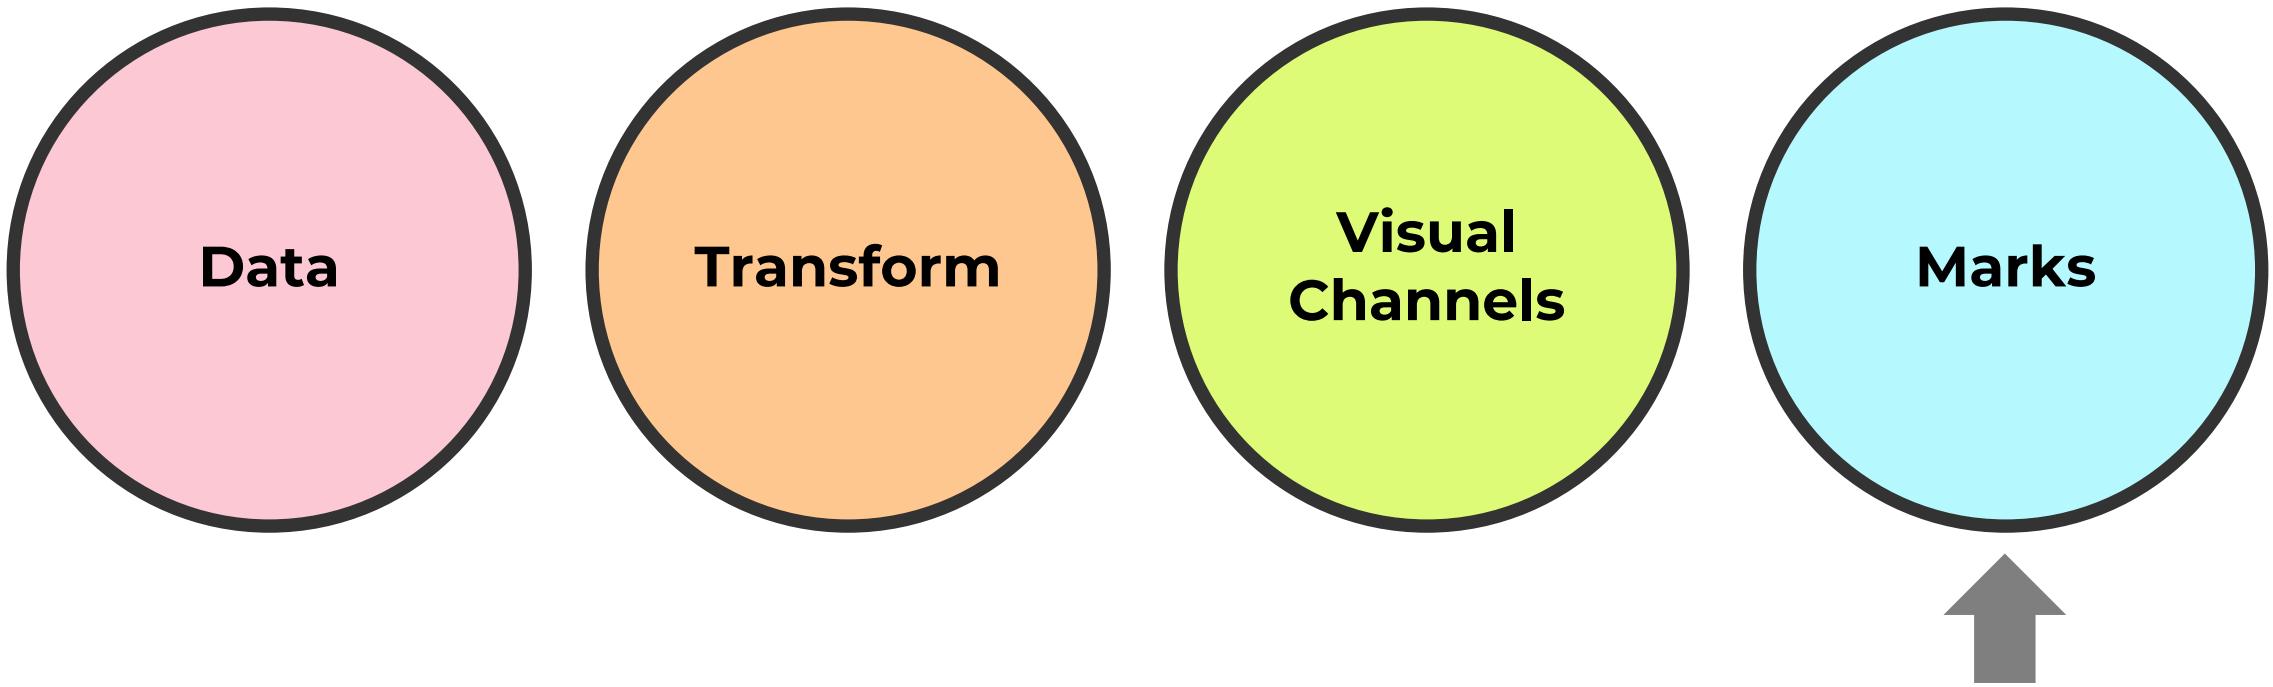
Is the red ball **smaller** than the yellow ball?

*Or...*

Is the red ball **further away** than the yellow ball?



# 3D Visualization Grammar



# Marks

- Points
- Curves/lines
- Triangles/quads
- Tetrahedra/Hexahedra