



# Support Vector Machines

Week 4

February 12, 2024

Spring 24 | CUSP-GX 7033 – Machine Learning for Cities | Dr. Anton Rozhkov

# Today's Outline

- From Linear to Non-Linear Classifiers with Support Vector Machines
- Linear decision boundaries
- Support vector machines (SVMs) for classification
- Moving from linear to non-linear decision boundaries with kernel SVM
- Lab: SVM in Python

# Support Vector Machines

# From Linear to Non-Linear Classifiers

Pre 1980: Almost all learning methods learned linear decision surfaces.

Linear learning methods have nice theoretical properties

1980's: Decision trees and NNs allowed efficient learning of non-linear decision surfaces.

Little theoretical basis, and all suffer from local minima

Starting from 1990's: Efficient learning algorithms for non-linear functions based on computational learning theory developed.

Nice theoretical properties.

---

1980

---

1990

---

# Support Vector Machines (SVMs)

**Support vector machines** are an optimization-based prediction approach used primarily for binary classification and are able to achieve state-of-the-art prediction accuracy on many real-world tasks.

Key idea 1: Learn a **decision boundary** that optimally separates positive and negative training examples. (But what does it mean to be optimal?)

Key idea 2: Learn a **linear** decision boundary in high dimensional space corresponding to a **non-linear** decision boundary for the original problem.


SVM assumes **real-valued** attributes on the **same scale**. Thus, it is very important to pre-process your data before training the model:

- Normalize real-valued attributes (scale either to  $[0,1]$  or to mean = 0 and variance = 1). Make sure to use the same scaling for training and test data.
- Replace discrete-valued attributes with **dummy variables**.




<u>Car</u>	<u>Weight</u>		<u>Car</u>	<u>Weight=Medium</u>	<u>Weight=Heavy</u>
1	Low		1	0	0
2	Medium	➔	2	1	0
3	Heavy		3	0	1

# SVMs: Basic Idea (Linearly Separable Case)




# SVMs: Basic Idea (Linearly Separable Case)



In general,  
there are many  
possible  
solutions (an  
infinite  
number!)

SVM finds an  
optimal  
solution

# SVMs: Basic Idea (Linearly Separable Case)




In general,  
there are many  
possible  
solutions (an  
infinite  
number!)


SVM finds an  
optimal  
solution




# SVMs: Basic Idea (Linearly Separable Case)



# SVMs: Basic Idea (Linearly Separable Case)

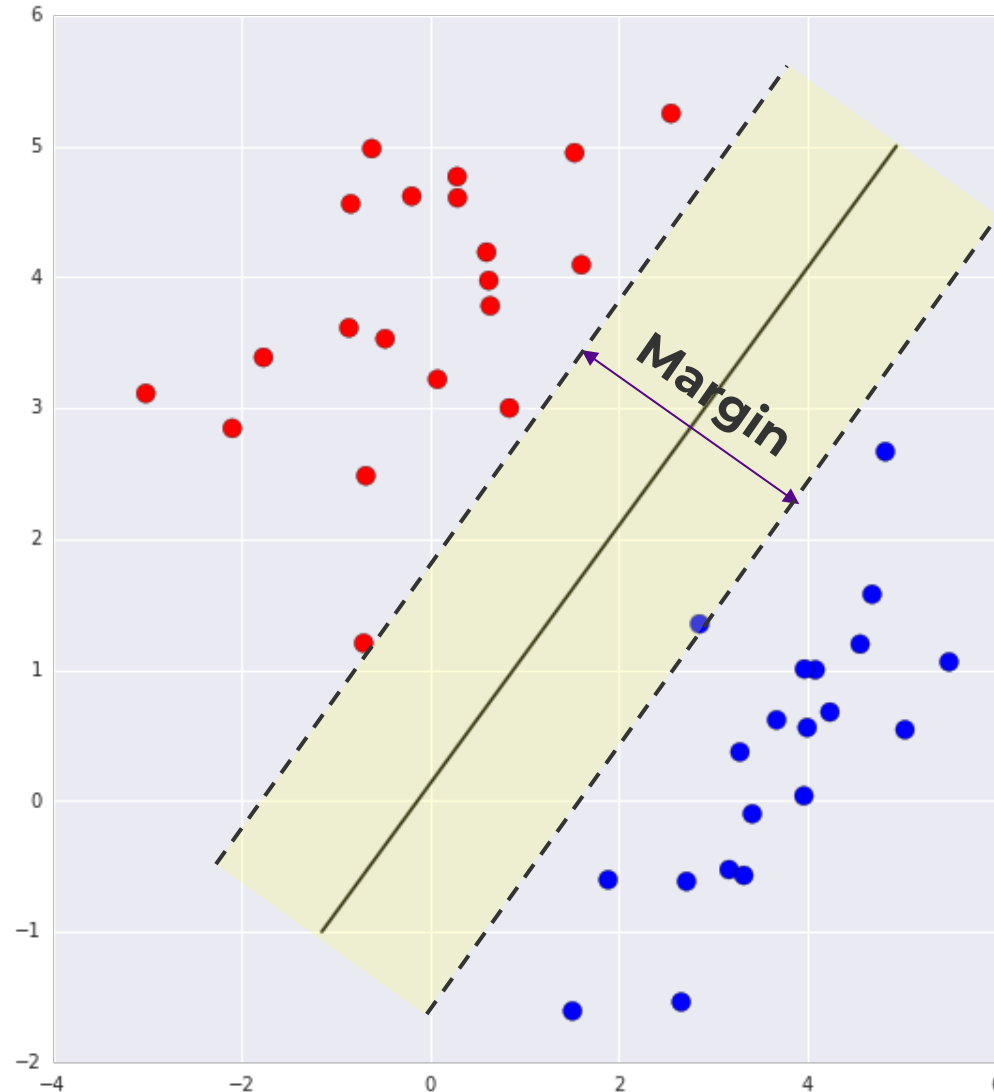


# Which Line to Choose?



This is an optimization problem

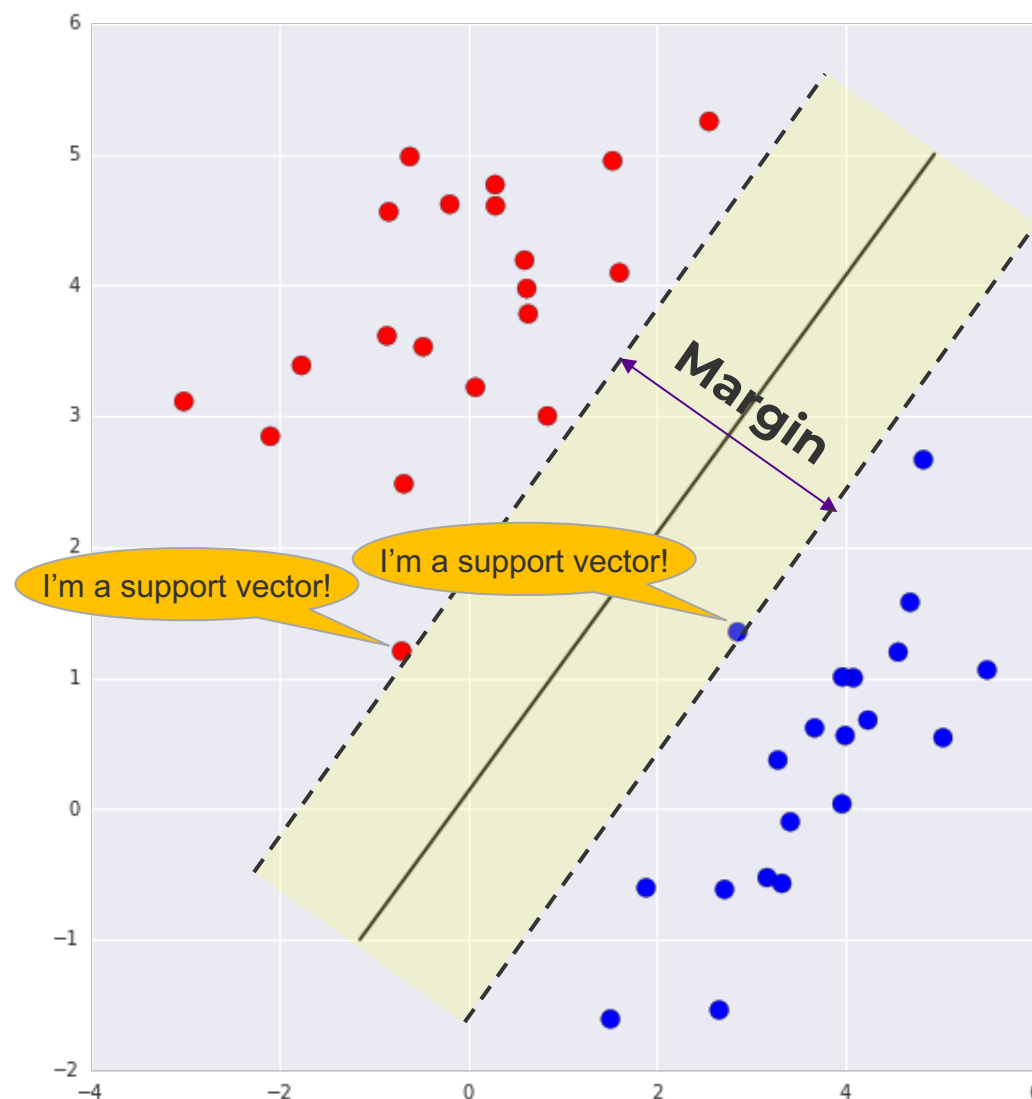
# Which Line to Choose?



Choose the line that **maximizes** the margin between classes.

Margin = how wide we could make the linear decision boundary before it contacts points from either class.

# Which Line to Choose?



Points on the margin are called **support vectors**.

The classifier can be defined entirely by the set of support vectors.


This fact has lots of useful implications:

- Fast classification of test points.
- Fast leave-one-out cross-validation.
  - Faster, but still expensive, training.

# Which Line to Choose?

Why is maximizing the margin a good idea?

1. Intuitively, this feels safest.
2. If we've made a small error in the location of the boundary, this gives us the least chance of causing a misclassification.
3. Backed up by statistical learning theory  $\rightarrow$  there are provable bounds on generalization error.
4. Empirically, it works very well.



# How to Maximize the Margin?

To separate, for all points  $j$ , we must have:

$$y_j(x_j^T w + b) > 0$$

For the margin, define:


$$M = \min_j y_j(x_j^T w + b) > 0$$

Then for  $y_j = 1$ , we have:

$$x_j^T w + b \geq M$$

For  $y_j = -1$ , we have:

$$x_j^T w + b \leq -M$$



Decision boundary (separating line or hyperplane)

Represent each point as  $(x_j, y_j)$ , where  $y_j$ , the class value we are trying to predict, is +1 or -1. Note that  $x_j$  is a vector of length 2 in this example.

# How to Maximize the Margin?

To separate, for all points  $j$ , we must have:

$$y_j(x_j^T w + b) > 0$$

For the margin, define:


$$M = \min_j y_j(x_j^T w + b) > 0$$

Then for  $y_j = 1$ , we have:

$$x_j^T w + b \geq M$$

For  $y_j = -1$ , we have:

$$x_j^T w + b \leq -M$$



Decision boundary (separating line or hyperplane)

Represent each point as  $(x_j, y_j)$ , where  $y_j$ , the class value we are trying to predict, is +1 or -1. Note that  $x_j$  is a vector of length 2 in this example.



# How to Maximize the Margin?

$$\text{Margin} = 2M / \|w\|.$$

This follows from computing the distance between parallel lines.

Goal: maximize  $2M / \|w\|$  subject to constraints, for all  $j$ :

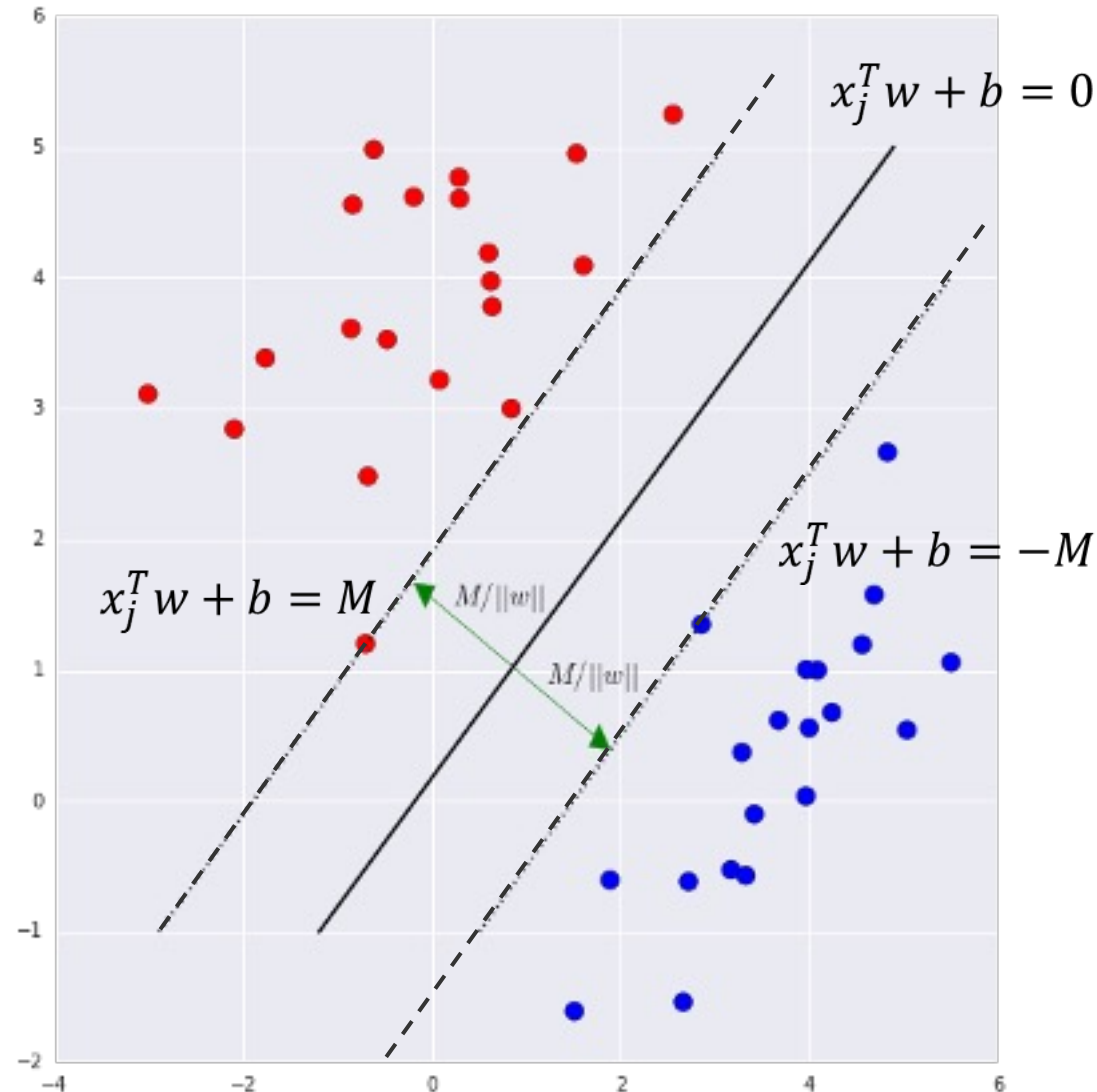
$$y_j(x_j^T w + b) \geq M$$

Simplify by change of variables, dividing  $w$  and  $b$  through by  $M$ .

New goal: minimize  $\|w\|$  subject to constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1$$

New margin:  $2 / \|w\|$



Decision boundary (separating line or hyperplane)

# How to Maximize the Margin?

$$\text{Margin} = 2M / \|w\|.$$

This follows from computing the distance between parallel lines.

Goal: maximize  $2M / \|w\|$  subject to constraints, for all  $j$ :


$$y_j(x_j^T w + b) \geq M$$

Simplify by change of variables, dividing  $w$  and  $b$  through by  $M$ .

New goal: minimize  $\|w\|$  subject to constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1$$

New margin:  $2 / \|w\|$



# How to Maximize the Margin?

$$\text{Margin} = 2M / \|w\|.$$

This follows from computing the distance between parallel lines.

Goal: maximize  $2M / \|w\|$  subject to constraints, for all  $j$ :


$$y_j(x_j^T w + b) \geq M$$

Simplify by change of variables, dividing  $w$  and  $b$  through by  $M$ .

New goal: minimize  $\|w\|$  subject to constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1$$

New margin:  $2 / \|w\|$



Decision boundary (separating line or hyperplane)

Question:  
What if the points are not linearly separable? Then there is no solution satisfying the constraints!

# Non-separable Case: Soft Margins


Goal (hard margin):  
minimize  $\|w\|$  subject to  
constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1$$



Goal (soft margin):  
minimize subject to  
constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1 - \xi_j$$
$$\xi_j \geq 0$$



Decision boundary (separating  
line or hyperplane)

# Non-separable Case: Soft Margins

Goal (hard margin):  
minimize  $\|w\|$  subject to  
constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1$$




Goal (soft margin):  
minimize subject to  
constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1 - \xi_j$$
$$\xi_j \geq 0$$

But what should we  
minimize?  $\|w\|$ ?

Answer: minimize

$$\frac{1}{2} \|w\|^2 + C \sum_j \xi_j$$



Decision boundary (separating  
line or hyperplane)

This is a quadratic  
programming (QP)  
problem, optimizing a  
quadratic function with  
linear constraints.

We can use off-the-  
shelf QP solvers or  
faster methods for large  
problems to find the  
optimal  $w$ ,  $b$ , and  $\xi$ .

Classify test points  
by  $\text{sign}(x_j^T w + b)$ .

# Soft Margins in Practice

**Goal (hard margin):**  
minimize  $\|w\|$  subject to  
constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1$$



**Goal (soft margin):**  
minimize subject to  
constraints, for all  $j$ :

$$y_j(x_j^T w + b) \geq 1 - \xi_j$$
$$\xi_j \geq 0$$


But what should we  
minimize?  $\|w\|$ ?

Answer: minimize


$$\frac{1}{2} \|w\|^2 + C \sum_j \xi_j$$

In practice, there may be many training points  
with  $\xi_j > 0$  (all of these are support vectors).

Training points with  $\xi_j > 1$  are misclassifications.




# Non-linear Decision Boundaries



Question:  
What do we do in cases like  
this one?

Any linear separator will  
perform terribly!





## Solution:


- 1) Map input space to a high-dimensional feature space.
- 2) Learn a linear decision boundary (hyperplane) in the high-dimensional space.
- 3) Map back to lower-dimensional space, giving a non-linear boundary.




# Non-linear Decision Boundaries



$$\phi : (x, y) \rightarrow (x, y, r)$$



The resulting classifier perfectly separates the training data.





# Non-linear Decision Boundaries

Non-linear QP problem:  $\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + C \sum_j \xi_j$  subject to:  $y_j(w^T \Phi(x_j) + b) \geq 1 - \xi_j$   
 $\xi_j \geq 0$

Problem: not efficiently computable, since  $\Phi(x_j)$  may be high- or infinite-dimensional!

Solution: transform to equivalent (“dual”) QP problem:

$$\min_{\alpha} \frac{1}{2} \alpha^T Q \alpha - \sum_j \alpha_j \quad \text{subject to: } 0 \leq \alpha_j \leq C \quad \text{where: } Q_{ij} = y_i y_j (\Phi(x_i) \cdot \Phi(x_j))$$
$$\sum_j \alpha_j y_j = 0 \quad \quad \quad = y_i y_j K(x_i, x_j)$$

Very cool trick (the “kernel trick”): instead of mapping both  $x_i$  and  $x_j$  into a high-dimensional space and computing the dot product in that space, we can just compute a function  $K(x_i, x_j)$  of the original data points.

This makes the QP efficiently solvable.  
To classify a test point  $x$ , we just need to compute  $\text{sign}(\sum_j \alpha_j y_j K(x_j, x) + \rho)$ .

Sum is just over the support vectors; other points have  $\alpha_j = 0$ .

# Some Common Kernel Functions

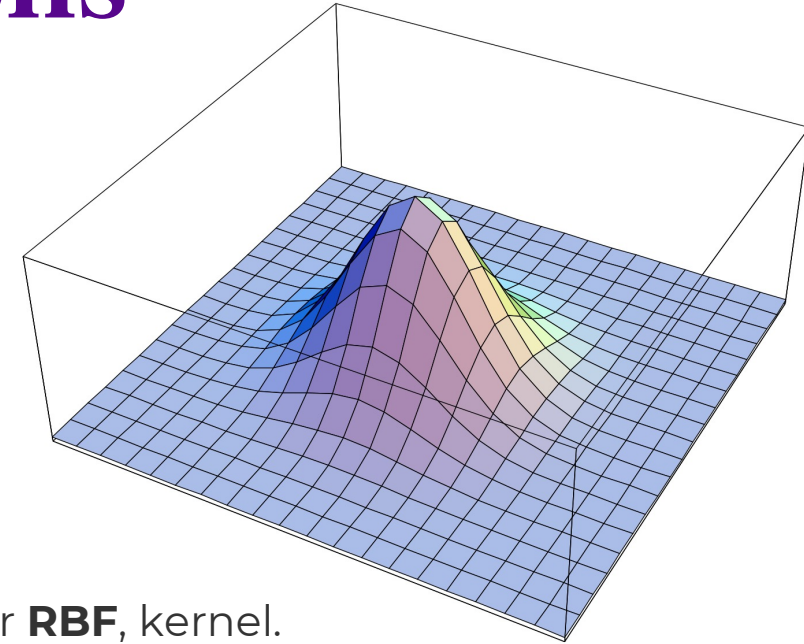
Non-linear  
kernels

Linear kernel:  $\phi : x \rightarrow x \quad K(x_i, x_j) = x_i \cdot x_j$

Polynomial kernel:  $K(x_i, x_j) = (\gamma(x_i \cdot x_j) + r)^d$

Sigmoid kernel:  $K(x_i, x_j) = \tanh(\gamma(x_i \cdot x_j) + r)$

Gaussian kernel:  $K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$



The Gaussian kernel is usually called the “radial basis function”, or **RBF**, kernel. It is one of the most widely used kernel choices and a good default option.

Very cool trick (the “kernel trick”): instead of mapping both  $x_i$  and  $x_j$  into a high-dimensional space and computing the dot product in that space, we can just compute a function  $K(x_i, x_j)$  of the original data points.

This makes the QP efficiently solvable. To classify a test point  $x$ , we just need to compute  $\text{sign}(\sum_j \alpha_j y_j K(x_j, x) + \rho)$ .

Sum is just over the support vectors; other points have  $\alpha_j = 0$ .

# Variants and Extensions of SVMs

SVMs are mainly used for **non-probabilistic, binary classification**.

To do multi-class classification:

For each class  $k$ , learn a binary classifier (class  $k$  vs. rest).

To predict the output for a new test example  $x$ , predict with each SVM.

Choose whichever one puts the prediction the furthest into the positive region.

To estimate class probabilities:

SVMs are not really the best for this, but they can do logistic regression using outputs of  $k(k-1)$  pairwise SVMs.

Lots of models + additional cross-validation are needed → this approach is very computationally expensive (Wu et al., 2004).

Support vector machines can also be used for **regression** (Smola and Schölkopf, 2003) and for **anomaly detection** (the “one-class SVM,” Schölkopf et al., 2001).

Both are implemented in scikit-learn but are beyond the scope of this class.


# Advantages of SVMs

- Very good performance: though lately outshined by convolutional neural networks on some benchmarks (e.g., the MNIST digit recognition dataset), they often beat basically everything else.
- Theoretical guarantees about their generalization performance (accuracy for labeling test data) based on statistical learning theory.
- SVMs rely on convex optimization and do not get stuck in suboptimal local minima (neural networks have a big problem with these; similarly, decision trees rely on greedy search).
- Fairly robust to the curse of dimensionality → can effectively solve prediction problems with a large number of features.
- Flexible: can choose kernel to fit very complex decision boundaries.
- Will generally avoid overfitting with well-chosen parameters (but can certainly be overfitted for poorly chosen values, e.g., if  $C$  is too large).
- Classification of test points relies only on the support vectors → fast and memory efficient, especially when # of support vectors is small.

# Disadvantages of SVMs

- Training the model is computationally expensive – dependent on # of support vectors, but typically quadratic to cubic in the number of data points.
- Sensitive to the choice of parameters, particularly the constant  $C$  and kernel bandwidth ( $\gamma$  for RBF kernel in sklearn).
  - $C$  trades off the misclassification rate against the simplicity of the decision surface.  
Low  $C \rightarrow$  smooth decision surface; High  $C \rightarrow$  more training examples classified correctly.
  - Larger  $\gamma$  = lower bandwidth (increased weight on nearest training examples).
  - Proper choice of  $C$  and  $\gamma$  is critical to the SVM's performance.
  - For sklearn, use GridSearchCV with  $C$  and  $\gamma$  spaced exponentially far apart.
- Not much interpretability for non-linear SVM: it can enumerate the support vectors or (in low dimensions) visualize the decision boundary, but actually obtaining these involves calling a black-box optimization routine.

# Non-Linear SVM Parameters



# Lab Time

# For the Next Week (Week 5)

No class on February 19, 2024 (Presidents Day).

1. Check readings (optional) and review them
2. Assignment 1  
Due: February 16, 2024 (11:59pm)



# References

1. Scikit-learn documentation: <http://scikit-learn.org/stable/modules/svm.html>
2. C.J.C. Burges. A tutorial on support vector machines for pattern recognition. *Data Mining & Knowledge Discovery*, 2: 955-974, 1998.
3. A.W. Moore. Support Vector Machines (tutorial slides). <https://www.cs.cmu.edu/~./awm/tutorials/svm.html>
4. V. Vapnik. *Statistical Learning Theory*. Wiley: 1998.
5. T.-F. Wu, C.-J. Lin, and R.C. Weng. Probability estimates for multi-class classification by pairwise coupling. *Journal of Machine Learning Research* 5: 975-1005, 2004.
6. A.J. Smola and B. Schölkopf. A tutorial on support vector regression, *Statistics and Computing*, 2003. <http://alex.smola.org/papers/2003/SmoSch03b.pdf>
7. B. Schölkopf et al. Estimating the support of a high-dimensional distribution. *Neural Computation* 13: 1443-1471, 2001.
8. Berwick. 2009. An Idiot's guide to Support vector machines (SVMs)